

Accelerate serverless deployments using HashiCorp Terraform on AWS



Anton Aleksandrov

Pr. Solutions Architect, Serverless
AWS



Debasis Rath

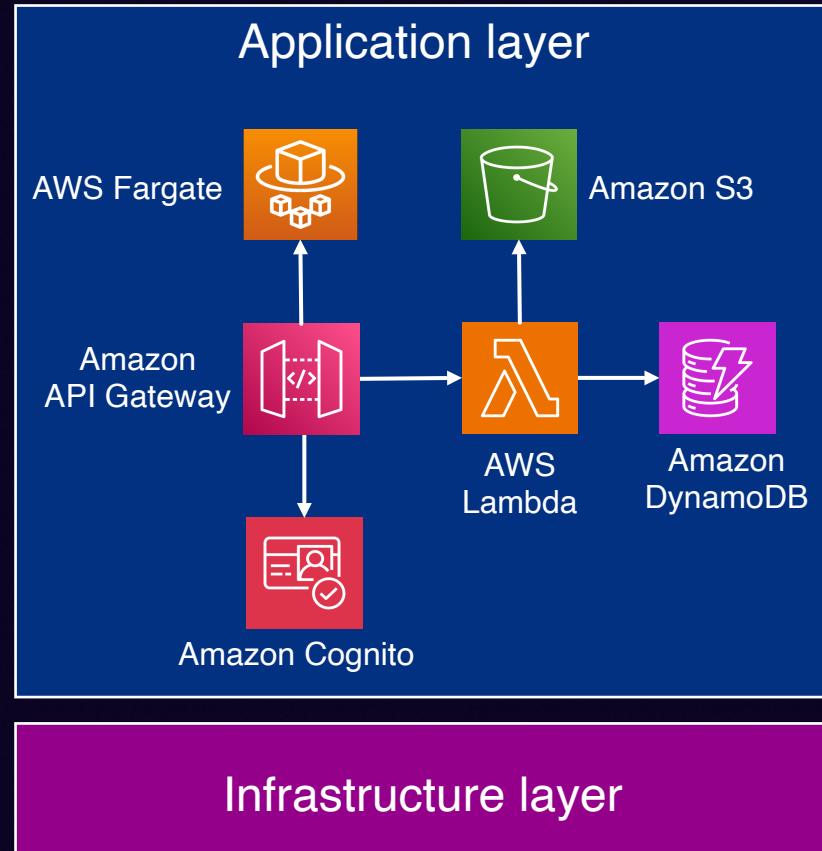
Sr. Solutions Architect, Serverless
AWS



Gautam Baghel

Partner Solutions Engineer
HashiCorp

What's so special about Serverless applications



- “**Infrastructure**” is redefined.
- A function, an event-source mapping, an event routing rule are **app resources** owned by the app team
- Updating app resources with IaC tools is commonly a part of **application developer's responsibilities**.

The ownership boundaries

This is an **application resource**! We need the flexibility to control it!



Application development teams

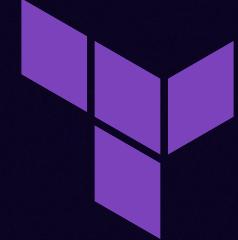


This is an **infrastructure resource**! We set the rules!



Infrastructure/ops team

Building Serverless Applications with Terraform



HashiCorp
Terraform



AWS Serverless

- A widely used IaC framework
- 450M+ downloads, 4,350+ customers
- A vast array of integrations, with support for over **3,000** providers
- Terraform AWS Provider is downloaded over **3 billion** times

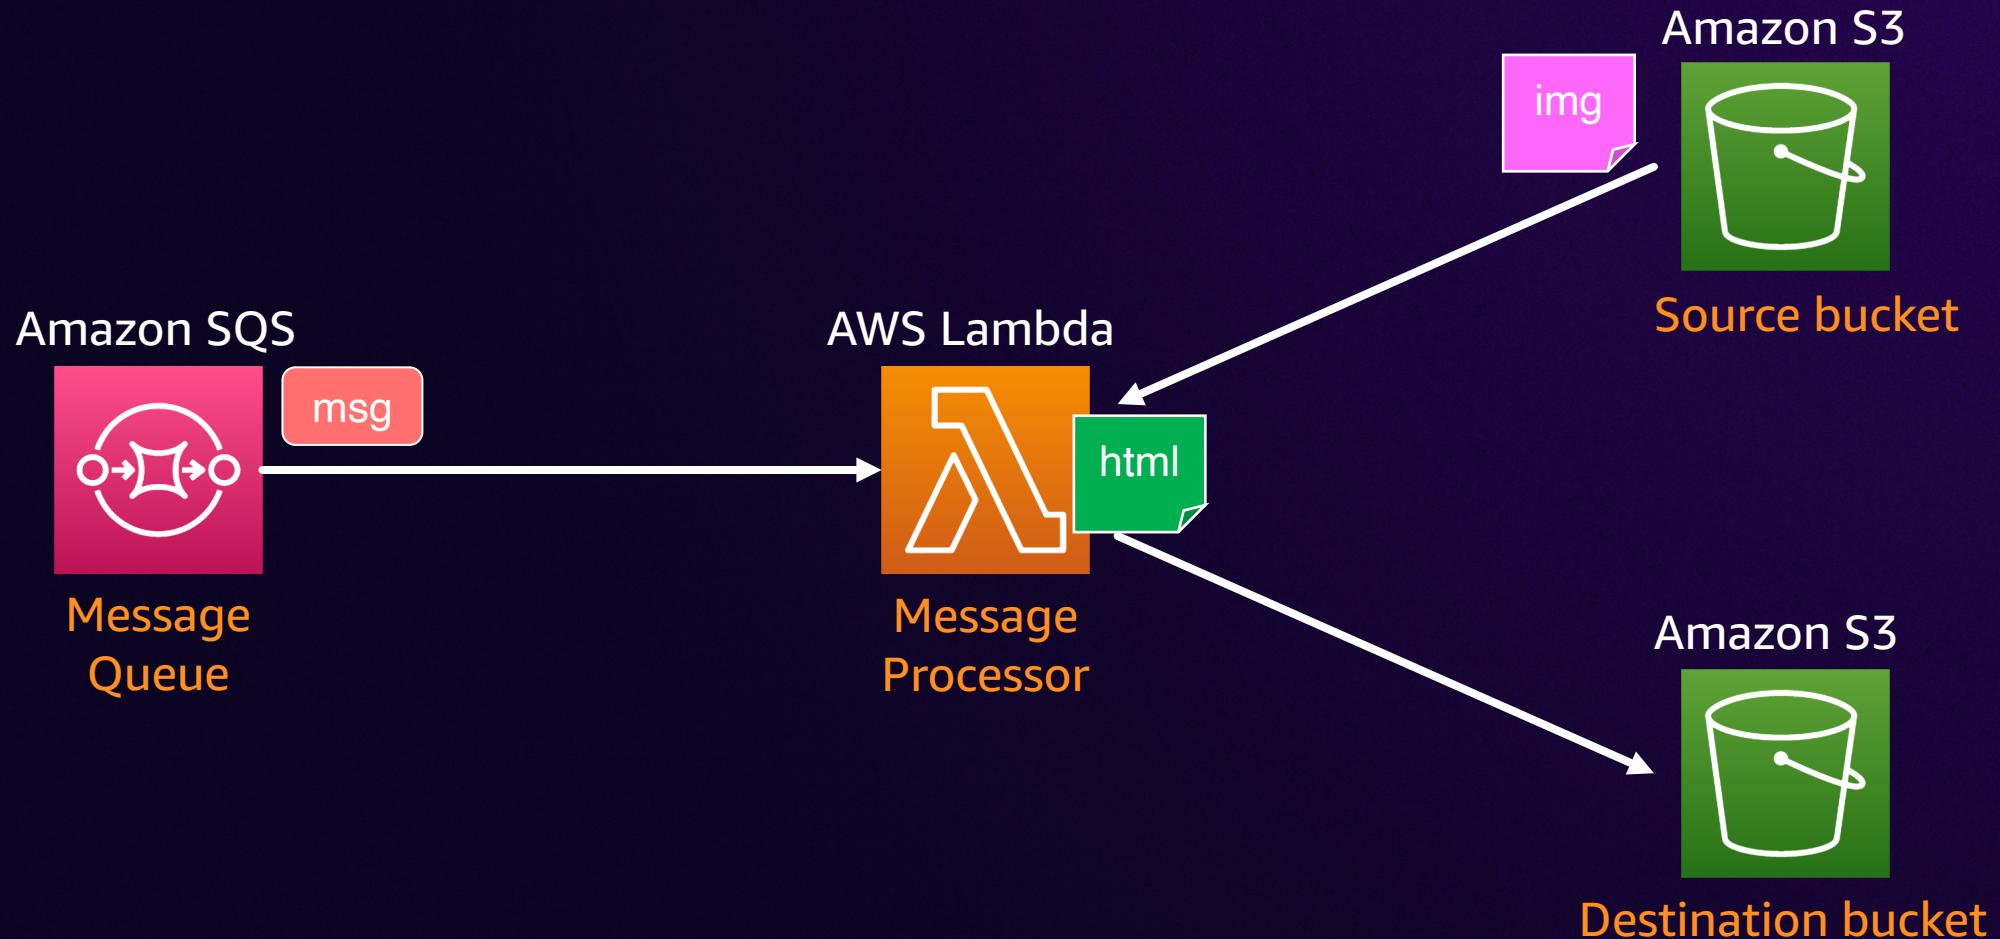
- Run code without thinking about servers or clusters
- Over **one million** customers are using AWS Serverless services every month
- Processes **10s of trillions** requests every month

Let's start building



© 2024, Amazon Web Services, Inc. or its affiliates. All rights reserved.

What are we building





Creating S3 buckets

```
resource "aws_s3_bucket" "src_bucket" {  
    bucket = "SourceBucket"  
    tags = {  
        environment = "Production"  
    }  
}
```

Bucket name

```
resource "aws_s3_bucket_ownership_controls" "src_bucket_ownership_controls" {  
    bucket = aws_s3_bucket.src_bucket.id  
    rule {  
        object_ownership = "BucketOwnerPreferred"  
    }  
}
```

Ownership controls

```
resource "aws_s3_bucket_acl" "src_bucket_acl" {  
    depends_on = [aws_s3_bucket_ownership_controls.src_bucket_ownership_controls]  
    bucket     = aws_s3_bucket.src_bucket.id  
    acl        = "private"  
}
```

Access controls

Creating SQS queues



```
resource "aws_sqs_queue" "this" {  
    name = "greetings_queue"  
    sqs_managed_sse_enabled = true  
}
```

Queue name



Creating Lambda Function Execution Role

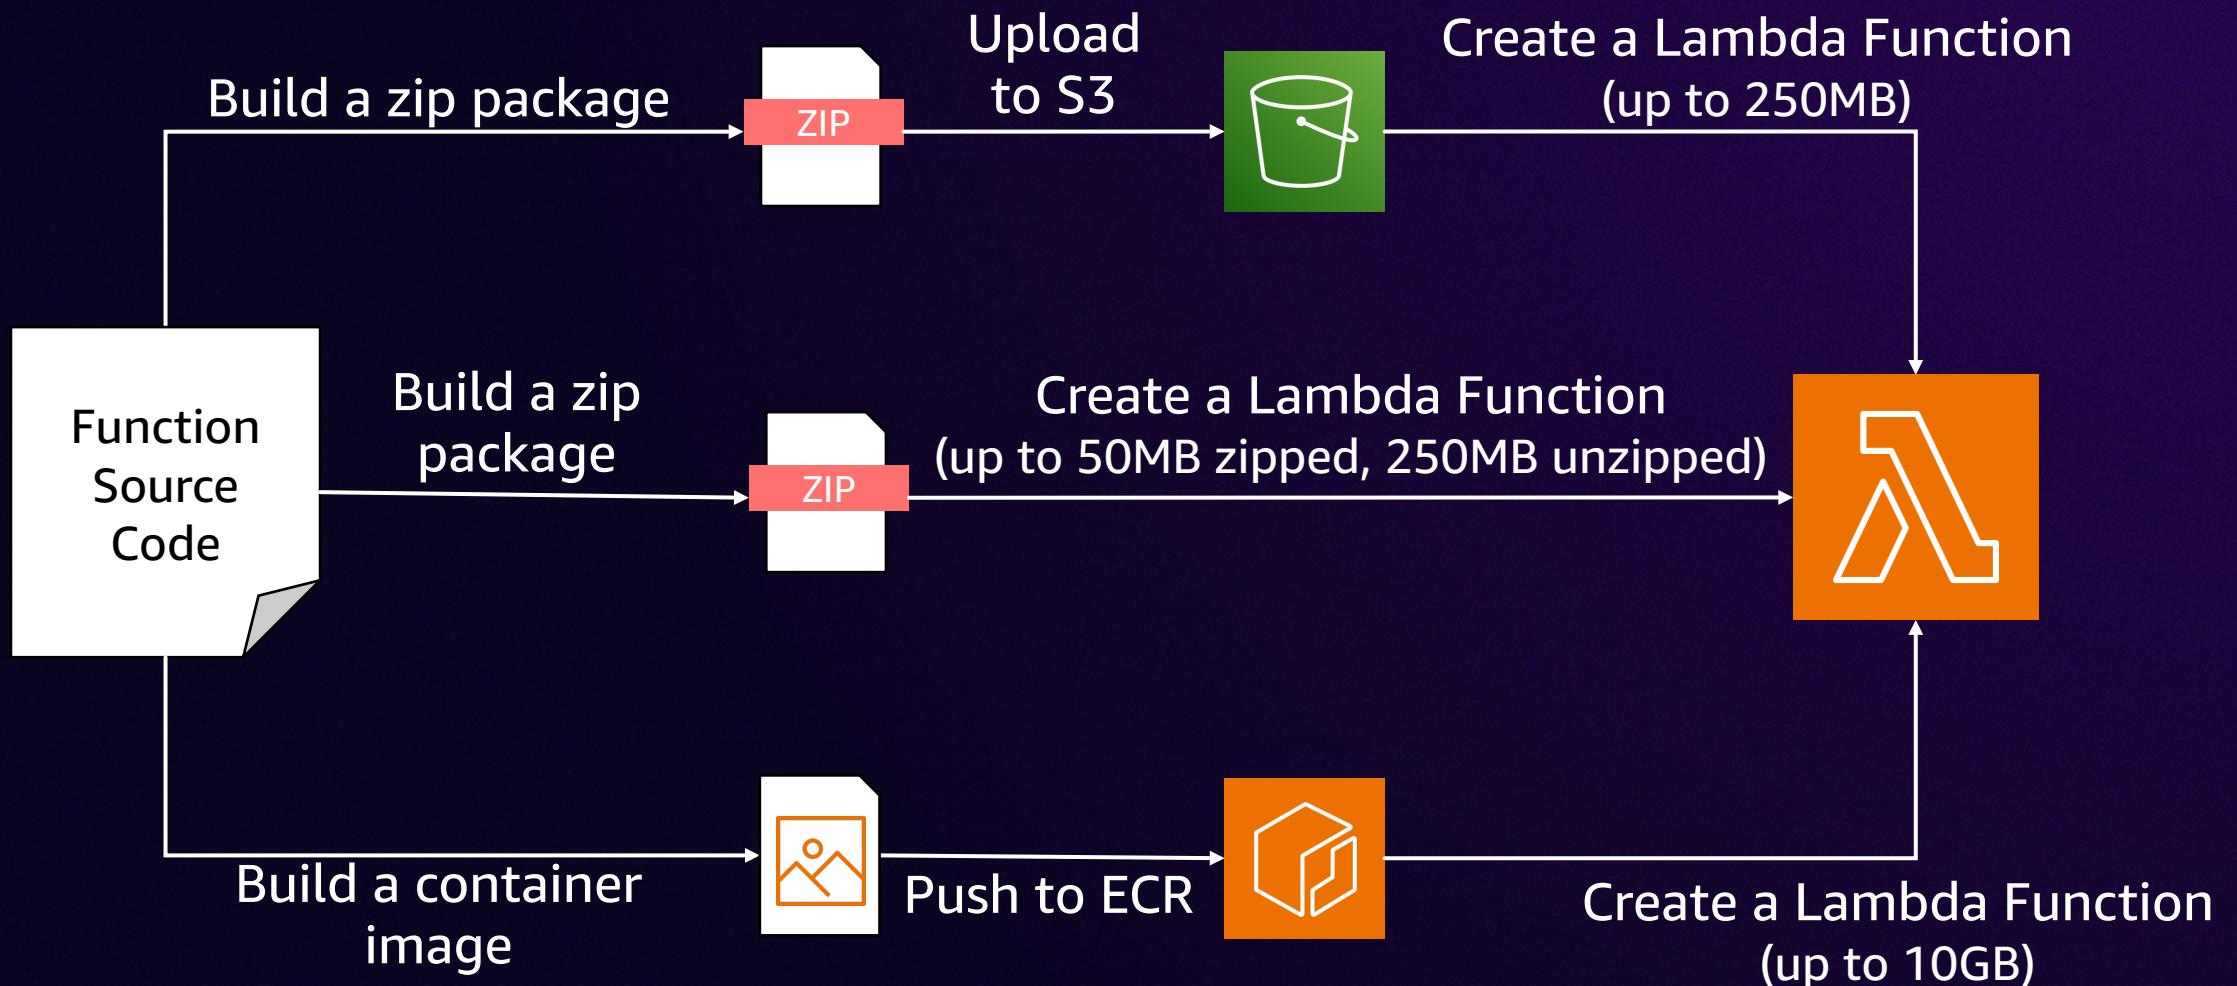


```
resource "aws_iam_policy" "this" {
  policy = jsonencode({
    Version = "2012-10-17"
    Statement = [
      {
        Action    = ["s3:GetObject"]
        Effect   = "Allow"
        Resource = aws_s3_bucket.src_bucket.arn
      },
      {
        Action    = ["s3:PutObject"]
        Effect   = "Allow"
        Resource = aws_s3_bucket.dst_bucket.arn
      },
    ]
  })
}
```

The principle of least privilege



Creating Lambda Functions





Creating Lambda Functions

```
data "archive_file" "this" {  
    type      = "zip"  
    source_dir = "../src/lambda/greetings"  
    output_path = "greetings_lambda.zip"  
}
```

Package source directory

```
resource "aws_lambda_function" "this" {  
    function_name = "greetings_lambda"  
    role          = aws_iam_role.this.arn  
  
    handler      = "index.handler"  
    runtime       = "nodejs20.x"  
    memory_size   = 256
```

Function configuration

```
filename        = data.archive_file.this.output_path  
source_code_hash = data.archive_file.this.output_base64sha256
```

Function code



Creating Lambda Functions

```
resource "aws_lambda_function" "this" {
    function_name = "greetings_lambda"
    role          = aws_iam_role.this.arn

    handler      = "index.handler"
    runtime       = "nodejs20.x"
    memory_size   = 256

    filename      = data.archive_file.this.output_path
    source_code_hash = data.archive_file.this.output_base64sha256

    # Alternatively, if getting the zip file from S3
    # s3_bucket = .....
    # s3_key    = .....

    #Alternatively, if using function image from ECR
    # image_uri = .....

}
```

Function configuration

Function code





Creating Lambda Functions

```
resource "aws_lambda_function" "this" {
    function_name = "greetings_lambda"
    role          = aws_iam_role.this.arn

    handler      = "index.handler"
    runtime       = "nodejs20.x"
    memory_size   = 256

    filename      = data.archive_file.this.output_path
    source_code_hash = data.archive_file.this.output_base64sha256

    # Alternatively, if getting the zip file from S3
    # s3_bucket = .....
    # s3_key    = .....

    #Alternatively, if using function image from ECR
    # image_uri = .....

}
```

Function configuration

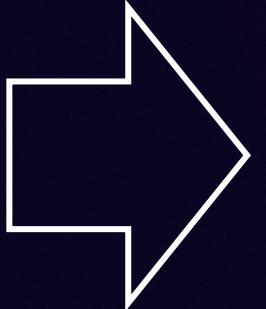
Function code



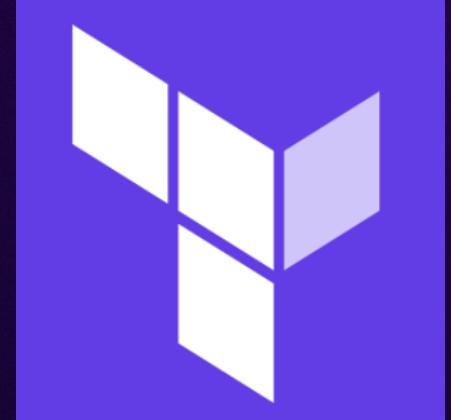
AWS SAM CLI Support for Terraform



Local development,
debugging, and testing



Seamless
development
and testing
experience



Cloud environment
templating and deployment

<https://aws.amazon.com/blogs/compute/aws-sam-support-for-hashicorp-terraform-now-generally-available/>



© 2024, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Local testing with AWS SAM

```
 sam - "ip-172-31- X +  
Admin:~/environment/serverless-tf-basic-app/terraform (master) $ echo -e "\nAWS CLI Version:\n"; aws --version; echo -e "\n\nTerraform Version:\n"; terraform --version; echo -e "\n\nAWS SAM CLI Version:\n"; sam --version  
  
AWS CLI Version:  
  
aws-cli/2.17.47 Python/3.11.9 Linux/6.1.106-116.188.amzn2023.x86_64 exe/x86_64.amzn.2023  
  
Terraform Version:  
  
Terraform v1.9.5  
on linux_amd64  
+ provider registry.terraform.io/hashicorp/archive v2.4.2  
+ provider registry.terraform.io/hashicorp/aws v5.54.1  
+ provider registry.terraform.io/hashicorp/random v3.6.2  
  
AWS SAM CLI Version:  
  
SAM CLI, version 1.112.0  
Admin:~/environment/serverless-tf-basic-app/terraform (master) $ █
```

Scaling



© 2024, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Growth challenges...

Use approved runtimes only

**Upload Zip
directly or via S3?**

**Different runtimes
have different
packaging tools**

**Uniformly apply
tags**

Project structure

**Applying best-
practices and
policies at scale**

**Principle of least
privileged access**

Do not reinvent the wheel



Modularization

DON'T

```
38016
38017   resource "aws_iam_role_policy_attachment"
38018     policy_arn = aws_iam_policy.greeting_lam
38019     role       = aws_iam_role.greeting_lambda
38020   }
38021
38022   resource "aws_lambda_event_source_mapping"
38023     event_source_arn = var.greeting_queue_ar
38024     function_name    = aws_lambda_function.g
38025     batch_size        = 1
38026
38027     depends_on = [aws_iam_role_policy_attach
38028   }
38029   # EOF
```

DO

Storage

Web Backend

Observability

Async
processing

Put all the IaC templates
in **one file**

Organize your IaC as a
collection of **reusable modules**

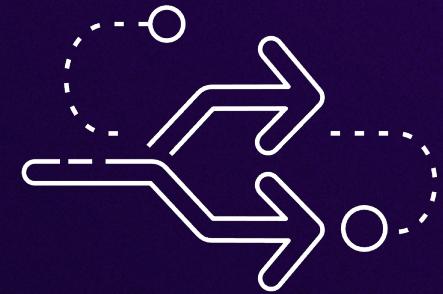
Modularization



Best practices



Reusability



Composability

A Terraform Module

Input (variables.tf)

```
variable "function_name" {  
    type = string  
}  
  
variable "runtime" {  
    type = string  
    default = "nodejs20.x"  
}
```

Cloud resources configuration (main.tf)

```
resource "aws_lambda_function" "lambda_function" {  
    function_name      = var.function_name  
    runtime           = var.runtime
```

A Terraform Module

Input (variables.tf)

```
variable "function_name" {  
    type = string  
}  
  
variable "runtime" {  
    type = string  
    default = "nodejs20.x"  
}  
  
variable "log_retention_days" {  
    type = number  
    default = 14  
}
```

Cloud resources configuration (main.tf)

```
resource "aws_lambda_function" "lambda_function" {  
    function_name      = var.function_name  
    runtime           = var.runtime  
    logging_config {  
        log_group = aws_cloudwatch_log_group.this.arn  
    }  
}  
  
resource "aws_cloudwatch_log_group" "this" {  
    name              = "/aws/lambda/${var.function_name}"  
    retention_in_days = var.log_retention_days  
}  
  
resource "aws_iam_policy" "this" {  
    ....  
}
```

A Terraform Module

Input (variables.tf)

```
variable "function_name" {
```

```
    resource "aws_lambda_function" "products" {
        function_name = local.function_name
        logging_config {
            application_log_level = "INFO"
            system_log_level      = "INFO"
            log_format              = "JSON"
            log_group                = aws_cloudwatch_log_group.lambda_log_group.name
        }
        tracing_config {
            mode = "Active"
        }
    }
}
```

Cloud resources configuration (main.tf)

```
resource "aws_lambda_function" "lambda_function" {
    function_name = var.function_name
```

A Terraform Module

```
Input  
variable "f  
    type =  
}  
  
variable "r  
    type =  
    default  
}  
  
variable "l  
    type =  
    default  
}  
  
resource "aws_lambda_function" "products" {  
    function_name = local.function_name  
  
    handler      = "index.handler"  
    runtime      = "nodejs20.x"  
    memory_size = 512  
    role         = aws_iam_role.lambda_role.arn  
    layers       = [local.powertools_layer_arn]  
  
    environment {  
        variables = {  
            POWERTOOLS_SERVICE_NAME      = "${var.resource_name_prefix}-products",  
            POWERTOOLS_METRICS_NAMESPACE = "${var.resource_name_prefix}-products",  
            # Below properties can help with debugging  
            POWERTOOLS_LOGGER_LOG_EVENT = false, # Logs incoming event  
            POWERTOOLS_LOG_LEVEL        = "INFO", # Changes log level  
            POWERTOOLS_DEV              = false # Increases JSON indentation  
        }  
    }  
}
```

main.tf)
unction" {

.this.arn

s" {
.function_name}"
days

A Terraform Module

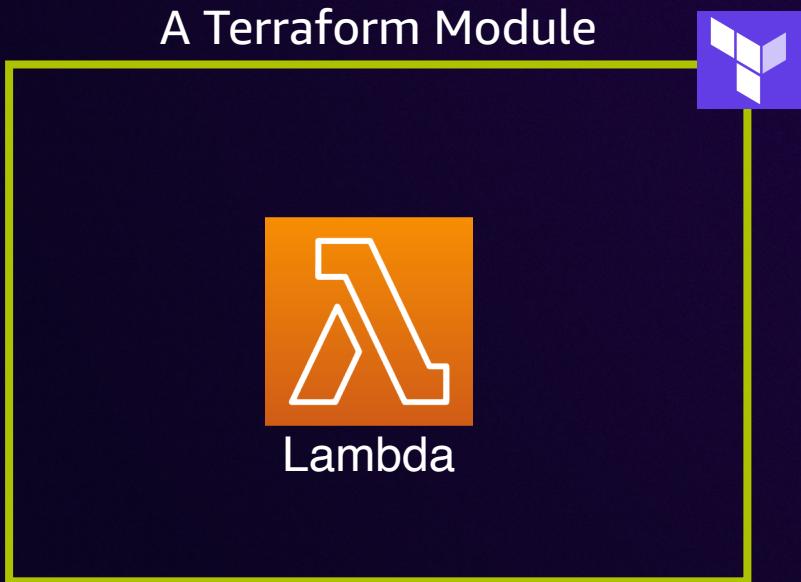
Input (variables.tf)

```
variable "function_name" {  
    type = string  
}  
  
variable "runtime" {  
    type = string  
    default = "nodejs20.x"  
}  
  
variable "log_retention_days" {  
    type = number  
    default = 14  
}
```

Cloud resources configuration (main.tf)

```
resource "aws_lambda_function" "lambda_function" {  
    function_name      = var.function_name  
    runtime           = var.runtime  
    logging_config {  
        log_group = aws_cloudwatch_log_group.this.arn  
    }  
}  
  
resource "aws_cloudwatch_log_group" "this" {  
    name              = "/aws/lambda/${var.function_name}"  
    retention_in_days = var.log_retention_days  
}  
  
resource "aws_iam_policy" "this" {  
    ....  
}
```

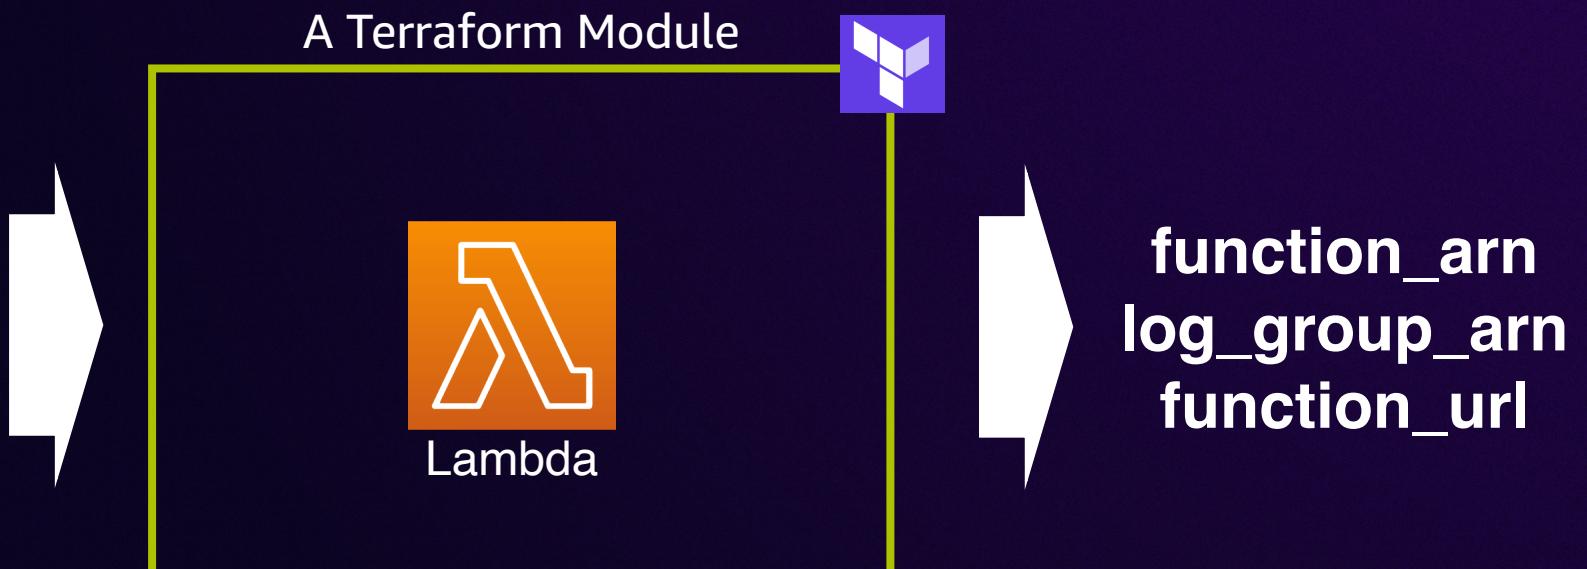
A Terraform Module



Best Practices

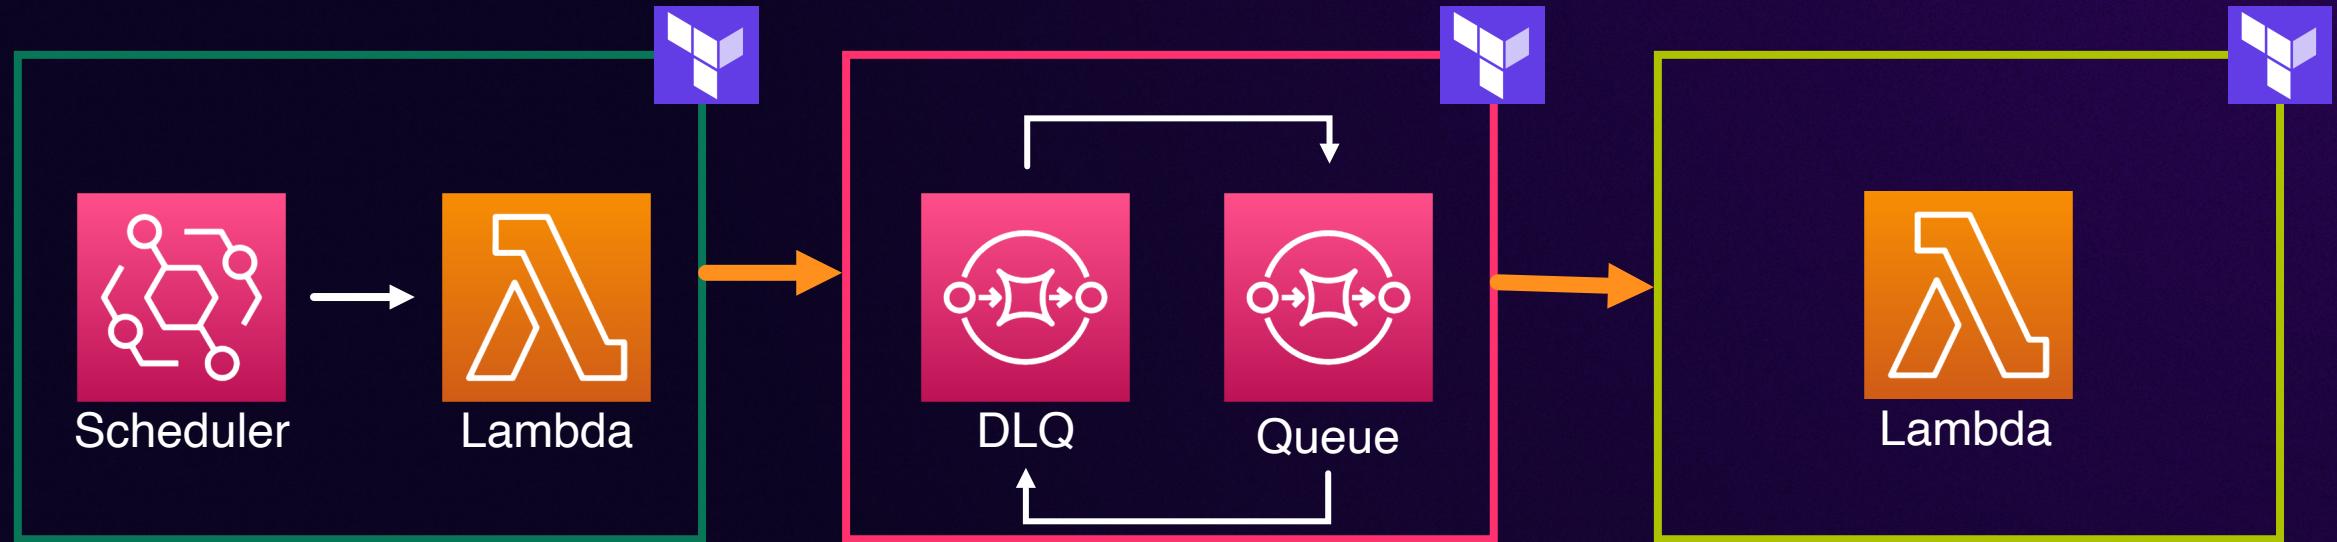
Variables

`function_name`
`memory_size`
`source_path`
`runtime`
`handler`
`attach_vpc`
`log_retention_days`
`enable_furl`
`etc...`



Outputs

Reusability and composability

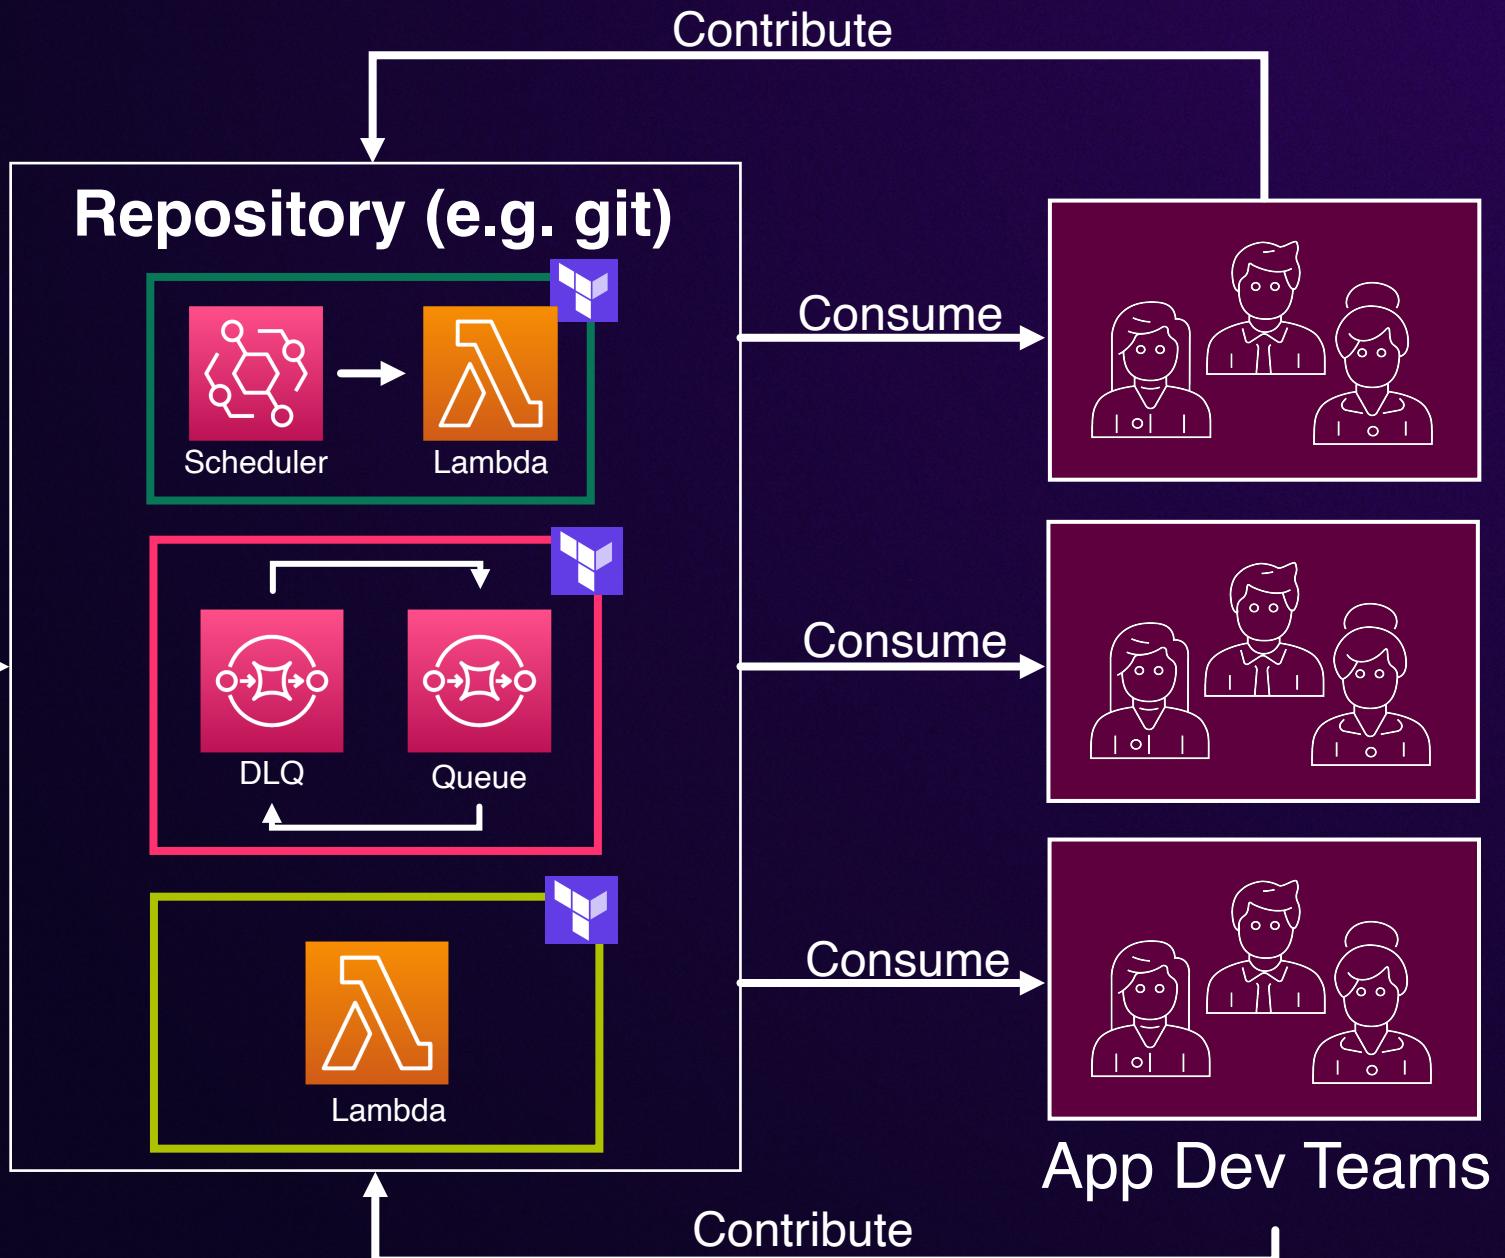


**Lambda function
with periodical
scheduler**

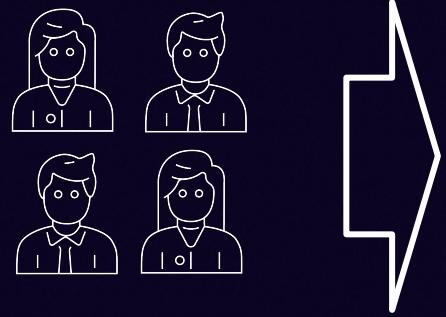
**SQS Queue with
redrive**

**Baseline Lambda
function**

Reusability



Evolve your IaC with Terraform Modules

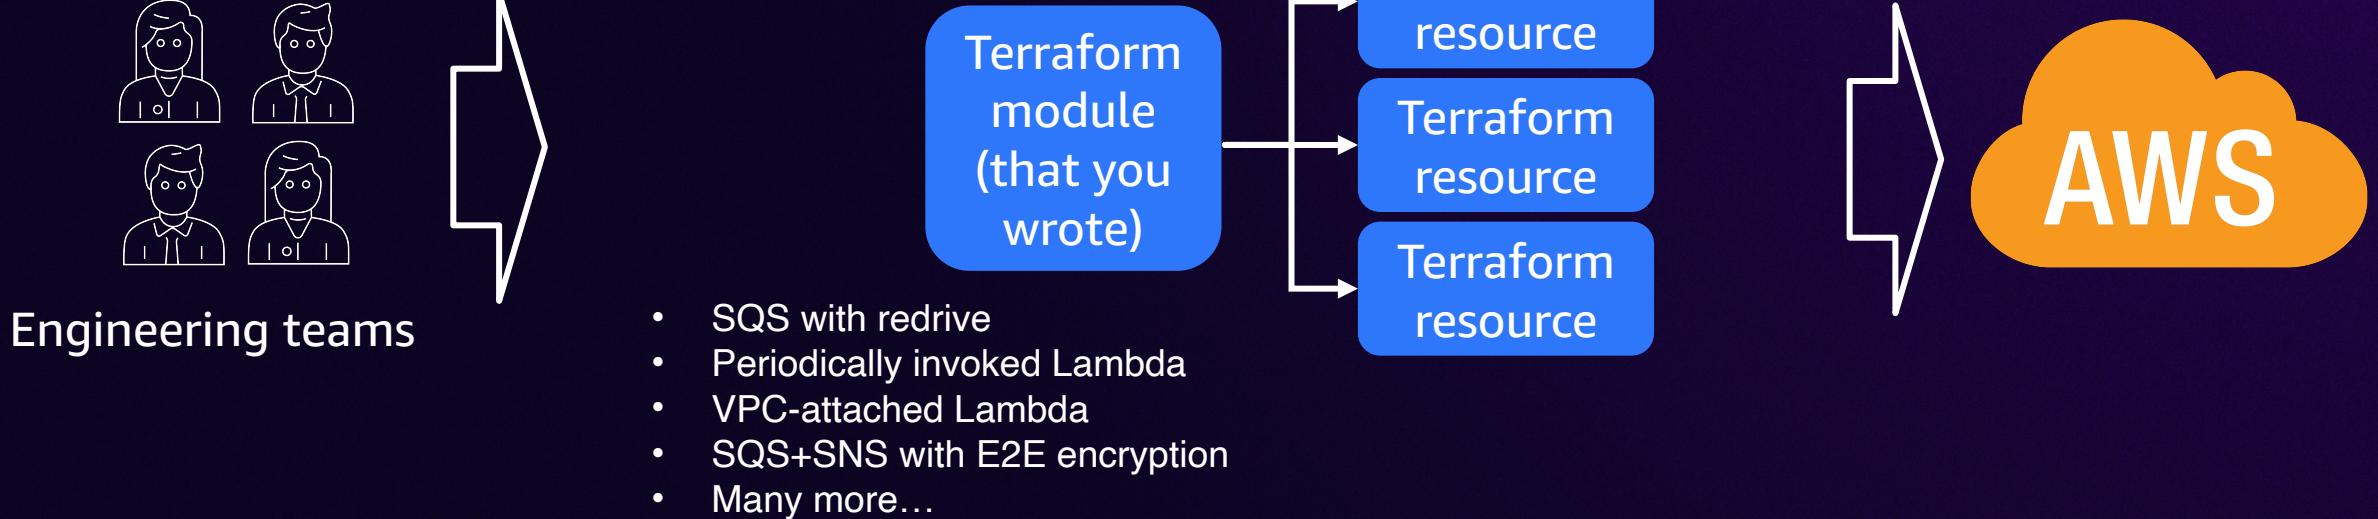


Engineering teams



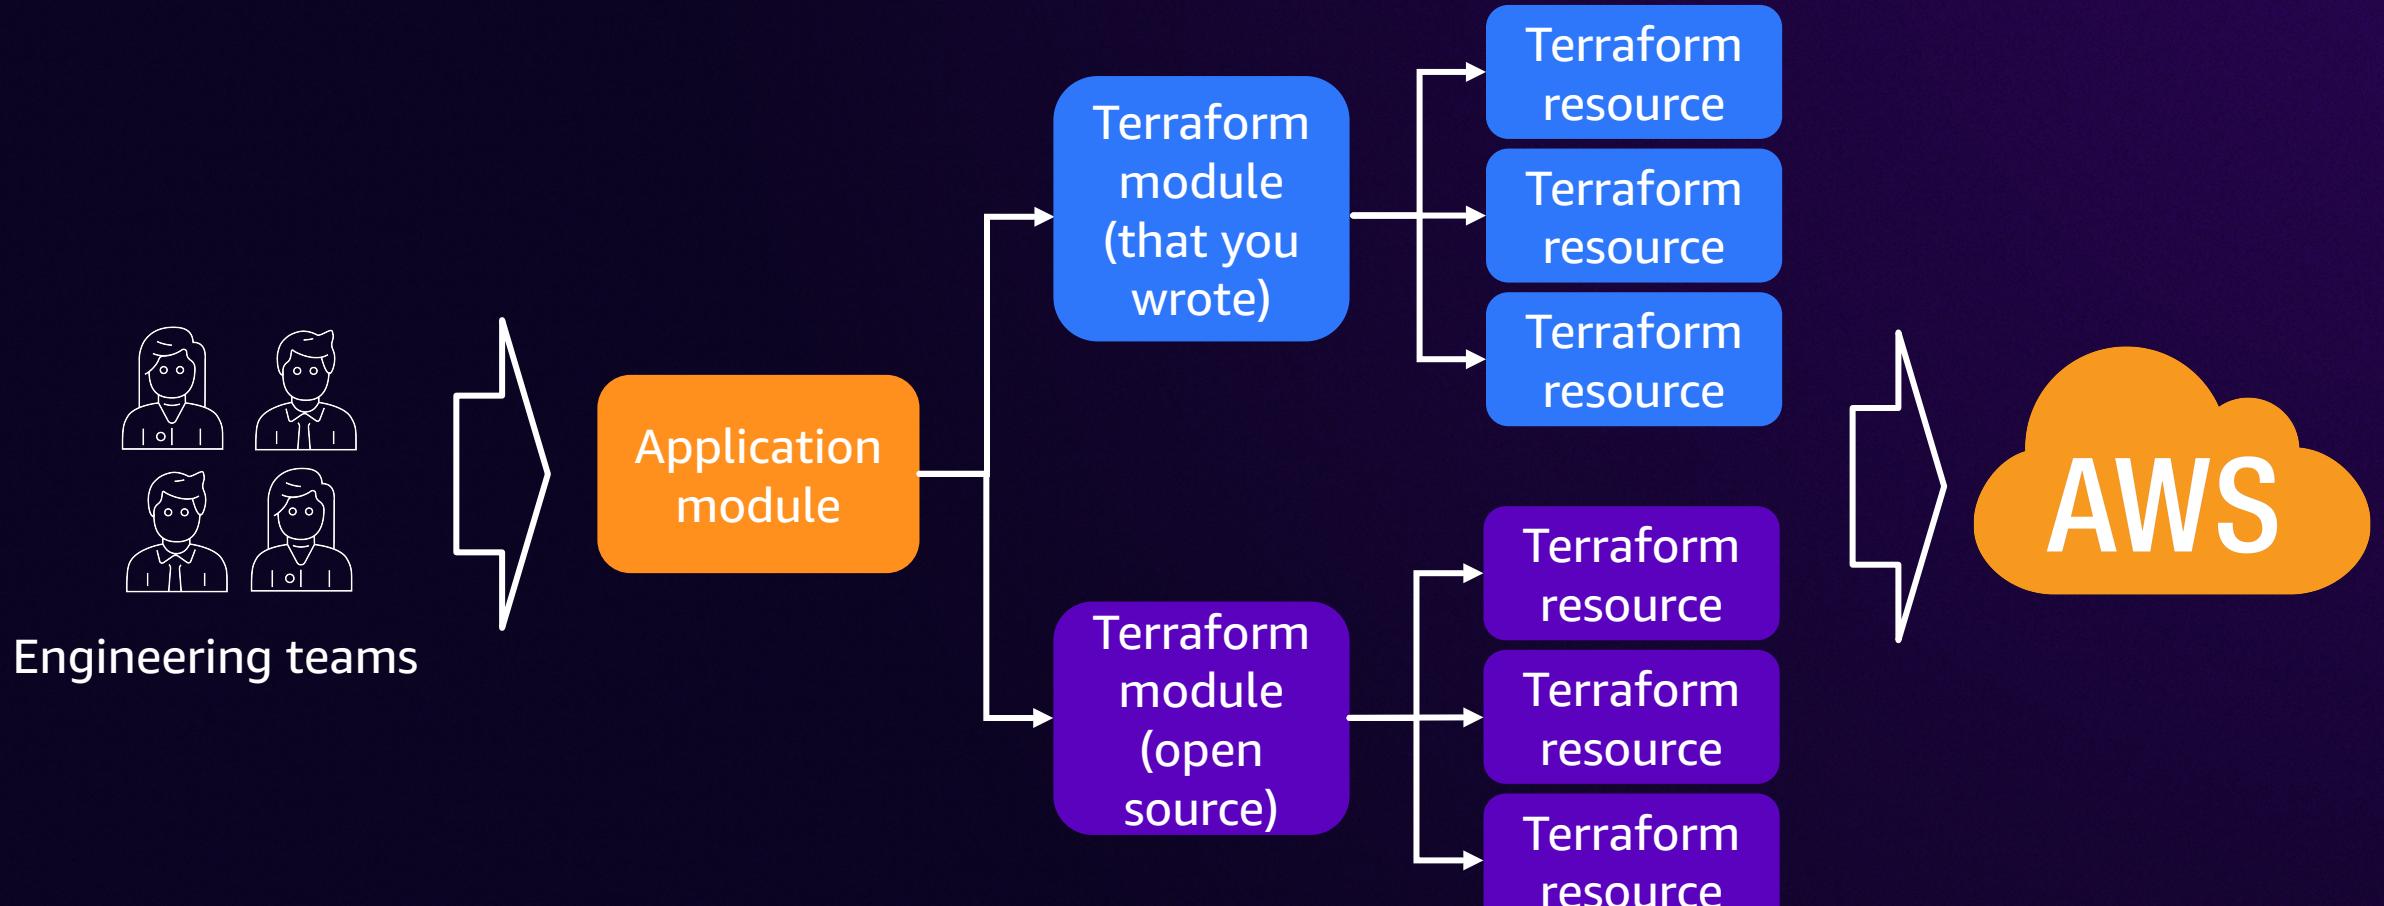
Maturity level: **Beginner**

Evolve your IaC with Terraform Modules



Maturity level: **Advanced**

Evolve your IaC with Terraform Modules



Maturity level: Power User

Serverless.tf – community project

An opinionated, 100% open-source community framework for developing, building, deploying, and securing serverless applications on AWS using Terraform.



Over **120 million** downloads

<https://serverless.tf>

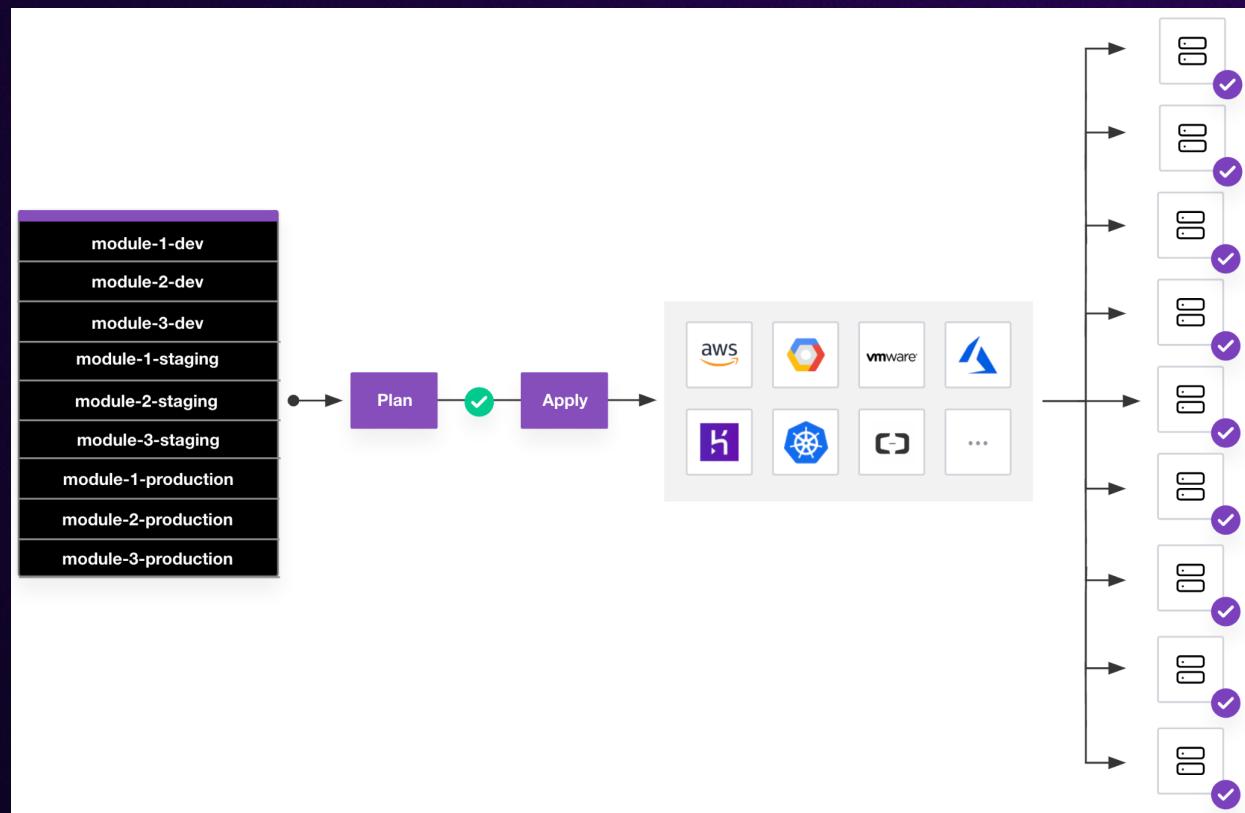
Consistent Deployments

The challenge – managing workspace complexity

HashiCorp recommends a **modular and composable** approach to structuring workspaces to optimize Terraform configurations for **speed, security, and maintainability**.

As organizations scale their Terraform usage with more complex workloads they begin to encounter challenges:

- Managing dependencies across **multiple workspaces**
- Repeating and managing **multiple instances of the same infrastructure**

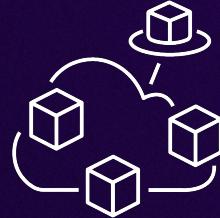


Common challenges

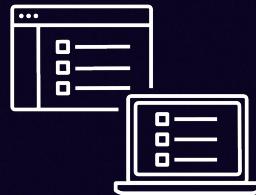
Deploy applications components reliably and consistently



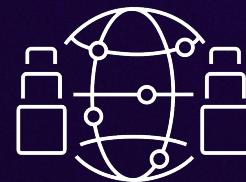
Multiple teams



Multiple environments



Multiple accounts

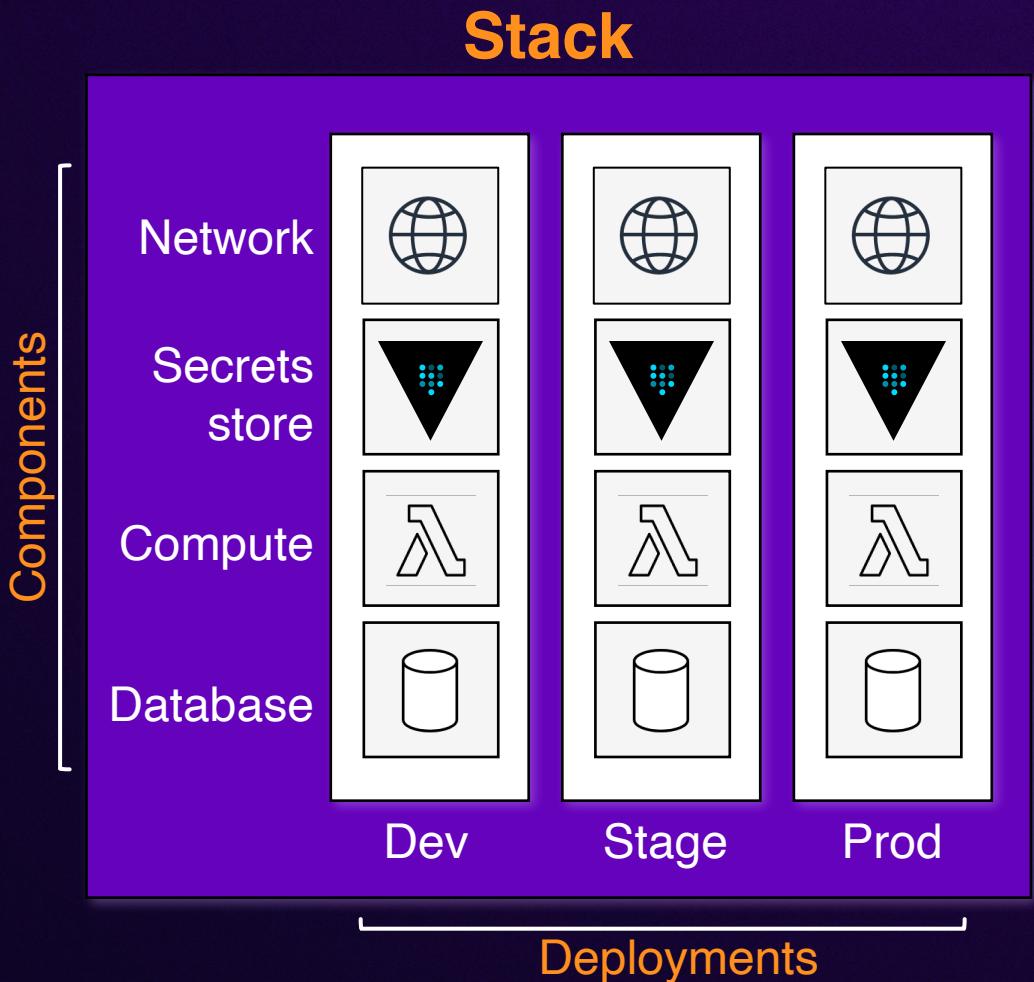


Multiple regions

Introducing Terraform Stacks

A new approach to provisioning and management of cloud application environments

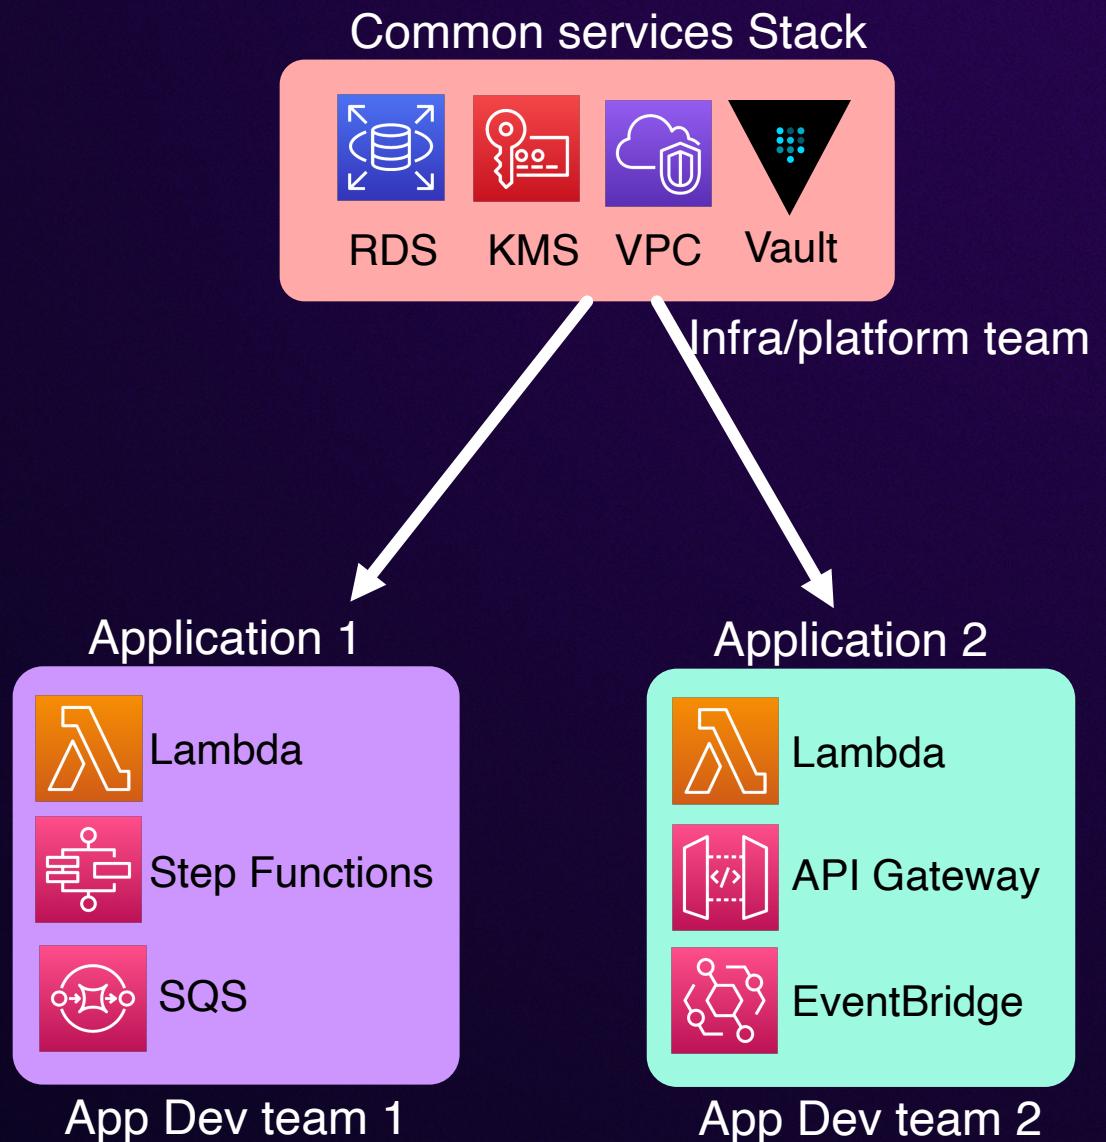
- Optimize the **coordination, provisioning, and management** of dependent Terraform configurations
- **Components**: standard Terraform modules with **dependencies defined in code**
- **Deployments**: used to repeat the components in the stack
- **Customize multiple deployments** with input variables, without duplicating code



Connecting Stacks via outputs

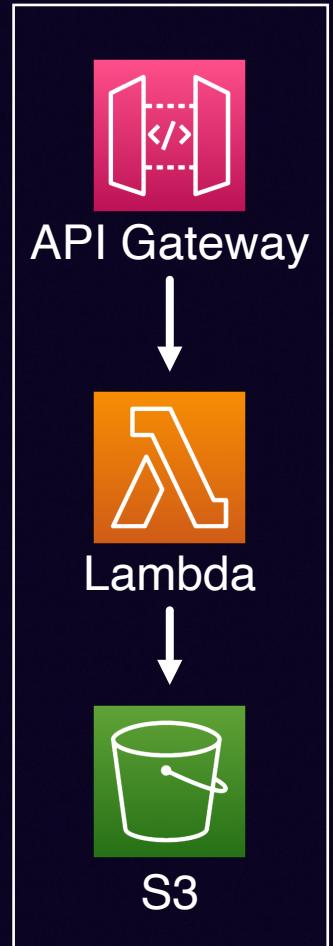
Separate stacks for based on persona

- **Outputs from one stack** deployment can be used as **inputs to another stack** deployment.
- **When an input in a downstream stack changes** as a result of an output from an upstream stack, **HCP Terraform automatically triggers plans** for any affected deployments.



Sample Serverless Application Stack

Stack Components



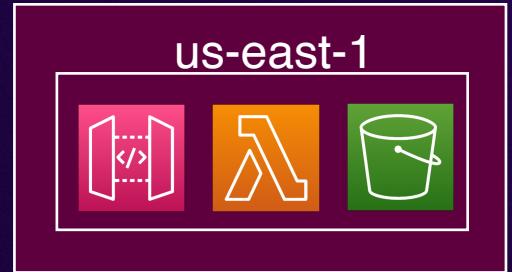
Deployment definitions

```
identity_token "aws" {
    audience = ["<Set to your AWS IAM assume-role audience>"]
}

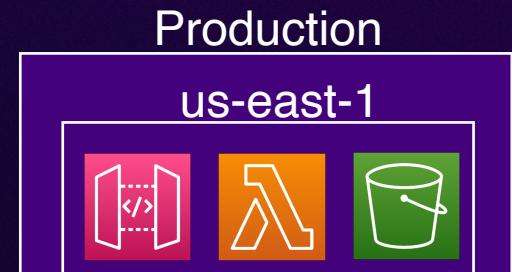
deployment "development" {
    variables = {
        regions          = ["us-east-1"]
        role_arn         = "<Your dev AWS account IAM role ARN>"
        identity_token_file = identity_token.aws.jwt_filename
    }
}

deployment "production" {
    variables = {
        regions          = ["us-east-1", "us-west-1"]
        role_arn         = "<Your prod AWS account IAM role ARN>"
        identity_token_file = identity_token.aws.jwt_filename
    }
}
```

Deployments in AWS Cloud



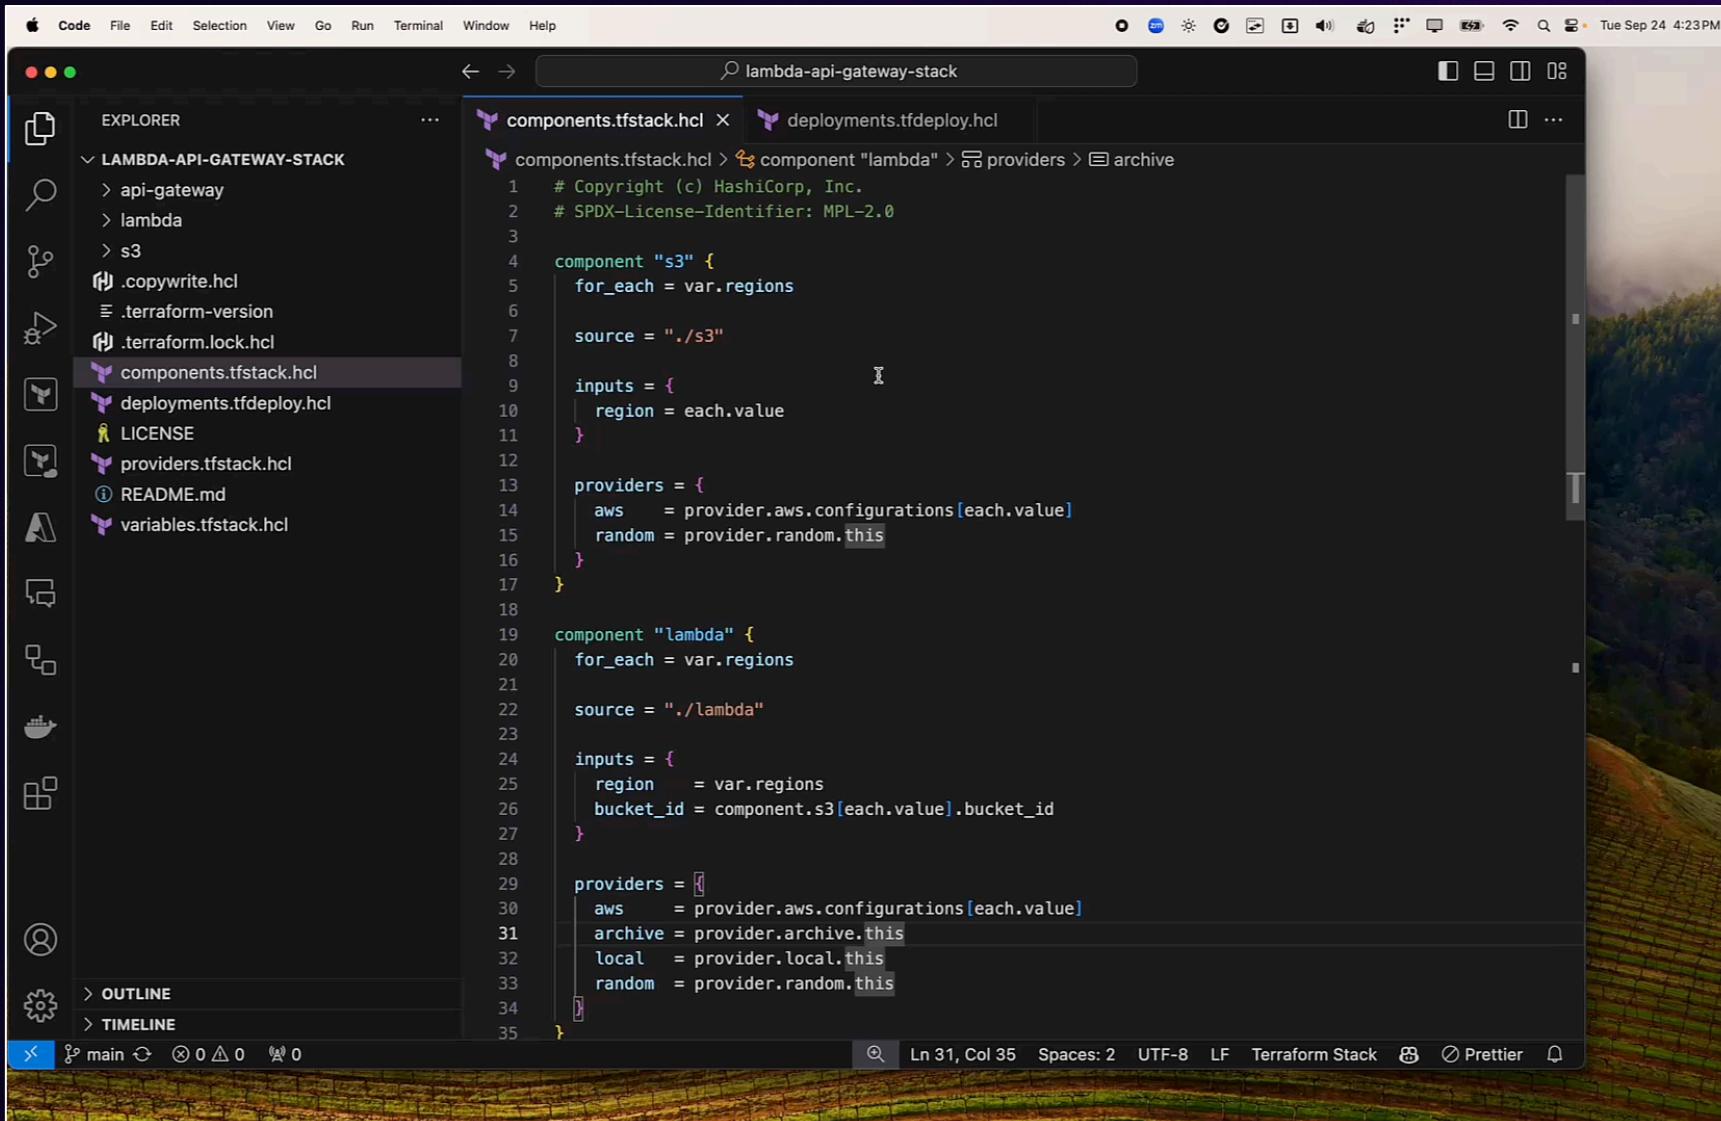
Development



Production



Serverless Deployments with Terraform Stacks



The screenshot shows a Mac OS X desktop environment with a dark-themed Code editor window open. The window title is "lambda-api-gateway-stack". The left sidebar of the editor shows a file tree for a directory named "LAMBDA-API-GATEWAY-STACK", which contains files like "api-gateway", "lambda", "s3", ".copywrite.hcl", ".terraform-version", ".terraform.lock.hcl", "components.tfstack.hcl" (which is selected), "deployments.tfdeploy.hcl", "LICENSE", "providers.tfstack.hcl", "README.md", and "variables.tfstack.hcl". The main editor area displays Terraform configuration code:

```
components.tfstack.hcl
component "s3" {
  for_each = var.regions
  source = "./s3"
  inputs = {
    region = each.value
  }
  providers = [
    aws = provider.aws.configurations[each.value]
    random = provider.random.this
  ]
}

component "lambda" {
  for_each = var.regions
  source = "./lambda"
  inputs = {
    region     = var.regions
    bucket_id = component.s3[each.value].bucket_id
  }
  providers = [
    aws = provider.aws.configurations[each.value]
    archive = provider.archive.this
    local   = provider.local.this
    random  = provider.random.this
  ]
}
```



Conclusion

1

Terraform + SAM enable you to efficiently build and test serverless applications with tools your organization already familiar with.

2

Modularizing your Terraform configurations and using existing OSS modules allows to build reusable components and patterns, so critical in the serverless world

3

Terraform Stacks and Workspaces help to streamline deployment of your serverless applications and resolve cross-team dependencies in a reliable way.

Next steps



**Building Serverless Applications
with Terraform **Guide****



**Building Serverless Applications
with Terraform **Workshop****



AWS Terraform Provider **docs**



Using **AWS SAM with Terraform**



**Governance for Serverless
Applications **Guide****



Sample **serverless patterns
with Terraform**

Thank you!



Anton Aleksandrov
Pr. Solutions Architect, Serverless
AWS



Debasis Rath
Sr. Solutions Architect, Serverless
AWS



Gautam Baghel
Partner Solutions Engineer
HashiCorp