



SERVERLESS EXPRESS LEAGUE

Building Multi-tenant SaaS Application using Serverless



Anubhav Sharma
Principal Solutions Architect
AWS SaaS Factory
 antonal80



Anubhav Sharma
Principal Solutions Architect
AWS SaaS Factory
 anubhavynr


SaaS architecture challenges



Tenant isolation




Observability



Noisy Neighbors



Identity Management



Tiering strategy



Cost per tenant

What do you need to build an application

Serverless Compute

Serverless Database

Serverless API Gateway

Serverless Orchestration

Serverless Messaging

Serverless Event Bus

Serverless Notifications

Serverless Observability



What is Serverless



No infrastructure provisioning,
no management



Automatic scaling

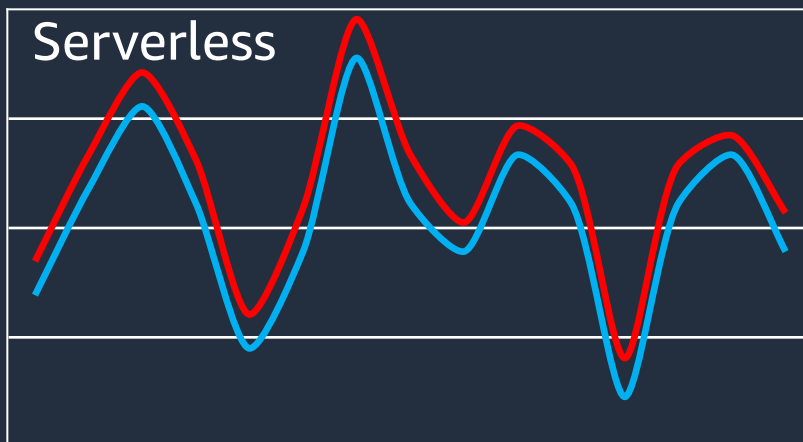
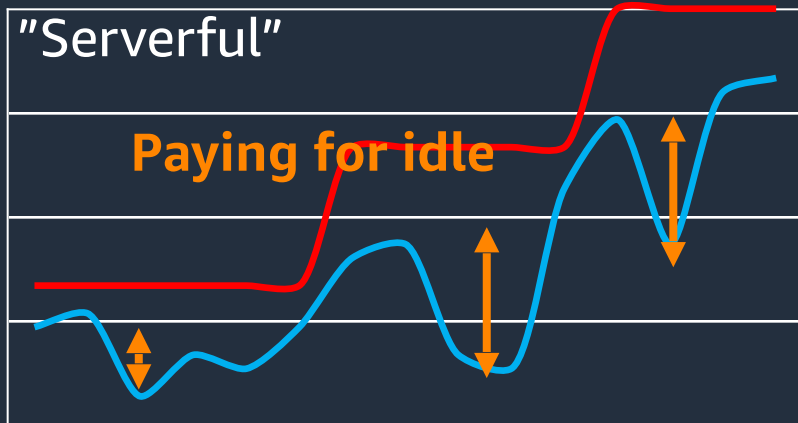
Pay for value




Highly available and secure



Serverless aligns SaaS costs with utilization



■ Capacity cost ■ Capacity utilization



Agility



Performance



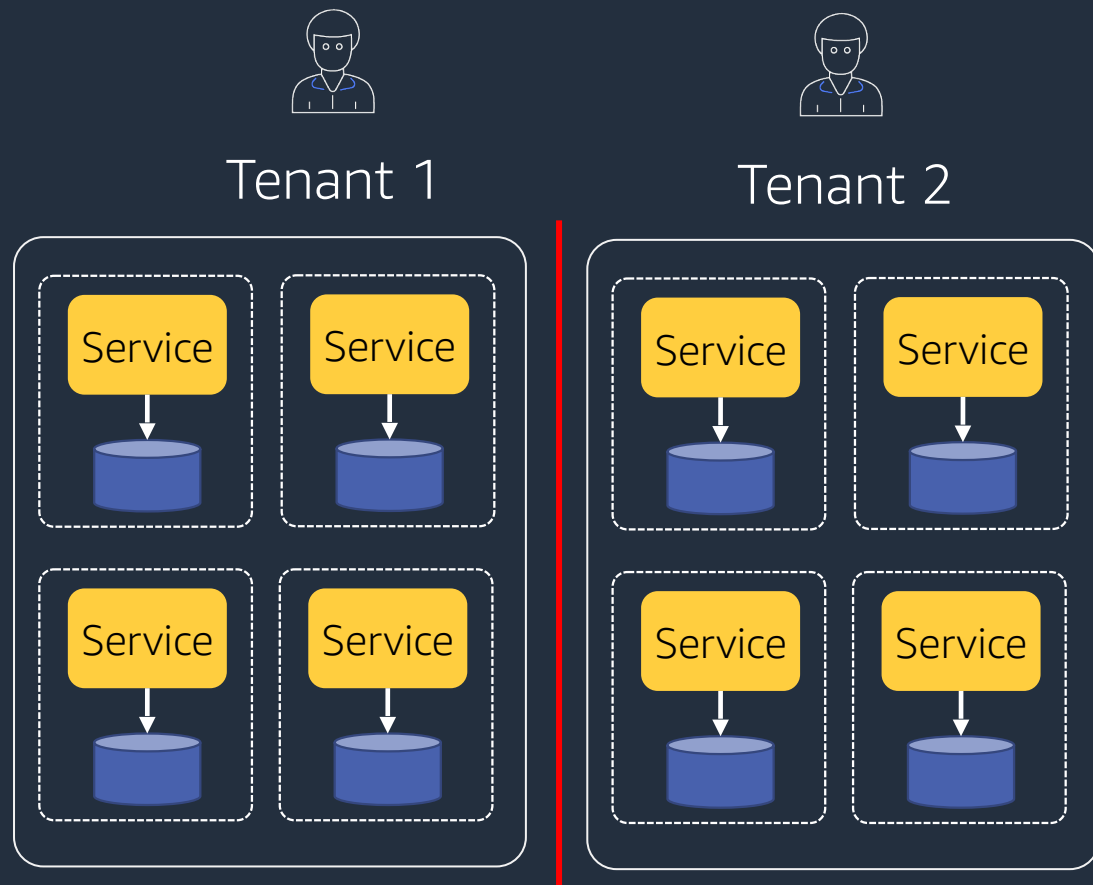
Cost



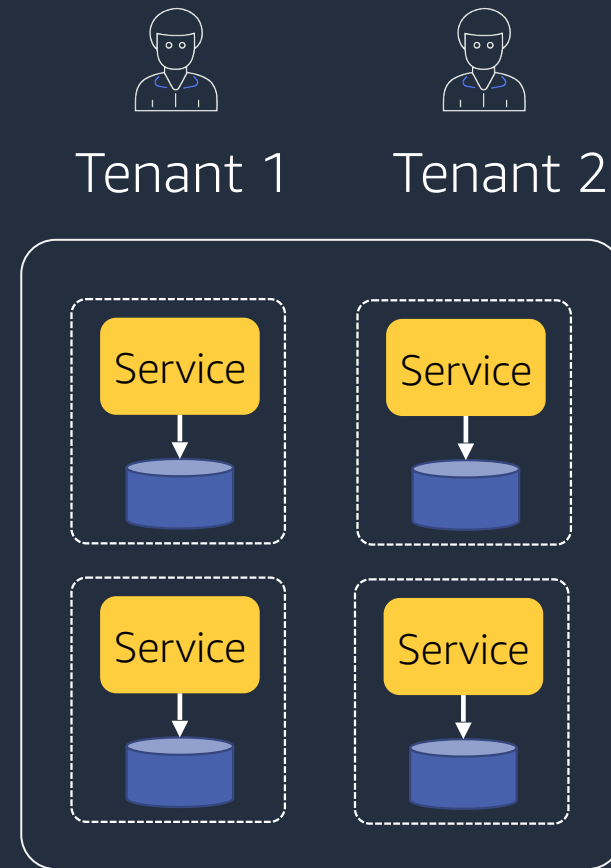
Security

Using Serverless allows you to have the best capacity cost to capacity utilization ratio.

SaaS deployment models



Dedicated resources for each tenant
(silo model)



Shared resources for all tenants
(pool model)

AWS Services & Features Used



API Gateway

Rest API, Lambda
Authorizer, Usage Plans,
API Keys



Cognito

User pools



Lambda

Fine-grained access
control with STS, Layers




DynamoDB

Data persistence



SAM



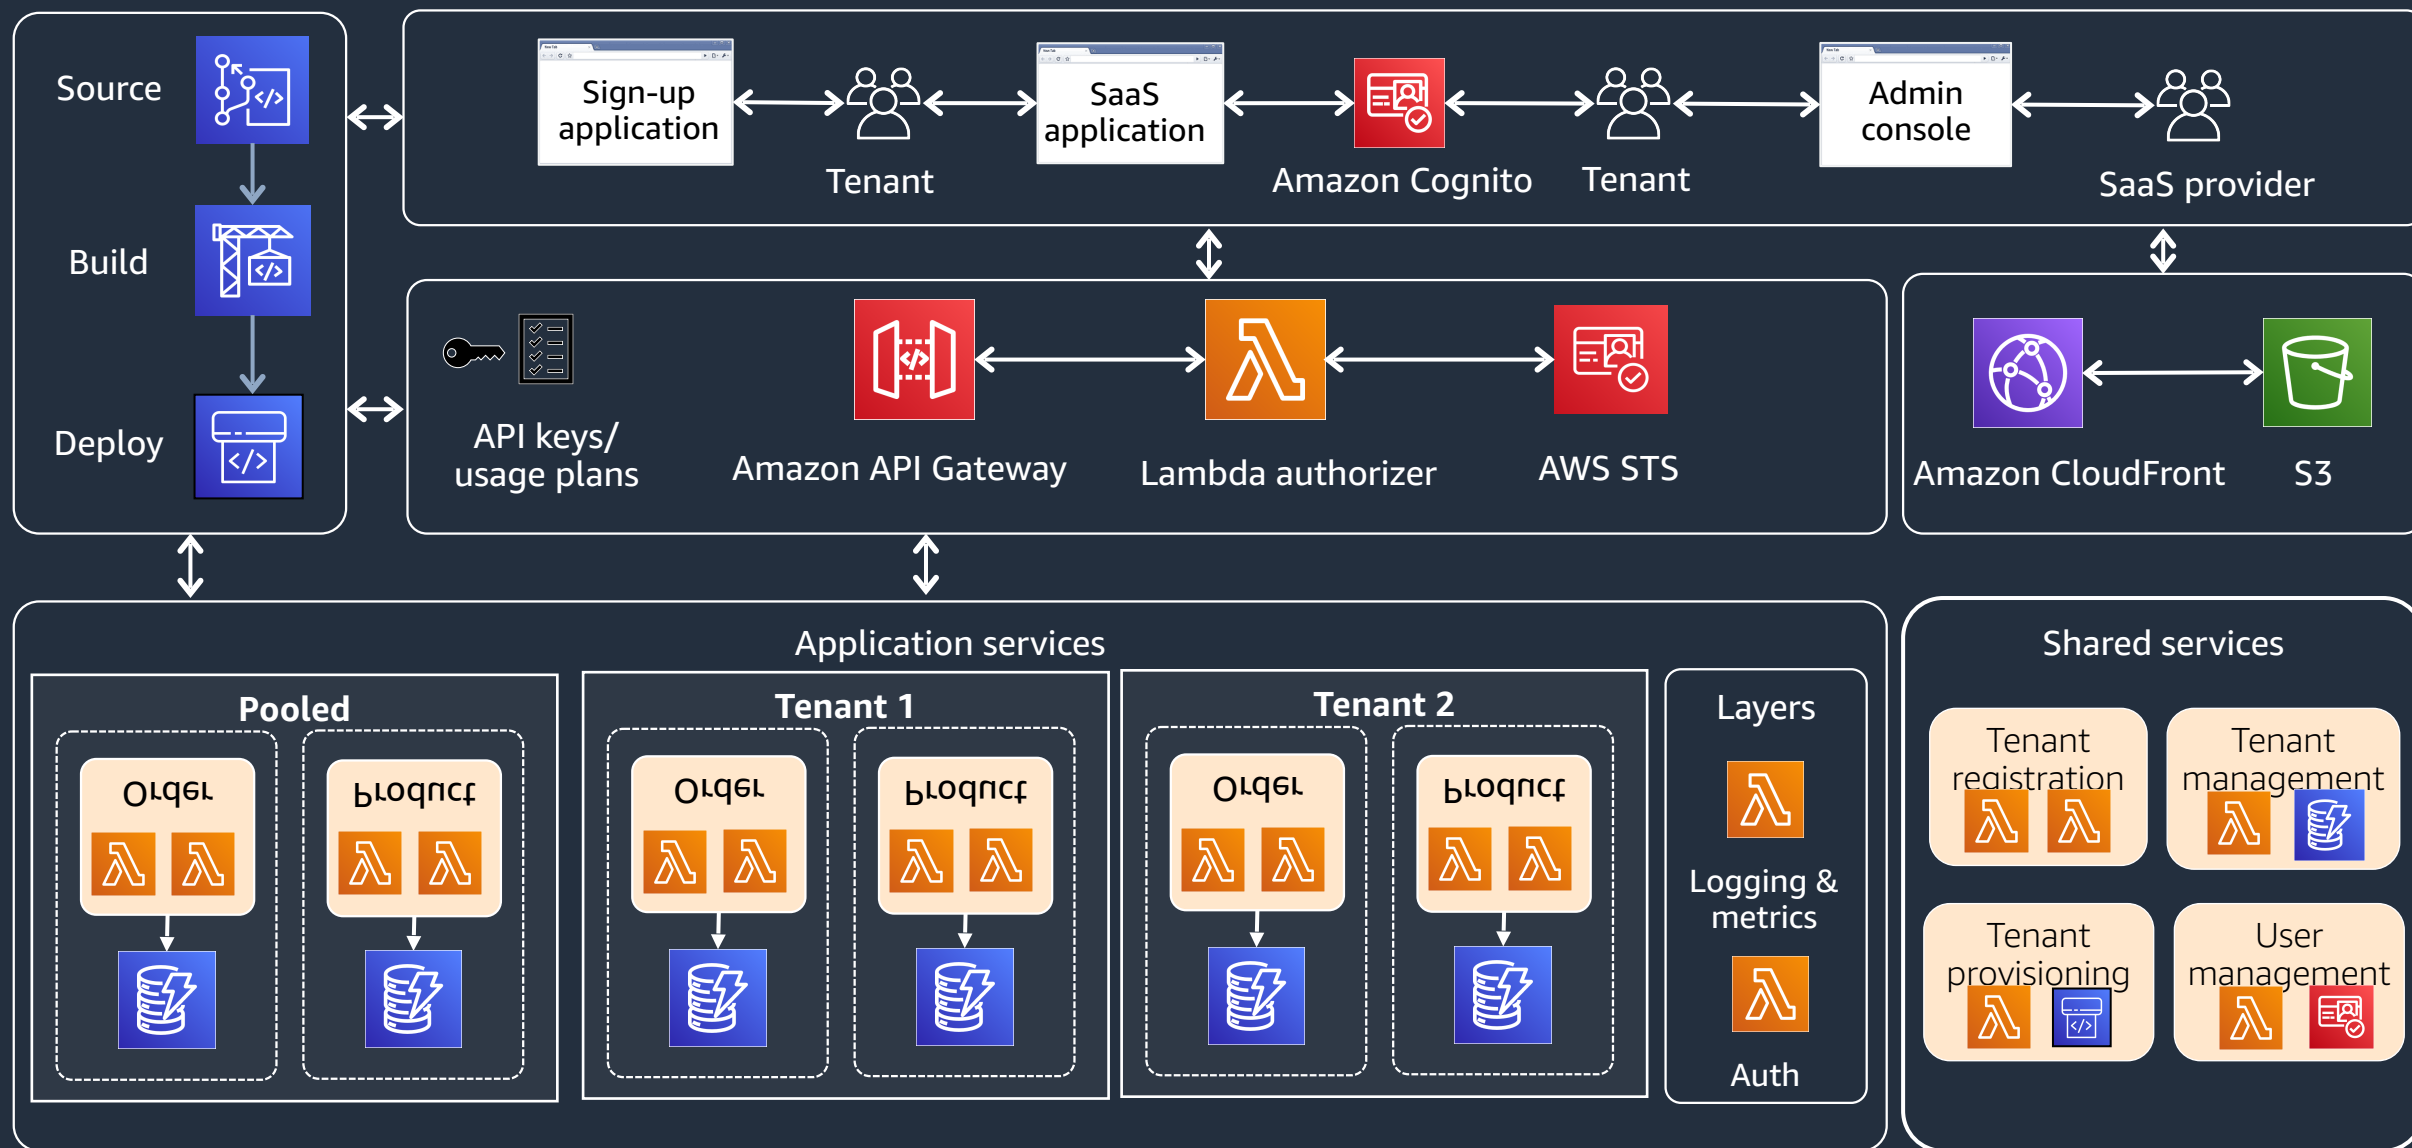
CDK



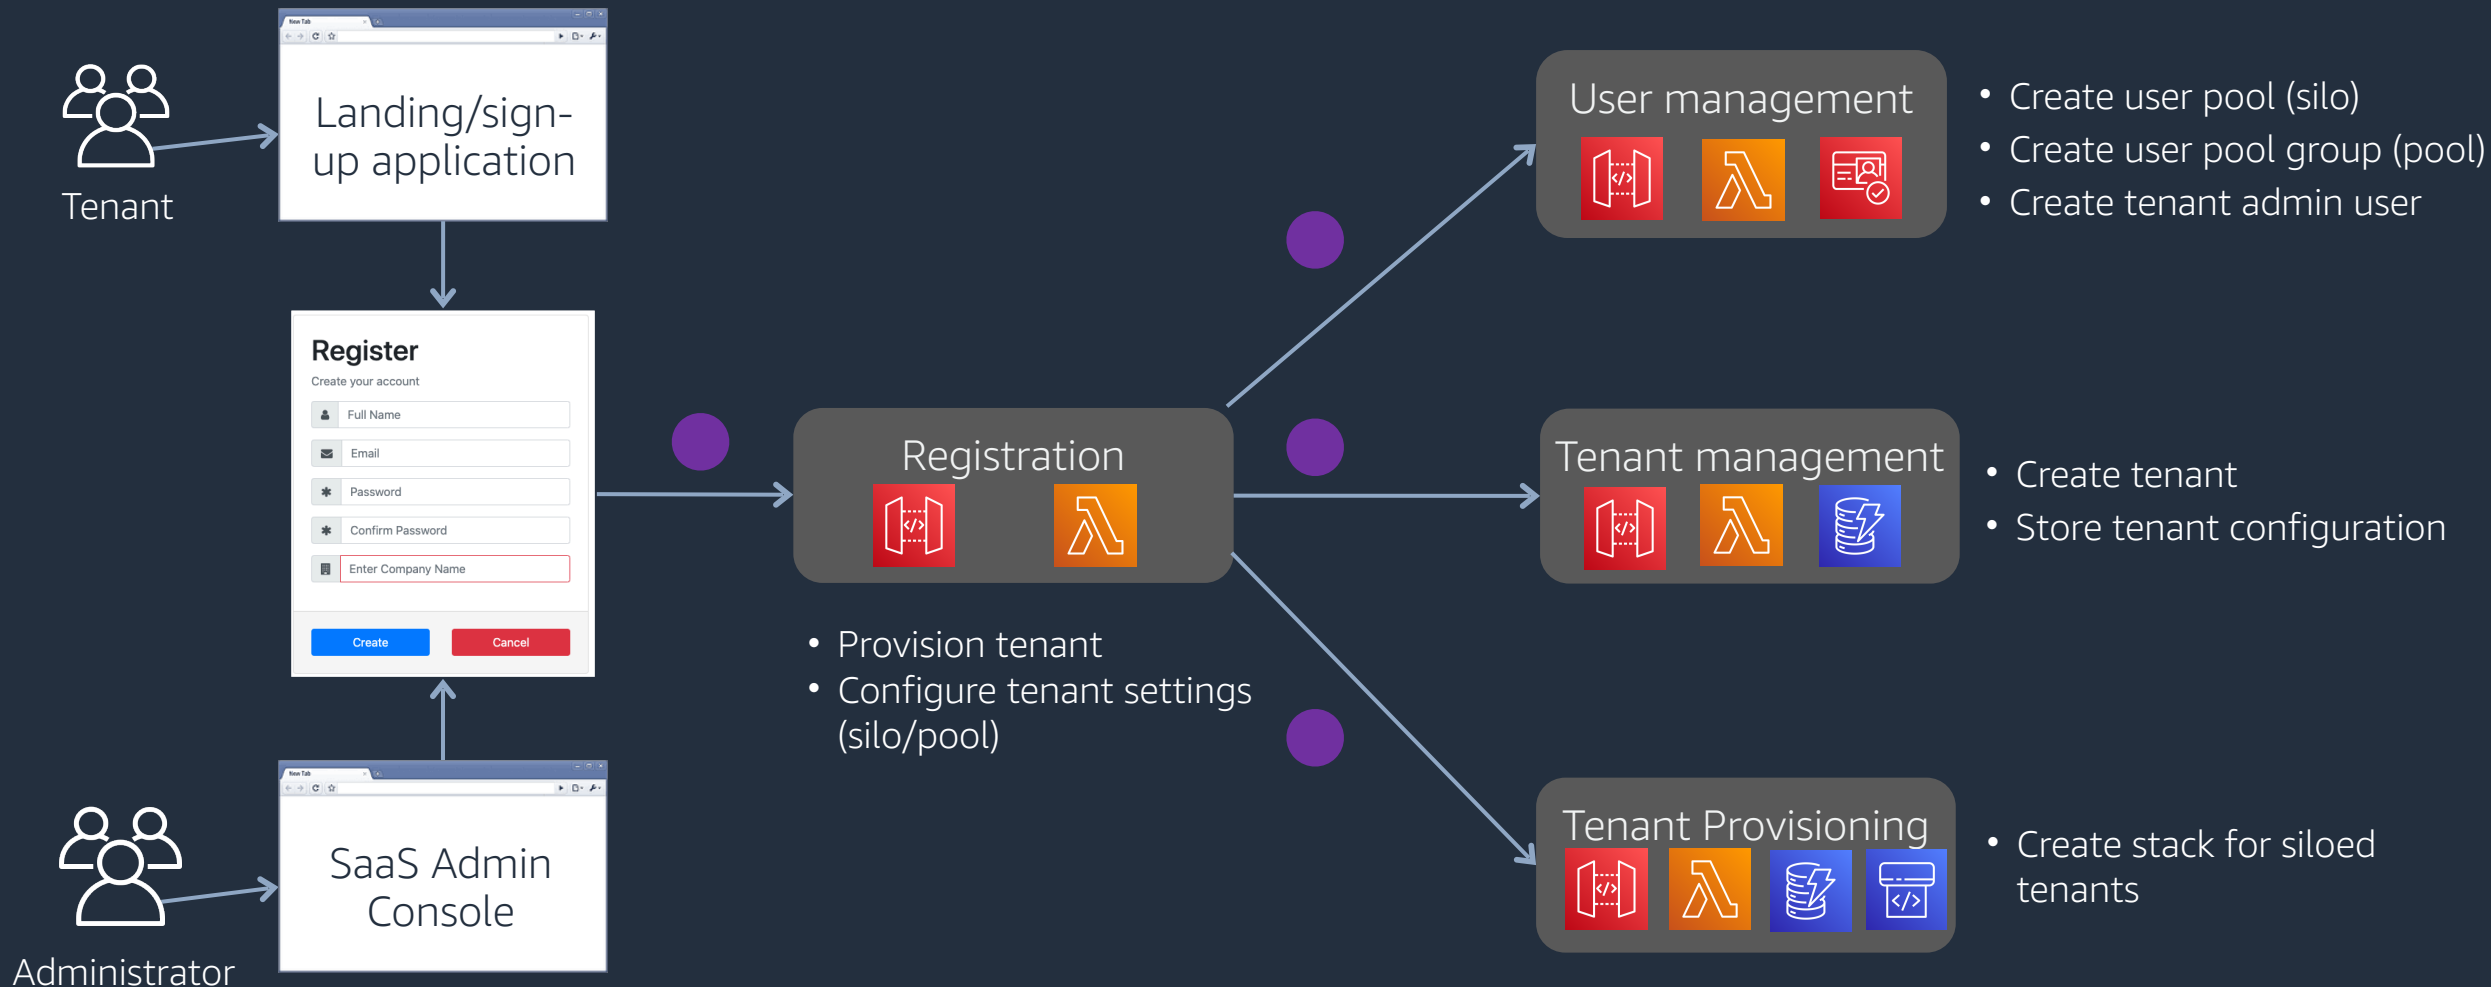
Code Pipeline

Canary deployments

Overall architecture



Registering new tenants

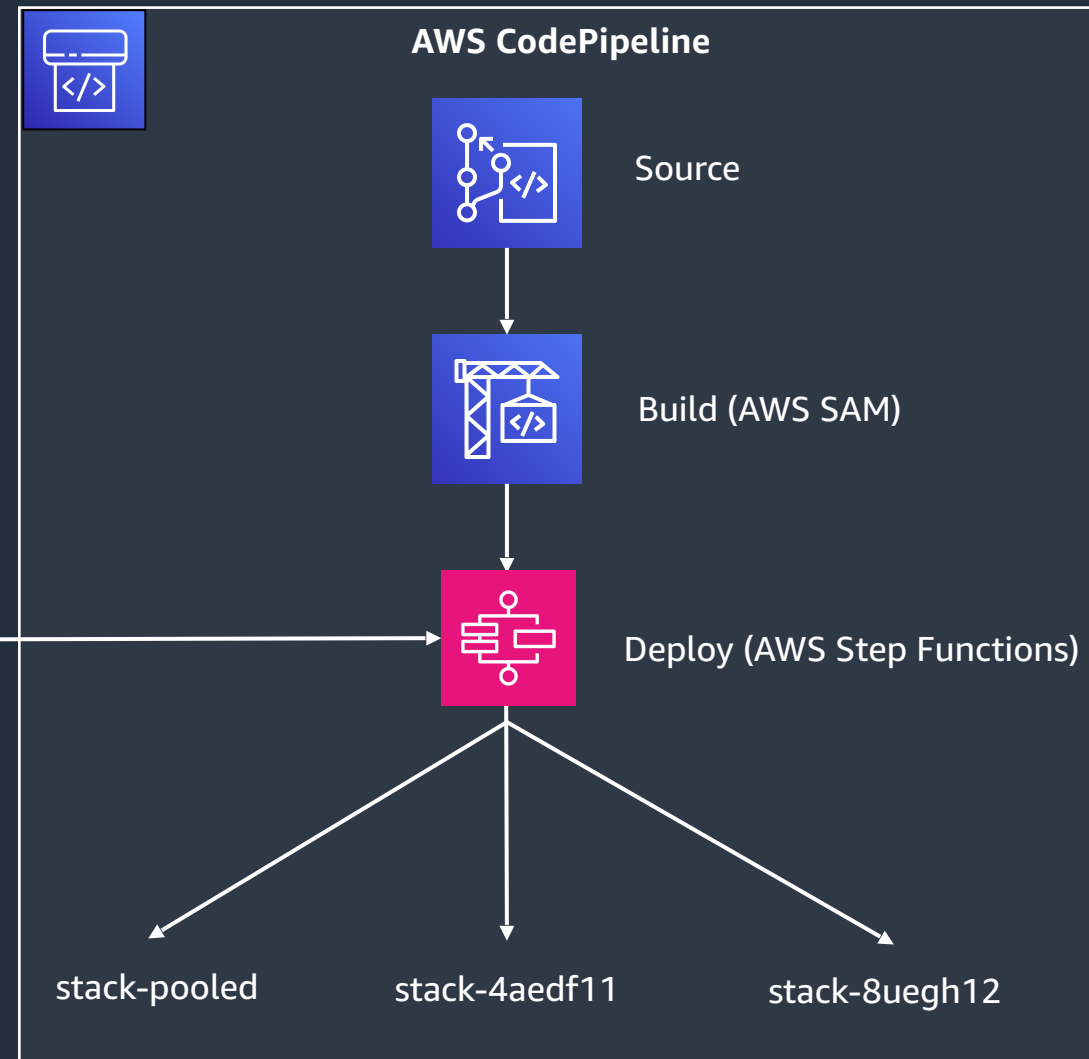


Tenant provisioning and deployment



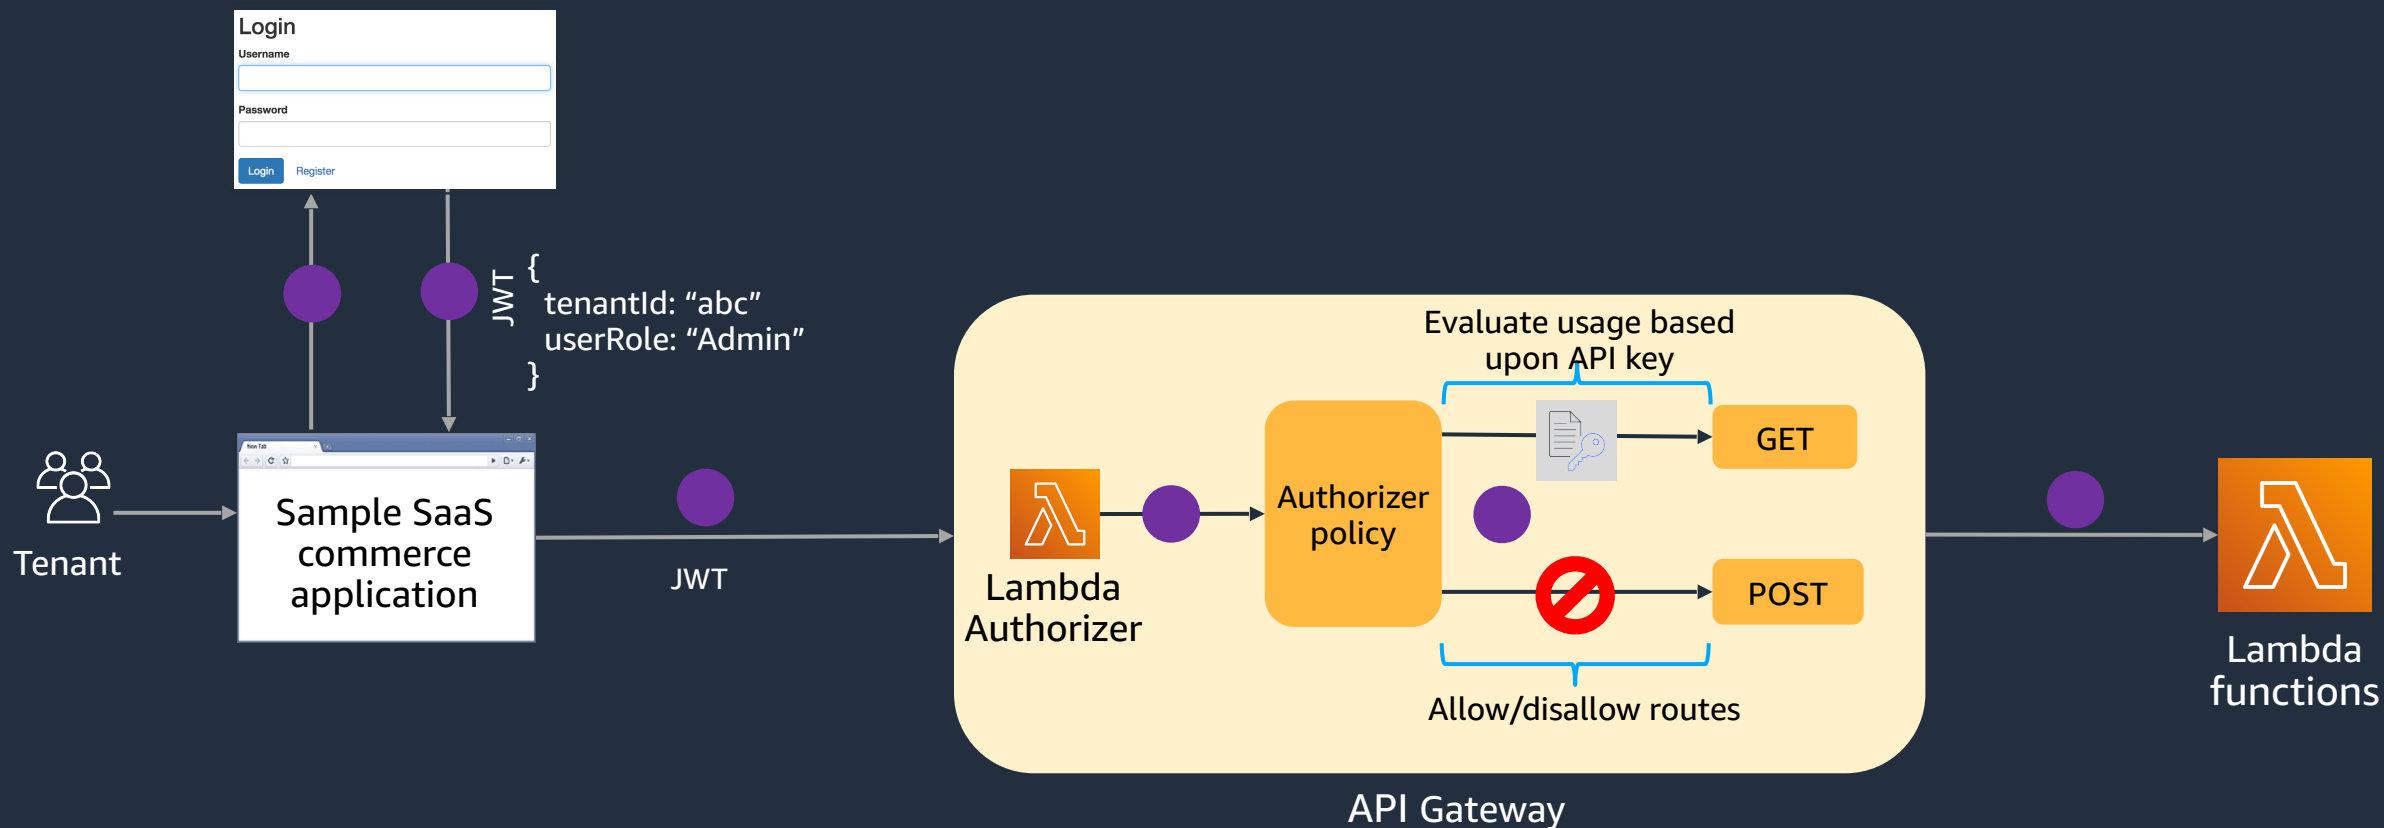
TenantStackMapping

tenantId	waveNumber	codeCommitId	stackName
pooled	1	0a65d	stack-pooled
tenant1	2	0a65d	stack-tenant1
tenant2	2	0a65d	stack-tenant2



Authentication and authorization

Cognito Amplify Library



Lambda Authorizer policy and context

```
#only tenant admin and system admin can do certain actions like create and disable users
if (auth_manager.isTenantAdmin(user_role) or auth_manager.isSystemAdmin(user_role)):
    policy.allowAllMethods()
    if (auth_manager.isTenantAdmin(user_role)):
        policy.denyMethod(HttpVerb.POST, "tenant-activation")
        policy.denyMethod(HttpVerb.GET, "tenants")
else:
    #if not tenant admin or system admin then only allow to get info and update info
    policy.allowMethod(HttpVerb.GET, "user/*")
    policy.allowMethod(HttpVerb.PUT, "user/*")

authResponse = policy.build()
```

Policy to allow/deny routes based on User Role

```
context = {
    #redacted
    'userName': user_name,
    'tenantId': tenant_id,
    'userPoolId': userpool_id,
    'apiKey': api_key,
    'userRole': user_role
}

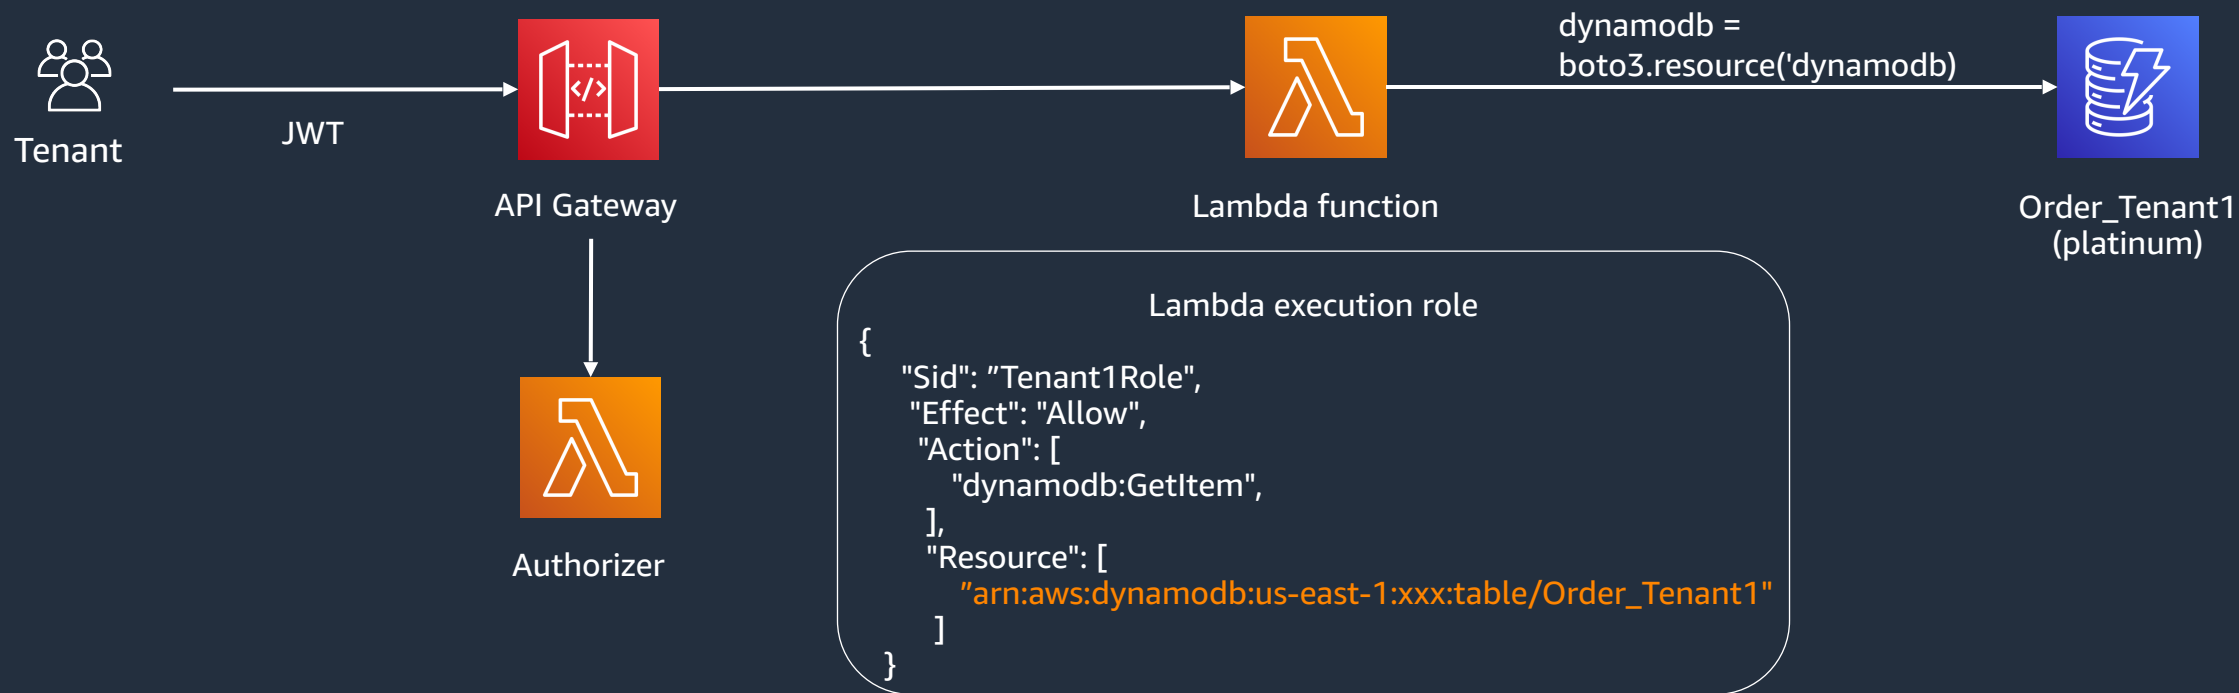
authResponse['context'] = context
authResponse['usageIdentifierKey'] = api_key

return authResponse
```

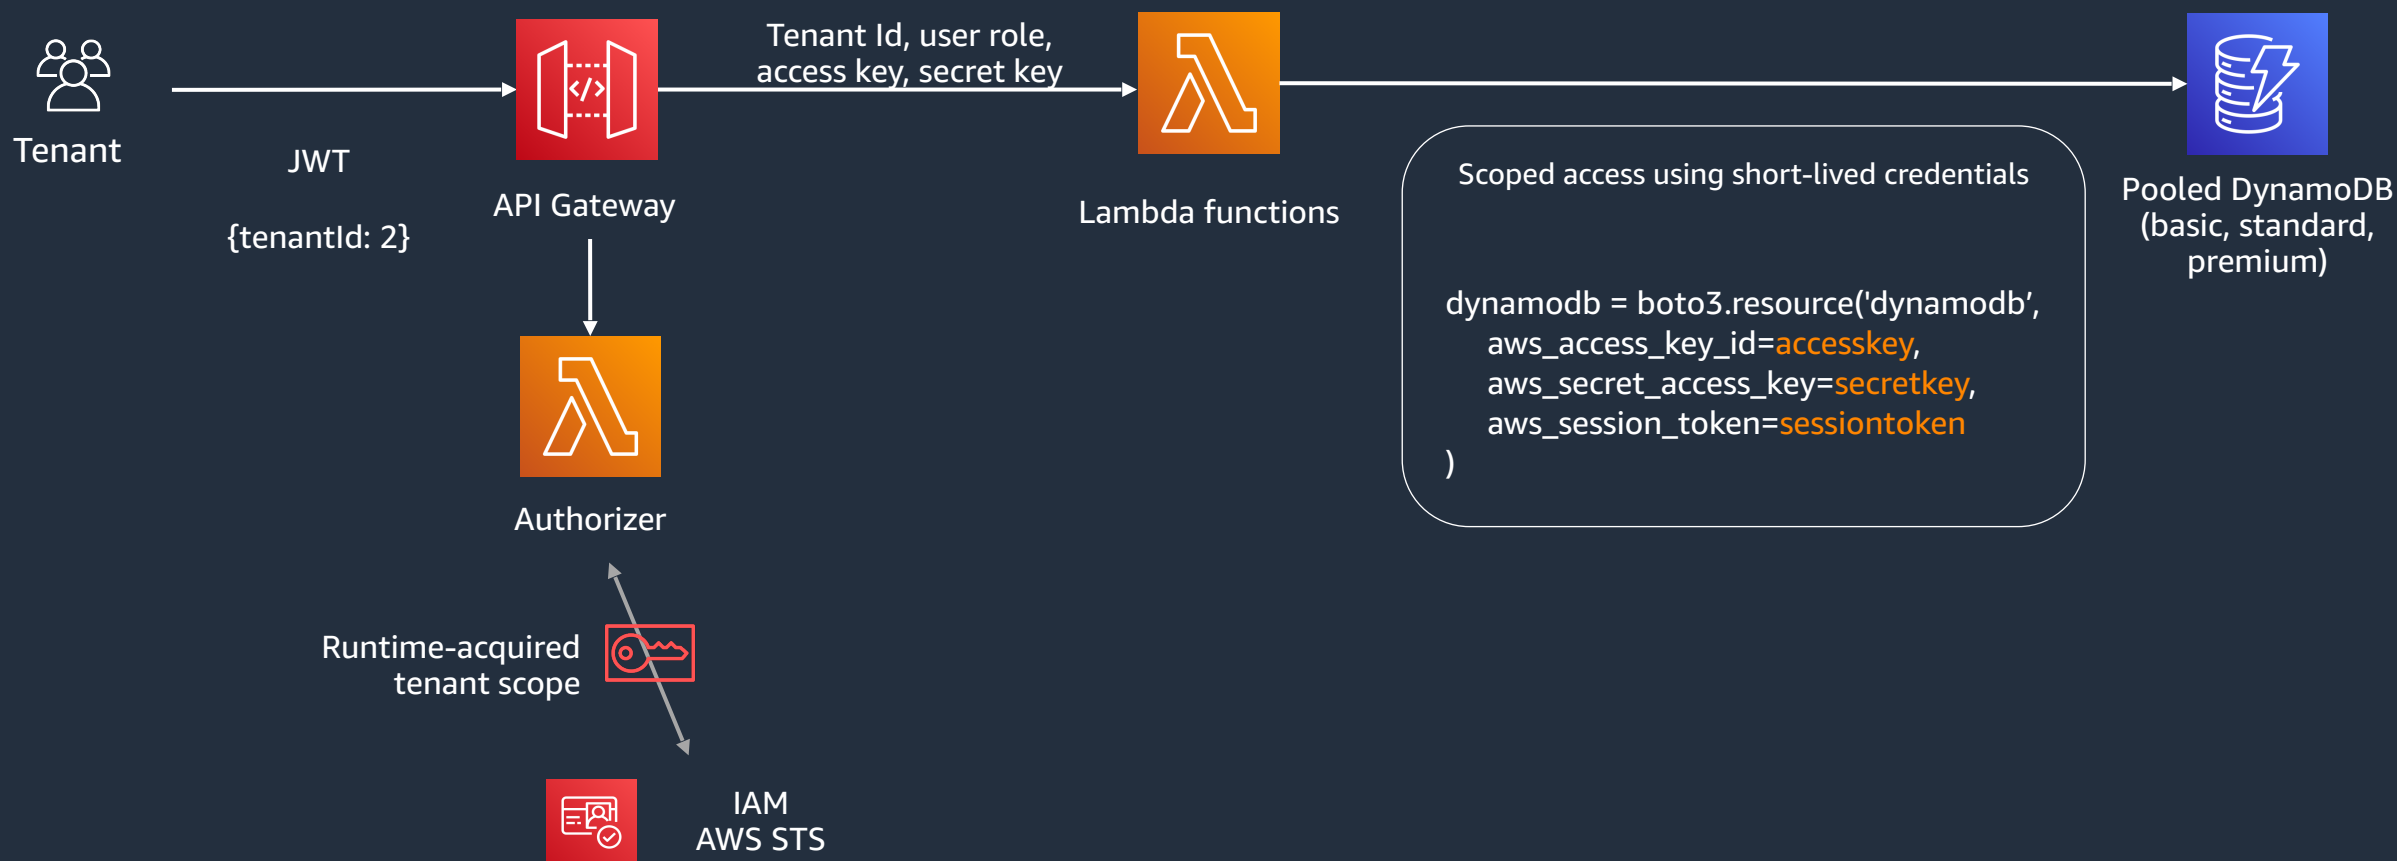
Context with tenant_id and api_key

x-amazon-apigateway-api-key-source : "AUTHORIZER"

Tenant isolation: Silo model



Tenant isolation: Pooled model



Dynamic policy

```
{
  "Effect": "Allow",
  "Action": [
    "dynamodb:UpdateItem",
    "dynamodb:GetItem",
    "dynamodb:PutItem",
    "dynamodb:DeleteItem",
    "dynamodb:Query"
  ],
  "Resource": [
    "arn:aws:dynamodb:{0}:{1}:table/Product-*".format(region, aws_account_id),
  ],
}
```

Code snippet: STS credentials

```
#get IAM policy
```

```
iam_policy = getPolicyForUser(user_role, tenant_id)
```

```
#use STS client to generate the credentials
```

```
assumed_role = sts_client.assume_role(  
    RoleArn=role_arn,  
    RoleSessionName="tenant-aware-session",  
    Policy=iam_policy,  
)
```

```
credentials = assumed_role["Credentials"]
```

```
#pass sts credentials to lambda
```

```
context = {  
    'accesskey': credentials['AccessKeyId'], # $context.authorizer.key -> value  
    'secretkey' : credentials['SecretAccessKey'],  
    'sessiontoken' : credentials["SessionToken"],  
}
```


Tier applied in data access layer

```
142 def __get_dynamodb_table(event, dynamodb):
143     if (is_pooled_deploy=='true'):
144         accesskey = event['requestContext']['authorizer']['accesskey']
145         secretkey = event['requestContext']['authorizer']['secretkey']
146         sessiontoken = event['requestContext']['authorizer']['sessiontoken']
147         dynamodb = boto3.resource('dynamodb',
148                                   aws_access_key_id=accesskey,
149                                   aws_secret_access_key=secretkey,
150                                   aws_session_token=sessiontoken
151                                   )
152     else:
153         if not dynamodb:
154             dynamodb = boto3.resource('dynamodb')
155
156     return dynamodb.Table(table_name)
157
```

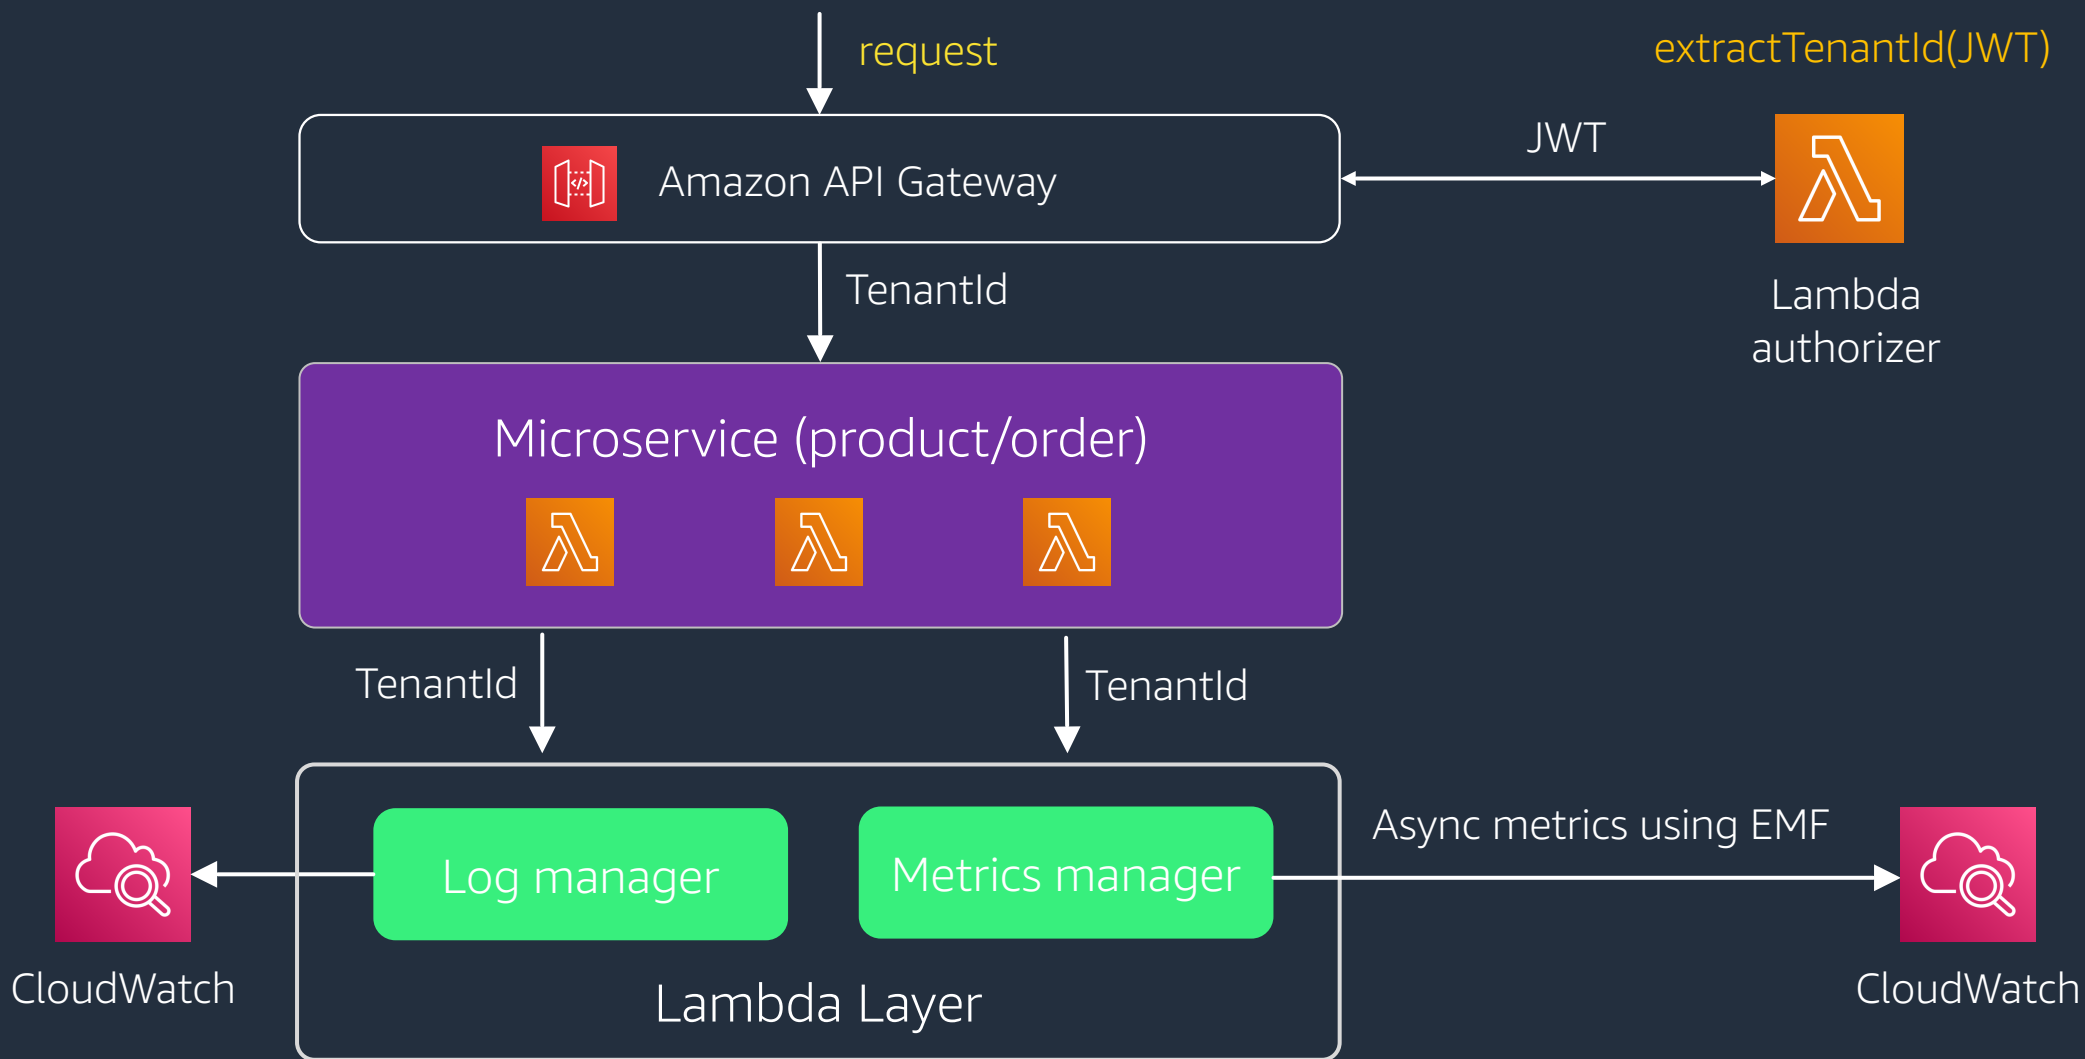
DynamoDB partition key structure

<input type="checkbox"/>	ShardID	ProductID	doc
<input type="checkbox"/>	tenant1-1	1	{"Description": "PRODUCT DESCRIPTION FOR PRODUCT ID : 1", "SKU": 750, "ProductPrice": 2709}
<input type="checkbox"/>	tenant1-4	2	{"Description": "PRODUCT DESCRIPTION FOR PRODUCT ID : 2", "SKU": 750, "ProductPrice": 1700}
<input type="checkbox"/>	tenant1-7	3	{"Description": "PRODUCT DESCRIPTION FOR PRODUCT ID : 3", "SKU": 41, "ProductPrice": 1885}
<input type="checkbox"/>	tenant1-9	8	{"Description": "PRODUCT DESCRIPTION FOR PRODUCT ID : 9", "SKU": 46, "ProductPrice": 1540}
<input type="checkbox"/>	tenant2-1	4	{"Description": "PRODUCT DESCRIPTION FOR PRODUCT ID : 4", "SKU": 766, "ProductPrice": 1081}
<input type="checkbox"/>	tenant2-4	6	{"Description": "PRODUCT DESCRIPTION FOR PRODUCT ID : 6", "SKU": 692, "ProductPrice": 677}
<input type="checkbox"/>	tenant2-5	5	{"Description": "PRODUCT DESCRIPTION FOR PRODUCT ID : 5", "SKU": 692, "ProductPrice": 677}
<input type="checkbox"/>	tenant2-6	7	{"Description": "PRODUCT DESCRIPTION FOR PRODUCT ID : 7", "SKU": 577, "ProductPrice": 3211}

<https://aws.amazon.com/blogs/apn/partitioning-pooled-multi-tenant-saas-data-with-amazon-dynamodb/>



Abstract Tenant Details with Lambda Layers



Observability

Logging with tenant context

```
▶ 2021-09-21T18:26:55.488-04:00 EXTENSION Name: cloudwatch_lambda_agent
▼ 2021-09-21T18:26:55.518-04:00 {"level":"INFO","location":"log_with_ten
{
  "level": "INFO",
  "location": "log_with_tenant_context:19",
  "message": "Request received to create a product",
  "timestamp": "2021-09-21 22:26:55,499+0000",
  "service": "ProductService",
  "tenant_id": "4fb565351b0a11ecab15250a6772c4a2",
  "xray_trace_id": "1-614a5c2e-76c4853c4e86383723b27a9d"
}
```

Multi-tenant tracing with annotations

▼ Query refiners

Refine query by **TenantId** ▼

Select rows to filter traces

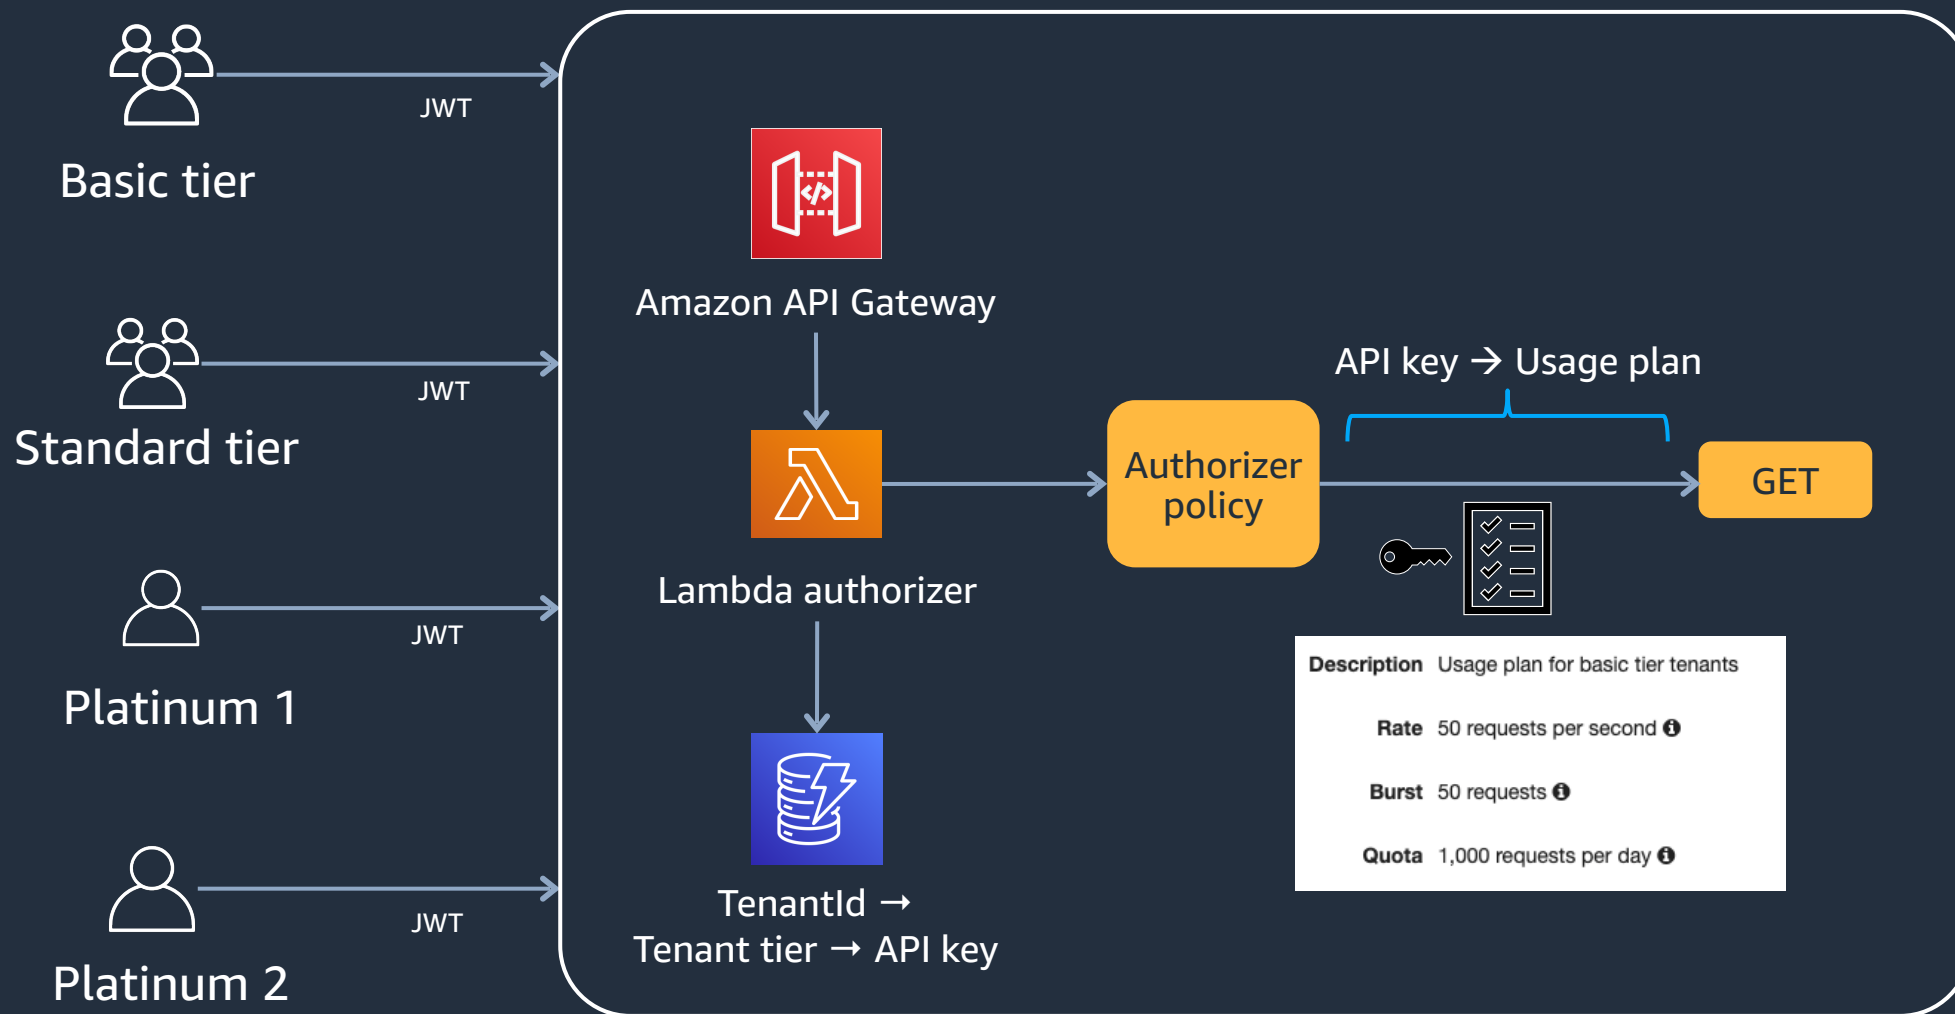
501c4578aa6e11edb87da24fb6091a4e X 1 match

<input checked="" type="checkbox"/>	TenantId	Successful requests
<input checked="" type="checkbox"/>	501c4578aa6e11edb87da24fb6091a4e	100.0%

Multi-tenant metrics (EMF)

```
▼ 2021-08-24T17:39:07.102-04:00 {"_aws": {"Timestamp": 1629841147102, "CloudWatchMetrics": [{"Name": "ProductCreated", "Unit": "Count"}]}, {"tenant_id": "fc651ed9052211eca250e538bb53ea77", "service": "ProductService", "ProductCreated": [{"tenant_id": "fc651ed9052211eca250e538bb53ea77", "service": "ProductService"}]}]}
```

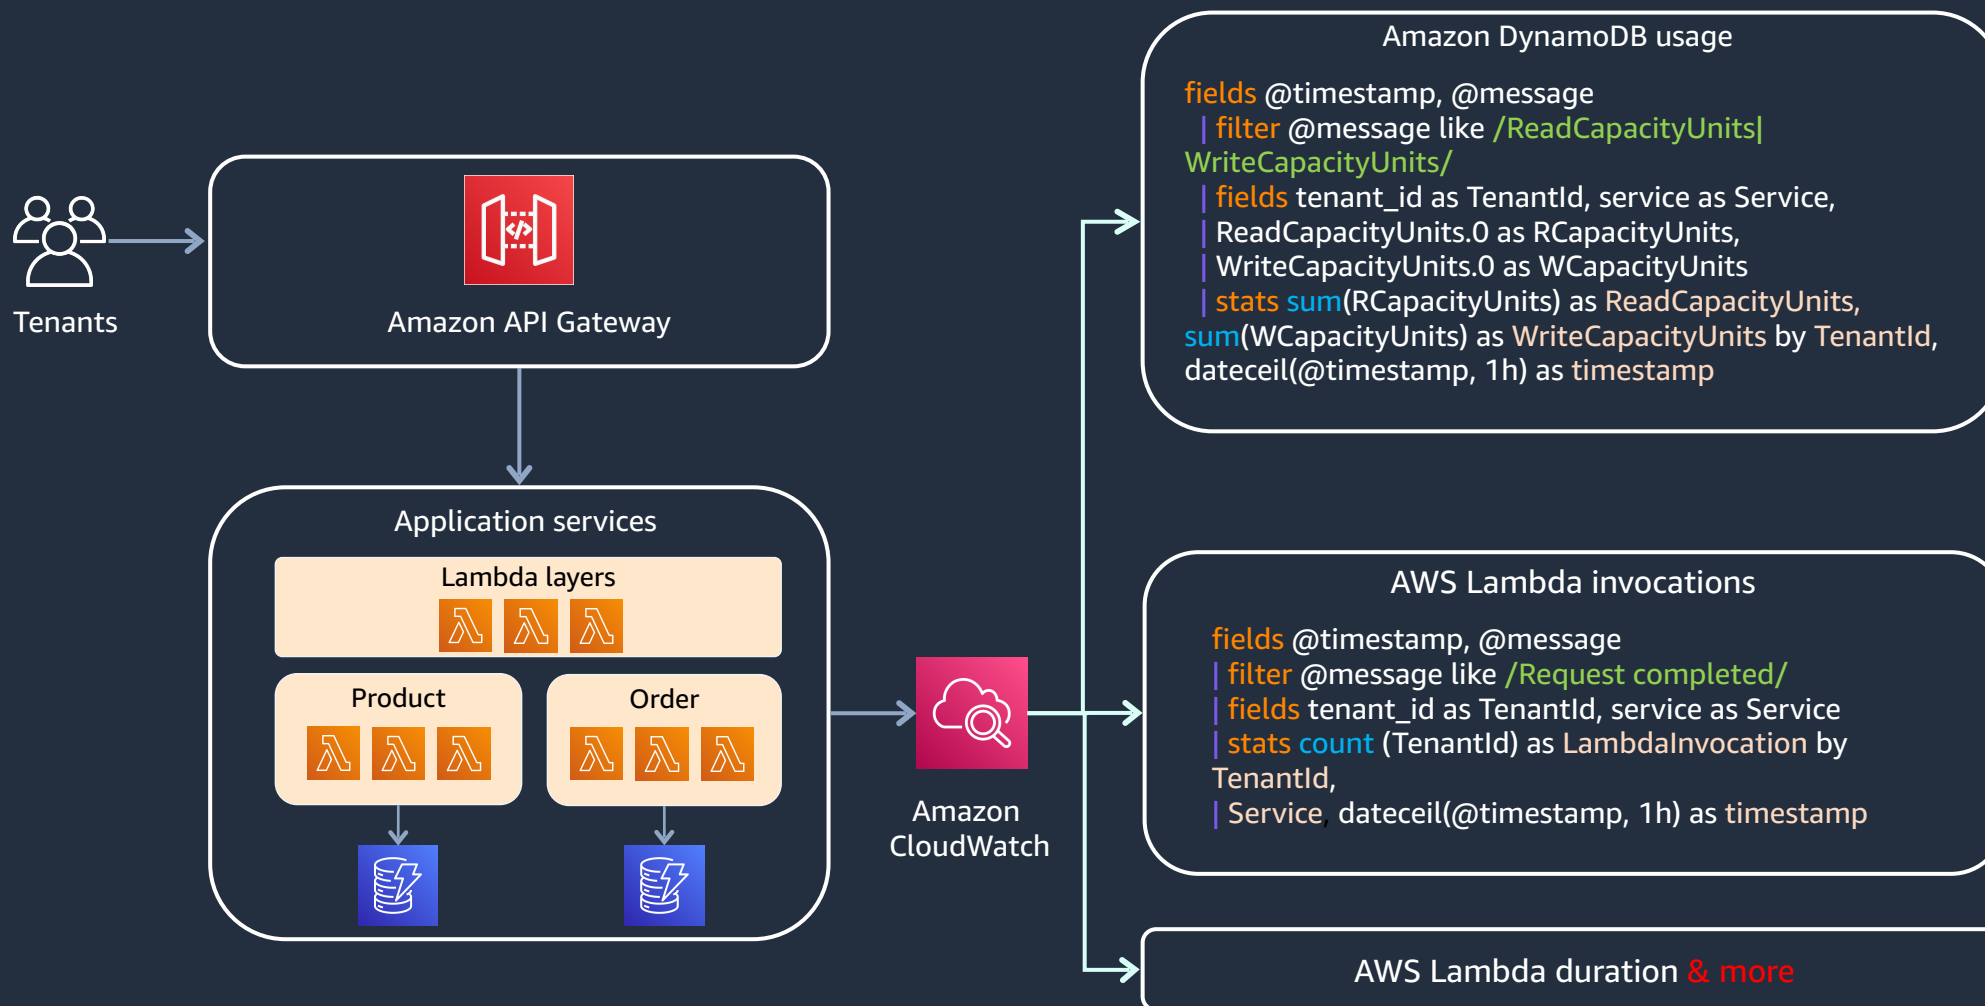
Tier-based throttling policies



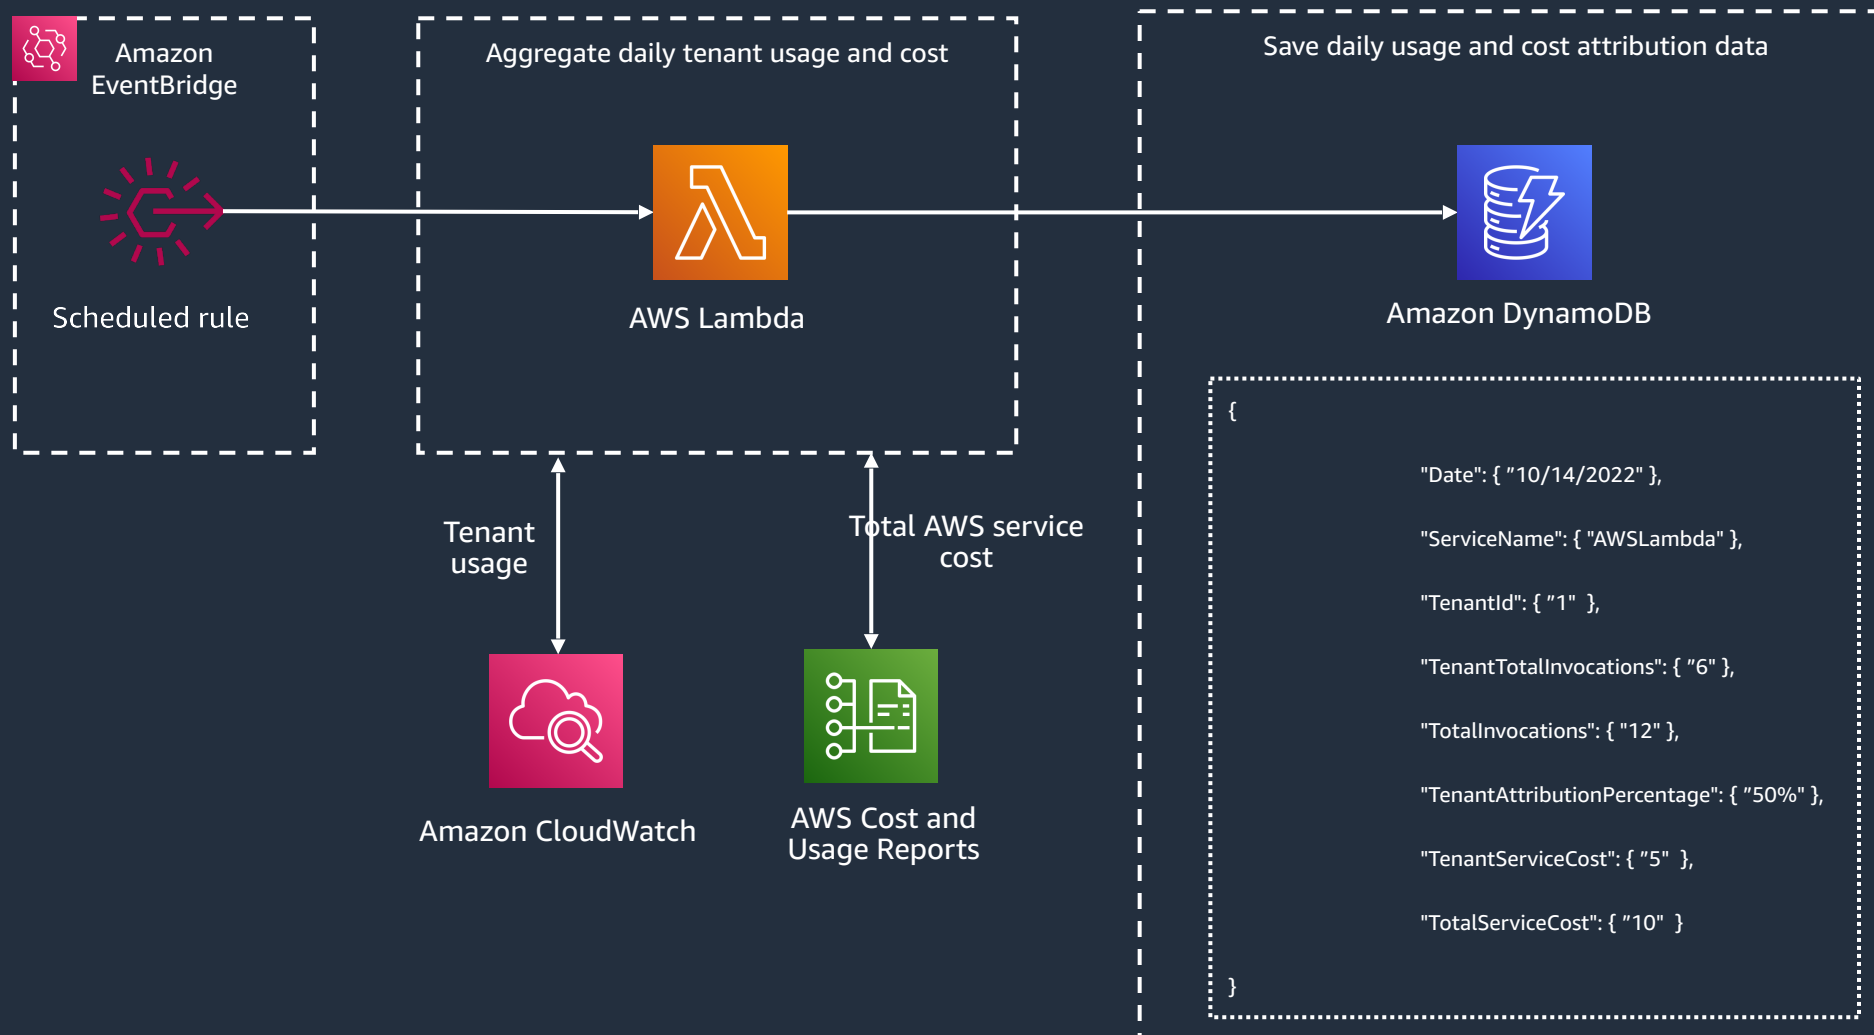
Cost attribution in a pooled model

- Or do you just need consumption?
- Capturing tenant level metrics is the key
- It will be always an approximation
- Capture metrics that are relevant

Cost attribution in a pooled model



Cost attribution in a pooled model



Serverless SaaS reference links

Building SaaS with Serverless blog



Serverless SaaS Workshop



Serverless SaaS Workshop (GitHub Repo)



Serverless SaaS Reference Solution



- <https://aws.amazon.com/blogs/apn/building-a-multi-tenant-saas-solution-using-aws-serverless-services/>
- <https://catalog.us-east-1.prod.workshops.aws/workshops/b0c6ad36-0a4b-45d8-856b-8a64f0ac76bb/en-US>
- <https://github.com/aws-samples/aws-serverless-saas-workshop>
- <https://github.com/aws-samples/aws-saas-factory-ref-solution-serverless-saas>

Meet us at re:Invent



- SAS404 | Designing and implementing a SaaS cost-per-tenant strategy
- SAS406 | SaaS DevOps deep dive: Automating multi-tenant deployments
- SAS403 | SaaS survivor: Building a rich multi-tenant operations experience
- SVS320 | Building multi-tenant applications with AWS Lambda and AWS Fargate
- API306 | Combining Step Functions and EventBridge: Use cases and best practices

Continue your AWS serverless learning

Learn at your
own pace



Expand your serverless
skills with our learning plan
on **AWS Skill Builder**

Increase your
knowledge



Use our **Ramp-Up Guides**
to build your serverless
knowledge

Earn AWS
Serverless badge



Demonstrate your
knowledge by achieving
digital badges




<https://s12d.com/serverless-learning>



Thank you



Anubhav Sharma
Principal Solutions Architect
AWS SaaS Factory
 antonal80



Anubhav Sharma
Principal Solutions Architect
AWS SaaS Factory
 anubhavynr