



Serverless Shuffle Sharding

Improving resiliency and fault isolation
in a multi-tenant SaaS environment



Anton Aleksandrov

Principal Solution Architect, Serverless
AWS

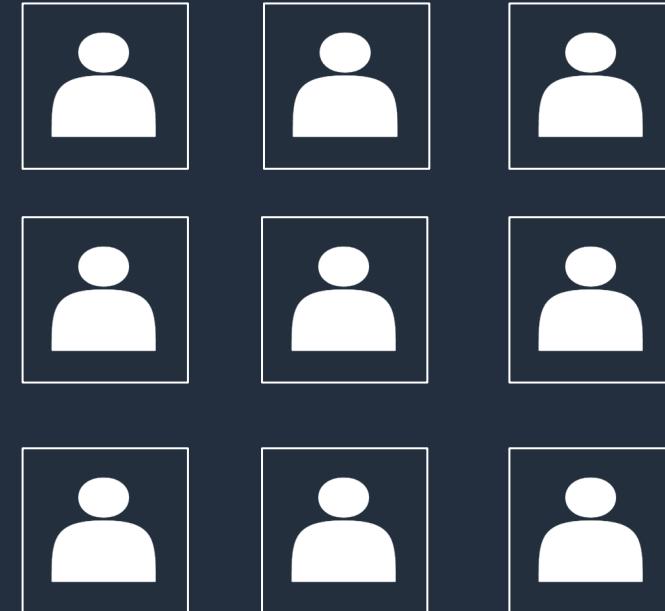
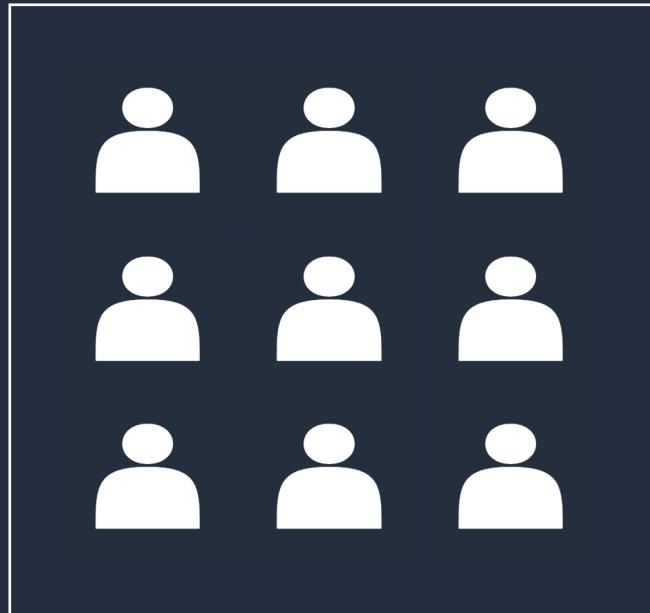
Reliability

The reliability pillar focuses on workloads performing their intended functions and how to **recover quickly from failure** to meet demands. Key topics include **distributed system design**, **recovery planning**, and **adapting to changing requirements**

AWS Well-Architected Framework

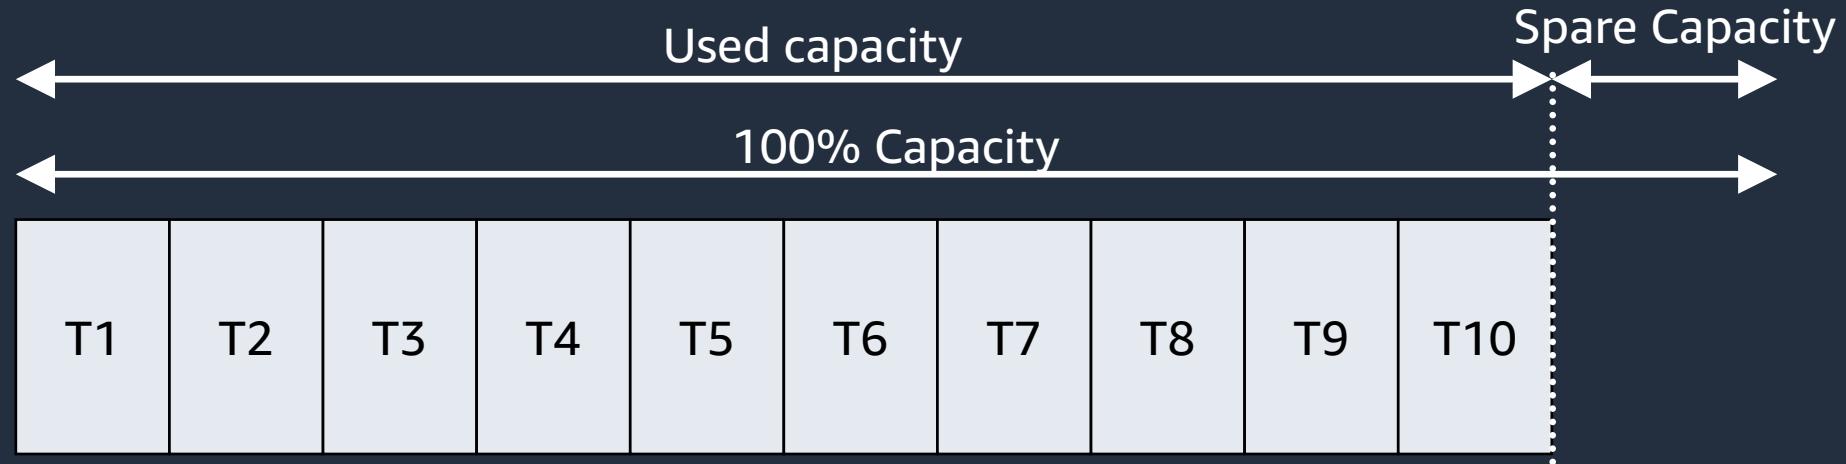
<https://docs.aws.amazon.com/wellarchitected/latest/reliability-pillar/welcome.html>

What is a multi-tenant SaaS environment



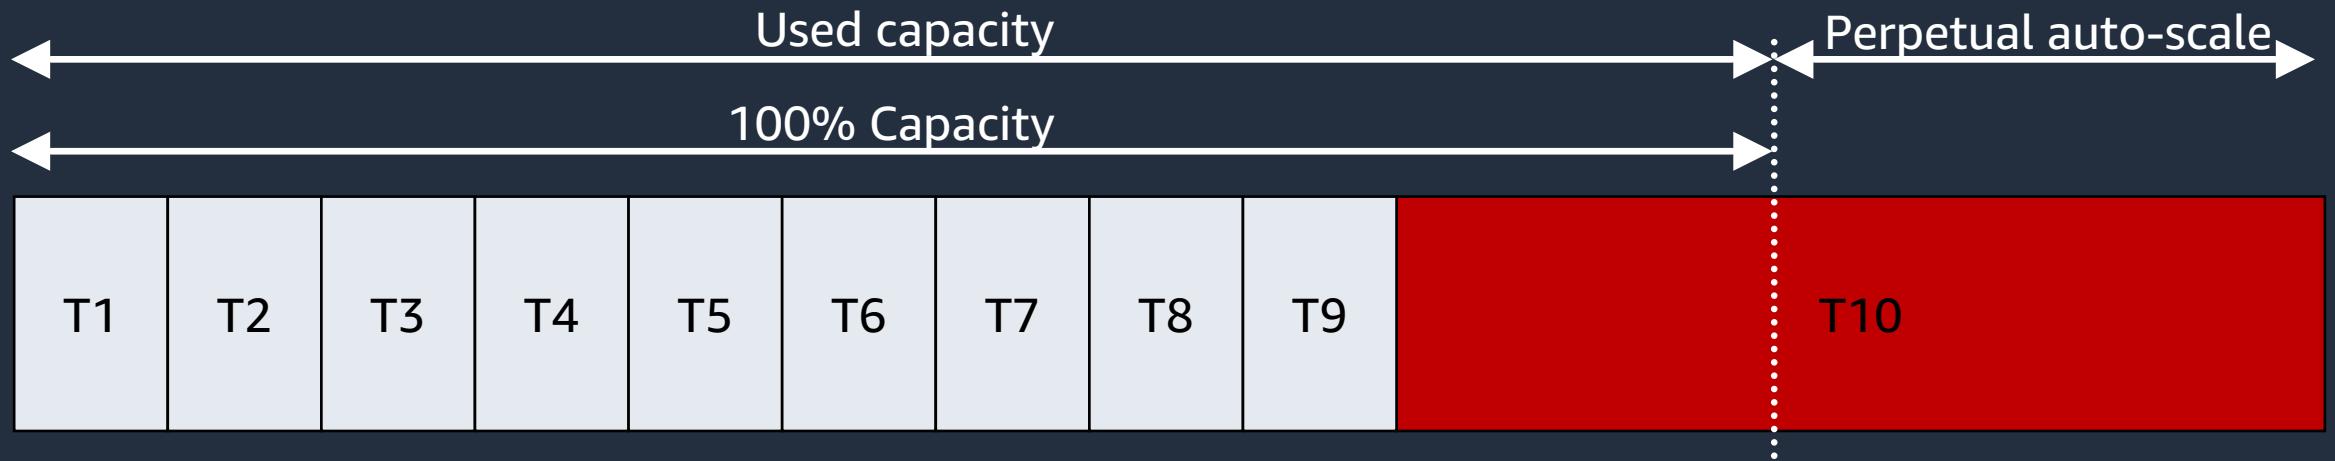
What is a noisy neighbors problem

Noisy neighbors



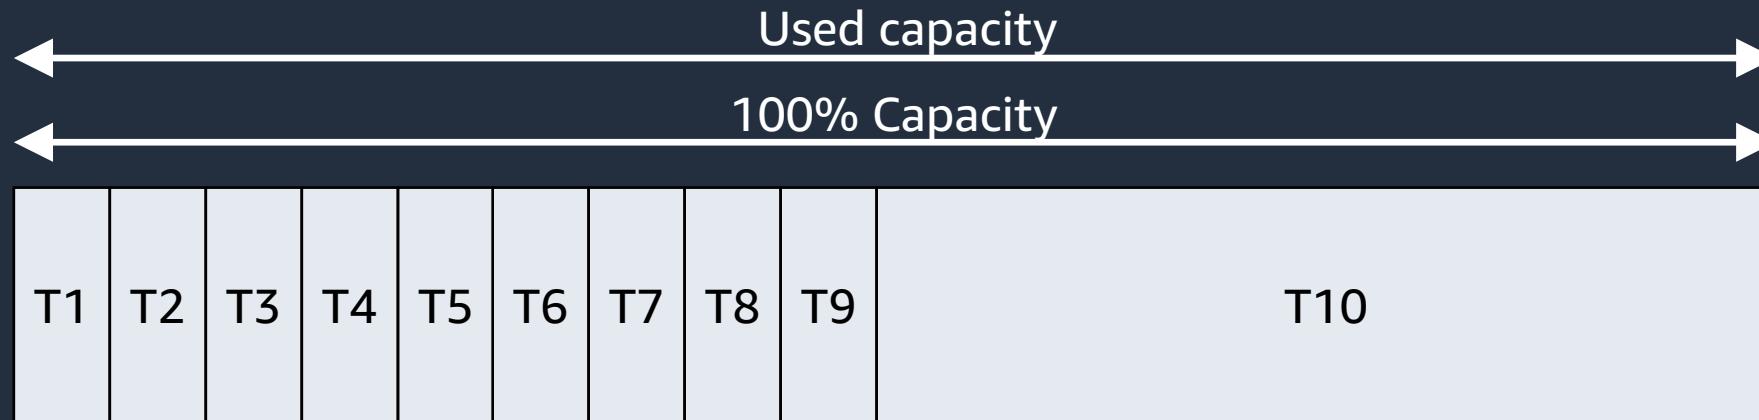
Utopia

Noisy neighbors



Expectation

Noisy neighbors

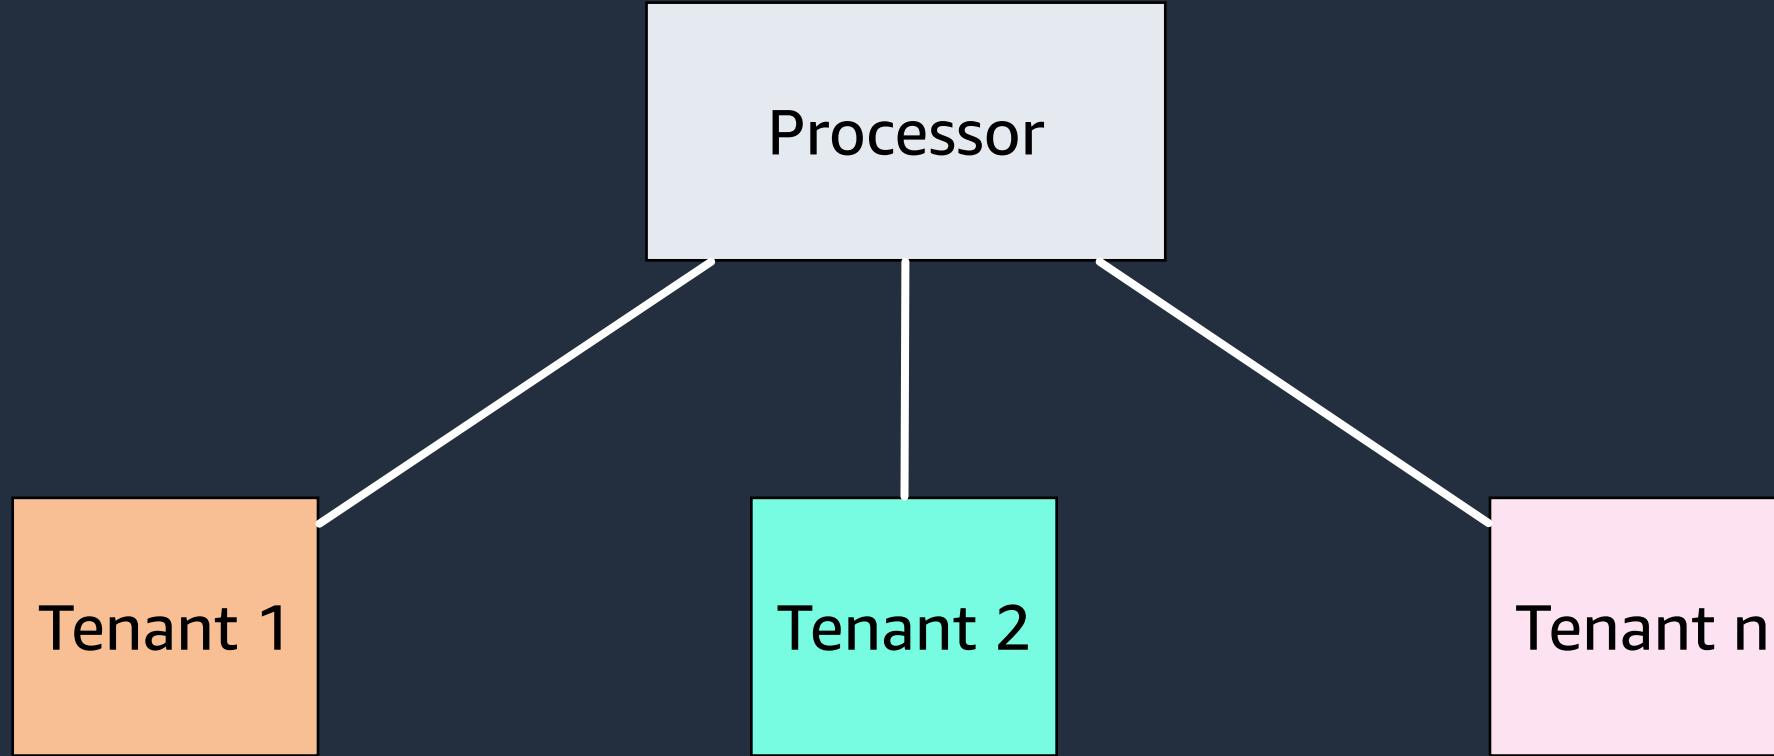


Reality

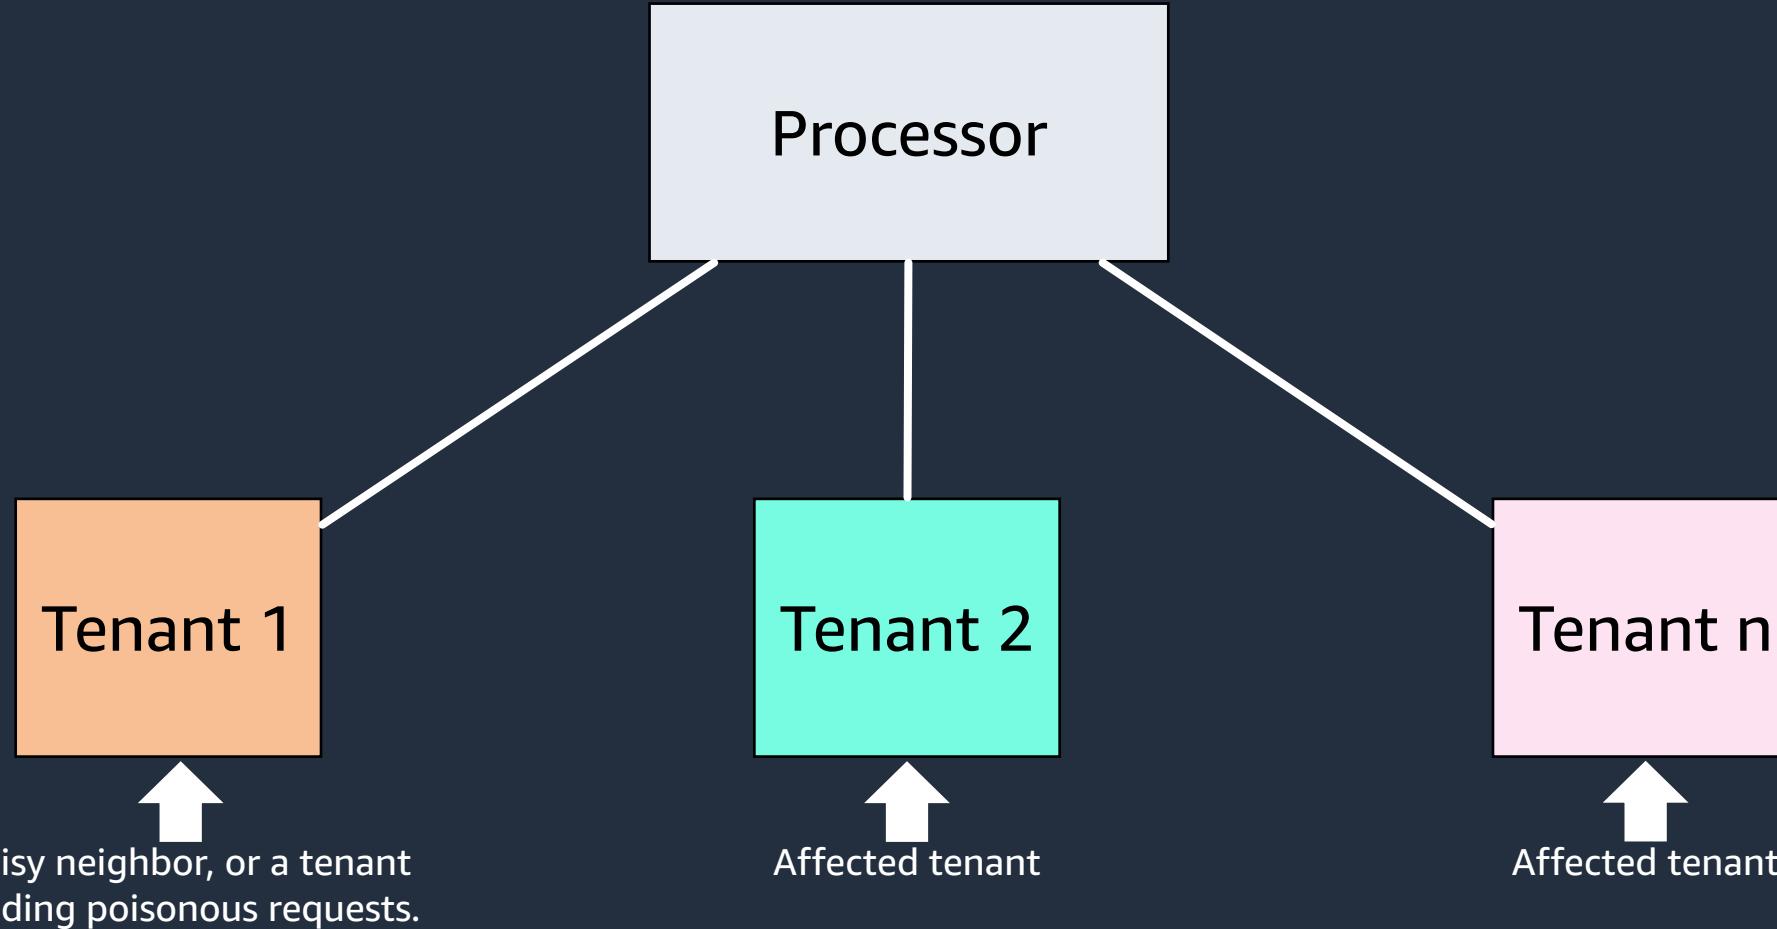
https://en.wikipedia.org/wiki/Cloud_computing_issues#Performance_interference_and_noisy_neighbors

Let's see this in action

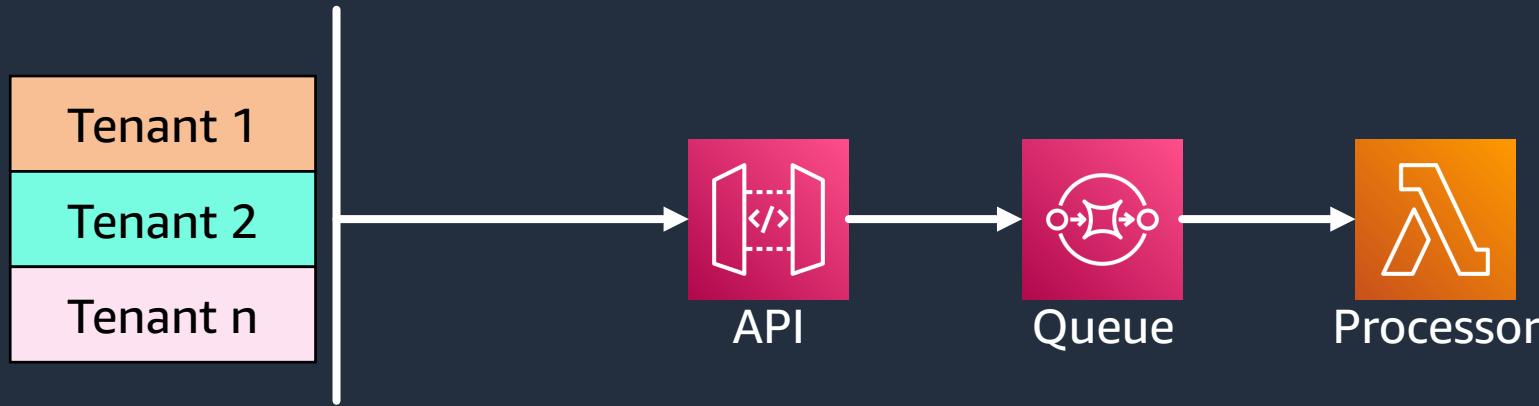
Single processor



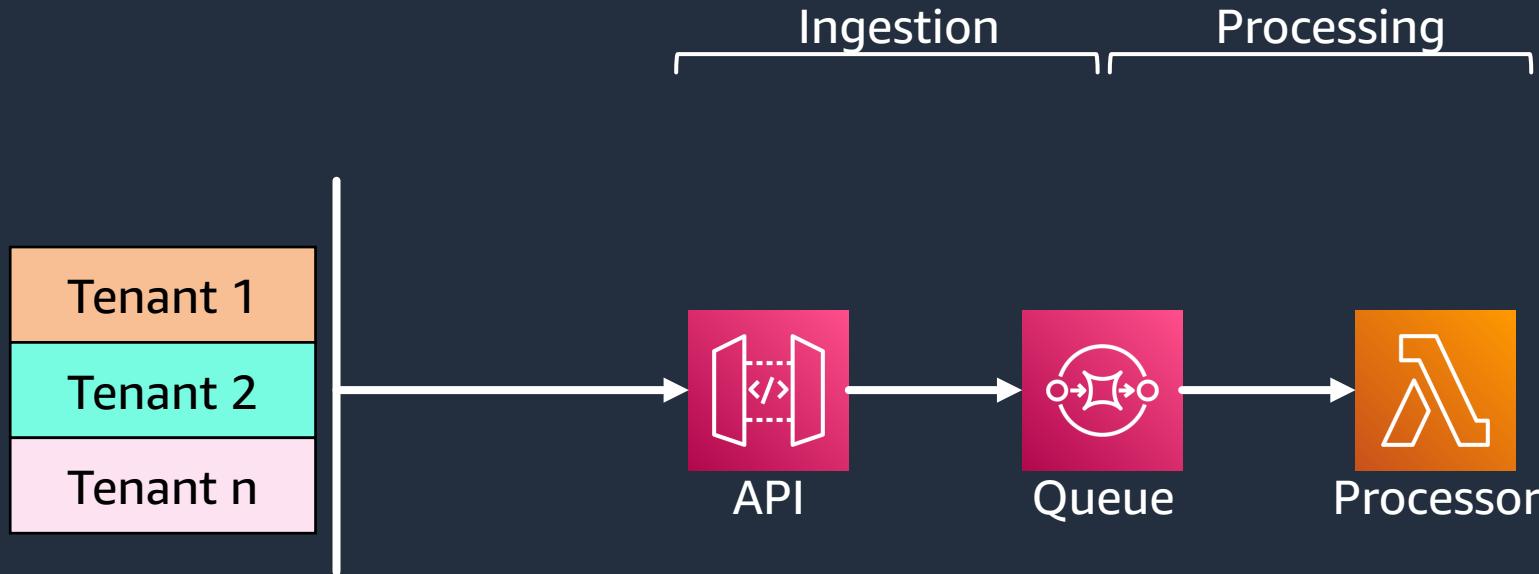
Single processor



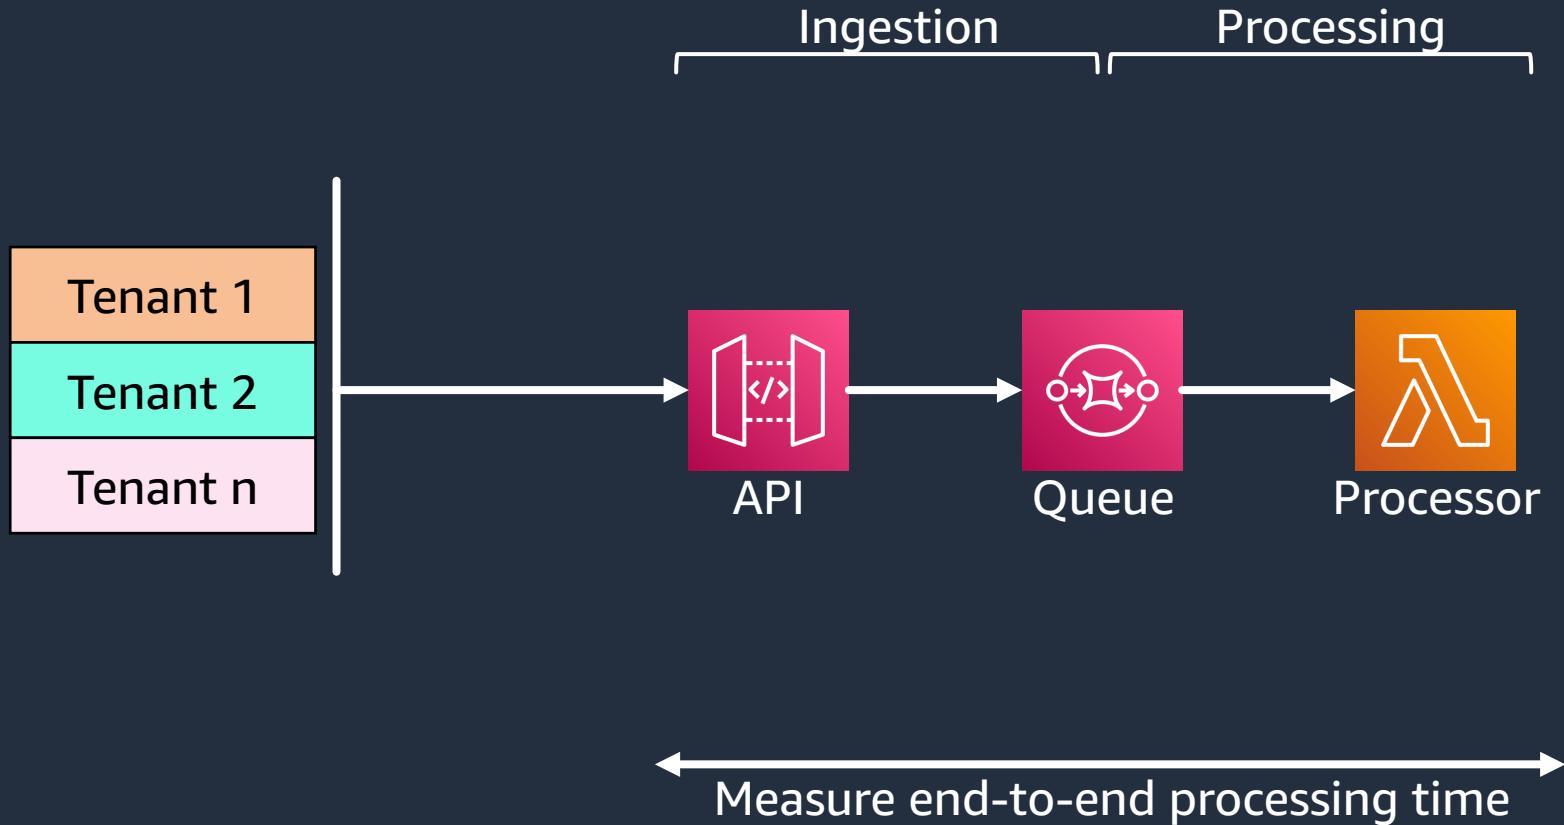
Single processor



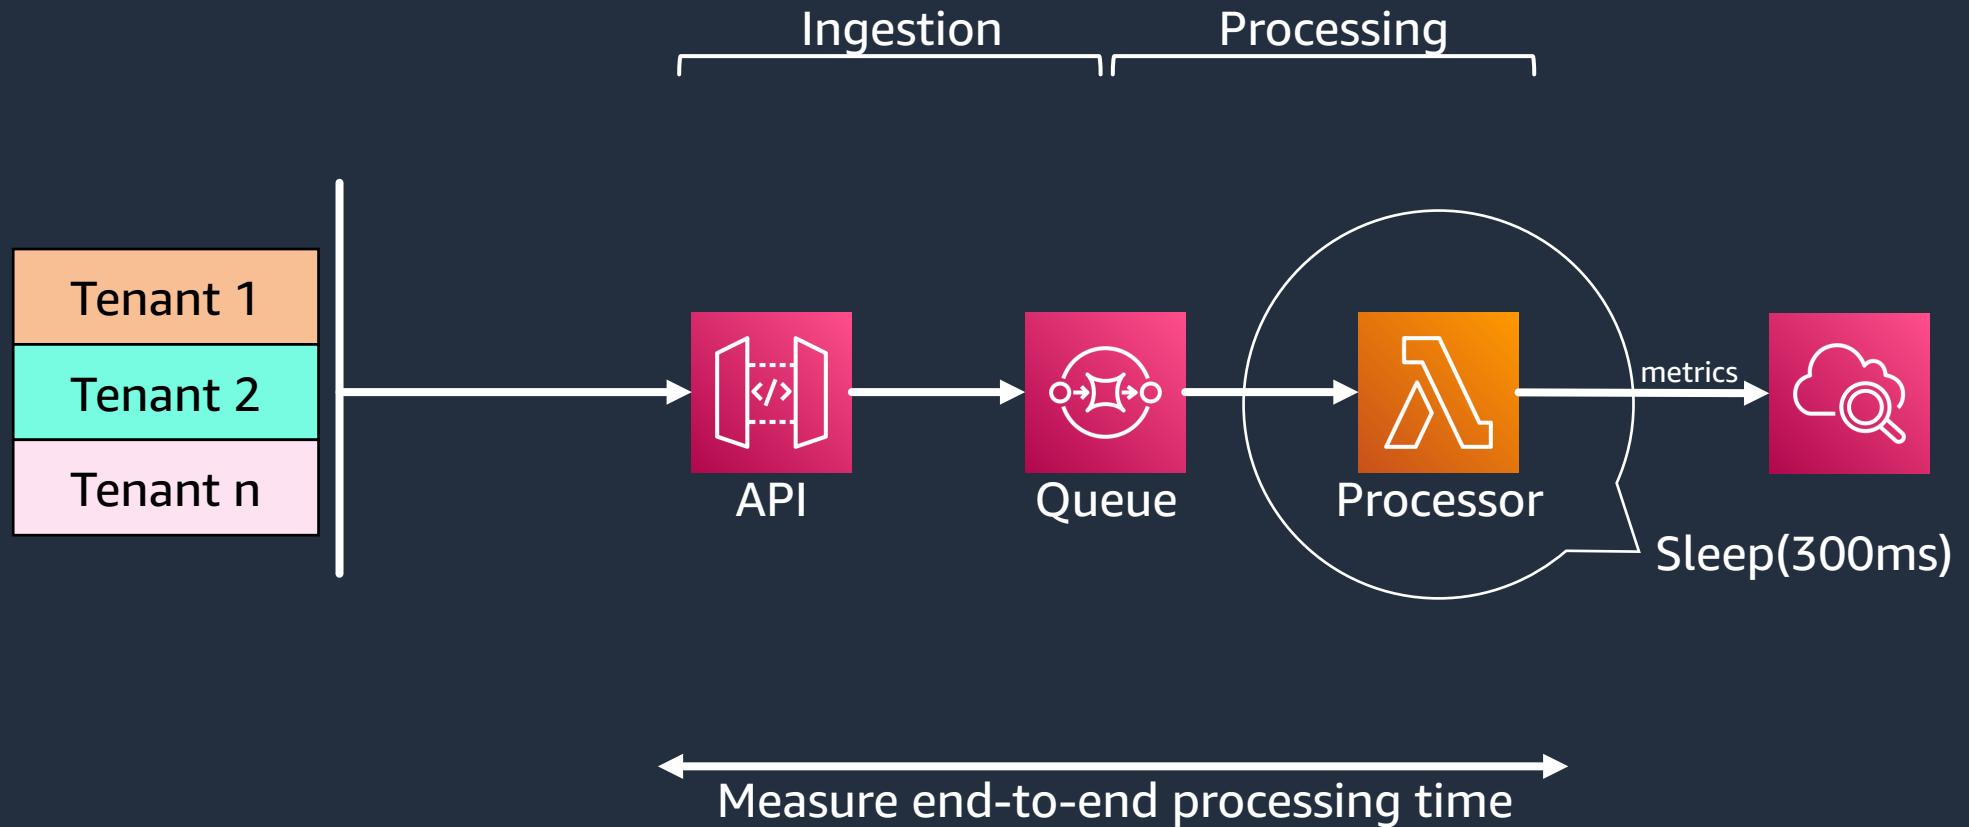
Single processor



Single processor



Single processor



Test setup

Step 1 – Peace and quiet

Thread Group

Name: Tenants 1-19

Comments:

Action to be taken after a Sampler error

Continue Start Next Thread Loop Stop Thread

Thread Properties

Number of Threads (users): **19**

Ramp-up period (seconds): 1

Loop Count: Infinite
 Same user on each iteration
 Delay Thread creation until needed
 Specify Thread lifetime

Duration (seconds):

Startup delay (seconds):

Thread Group

Name: Tenant 20

Comments:

Action to be taken after a Sampler error

Continue Start Next Thread Loop Stop Thread

Thread Properties

Number of Threads (users): **1**

Ramp-up period (seconds): 1

Loop Count: Infinite
 Same user on each iteration
 Delay Thread creation until needed
 Specify Thread lifetime

Duration (seconds):

Startup delay (seconds):

Test setup

Step 2 – Noisy neighbor!

Thread Group

Name: Tenants 1-19

Comments:

Action to be taken after a Sampler error
 Continue Start Next Thread Loop Stop Th

Thread Properties

Number of Threads (users): **19**

Ramp-up period (seconds): 1

Loop Count: Infinite
 Same user on each iteration
 Delay Thread creation until needed
 Specify Thread lifetime

Duration (seconds):

Startup delay (seconds):

Thread Group

Name: Tenant 20

Comments:

Action to be taken after a Sampler error
 Continue Start Next Thread Loop Stop Th

Thread Properties

Number of Threads (users): **15**

Ramp-up period (seconds): 1

Loop Count: Infinite
 Same user on each iteration
 Delay Thread creation until needed
 Specify Thread lifetime

Duration (seconds):

Startup delay (seconds):

Test setup

Step 3 – Cooldown and recovery

Thread Group

Name: Tenants 1-19

Comments:

Action to be taken after a Sampler error

Continue Start Next Thread Loop Stop Thread

Thread Properties

Number of Threads (users): **19**

Ramp-up period (seconds): 1

Loop Count: Infinite
 Same user on each iteration
 Delay Thread creation until needed
 Specify Thread lifetime

Duration (seconds):

Startup delay (seconds):

Thread Group

Name: Tenant 20

Comments:

Action to be taken after a Sampler error

Continue Start Next Thread Loop Stop Thread

Thread Properties

Number of Threads (users): **1**

Ramp-up period (seconds): 1

Loop Count: Infinite
 Same user on each iteration
 Delay Thread creation until needed
 Specify Thread lifetime

Duration (seconds):

Startup delay (seconds):

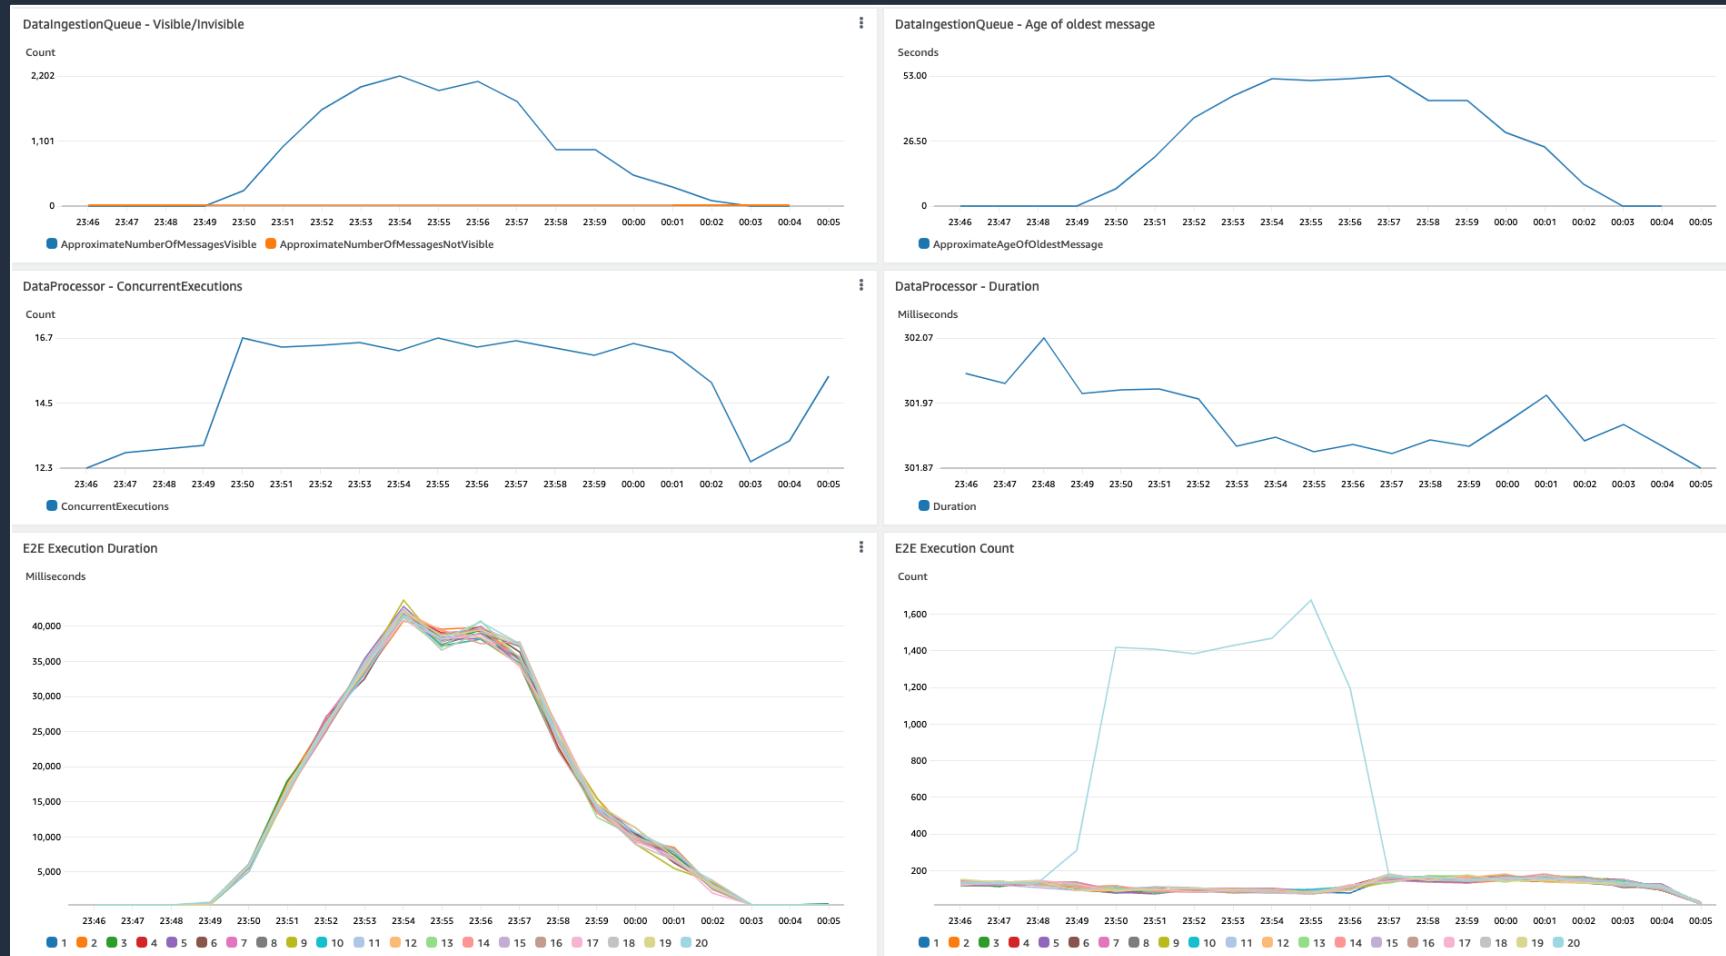
Step 1 - Peace and quiet

| Label | # Samples | Average | Median | 90% Line | 95% Line | 99% Line | Min | Maximum | Error % | Throughput | Re |
|---------------------|-----------|---------|--------|----------|----------|----------|-----|---------|---------|------------|----|
| Request - Tenant 1 | 683 | 325 | 282 | 458 | 554 | 669 | 237 | 1216 | 0.00% | 2.4/sec | |
| Request - Tenant 10 | 644 | 322 | 281 | 461 | 548 | 688 | 223 | 852 | 0.00% | 2.2/sec | |
| Request - Tenant 11 | 666 | 322 | 282 | 452 | 540 | 741 | 238 | 1378 | 0.00% | 2.3/sec | |
| Request - Tenant 12 | 678 | 321 | 283 | 456 | 539 | 606 | 235 | 1129 | 0.00% | 2.3/sec | |
| Request - Tenant 13 | 657 | 325 | 284 | 453 | 558 | 783 | 233 | 990 | 0.00% | 2.3/sec | |
| Request - Tenant 14 | 724 | 324 | 282 | 455 | 553 | 727 | 232 | 1327 | 0.00% | 2.5/sec | |
| Request - Tenant 15 | 711 | 318 | 282 | 426 | 529 | 664 | 230 | 931 | 0.00% | 2.5/sec | |
| Request - Tenant 16 | 659 | 319 | 281 | 454 | 542 | 646 | 232 | 945 | 0.00% | 2.3/sec | |
| Request - Tenant 17 | 693 | 325 | 282 | 456 | 554 | 743 | 229 | 982 | 0.00% | 2.4/sec | |
| Request - Tenant 18 | 687 | 323 | 282 | 454 | 556 | 673 | 237 | 1138 | 0.00% | 2.4/sec | |
| Request - Tenant 19 | 697 | 329 | 285 | 504 | 553 | 619 | 235 | 1283 | 0.00% | 2.4/sec | |
| Request - Tenant 2 | 691 | 325 | 284 | 456 | 544 | 684 | 226 | 968 | 0.00% | 2.4/sec | |
| Request - Tenant 20 | 679 | 324 | 282 | 455 | 547 | 701 | 236 | 1129 | 0.00% | 2.3/sec | |
| Request - Tenant 3 | 658 | 327 | 283 | 465 | 550 | 648 | 234 | 1158 | 0.00% | 2.3/sec | |
| Request - Tenant 4 | 675 | 319 | 283 | 438 | 529 | 612 | 231 | 1261 | 0.00% | 2.3/sec | |
| Request - Tenant 5 | 679 | 323 | 284 | 453 | 542 | 630 | 233 | 1273 | 0.00% | 2.4/sec | |
| Request - Tenant 6 | 699 | 319 | 283 | 426 | 534 | 632 | 234 | 1254 | 0.00% | 2.4/sec | |
| Request - Tenant 7 | 674 | 321 | 283 | 461 | 540 | 657 | 238 | 900 | 0.00% | 2.3/sec | |
| Request - Tenant 8 | 706 | 325 | 282 | 458 | 551 | 760 | 234 | 1379 | 0.00% | 2.4/sec | |
| Request - Tenant 9 | 639 | 316 | 280 | 446 | 518 | 666 | 235 | 936 | 0.00% | 2.2/sec | |
| TOTAL | 13599 | 323 | 283 | 456 | 546 | 688 | 223 | 1379 | 0.00% | 46.9/sec | |

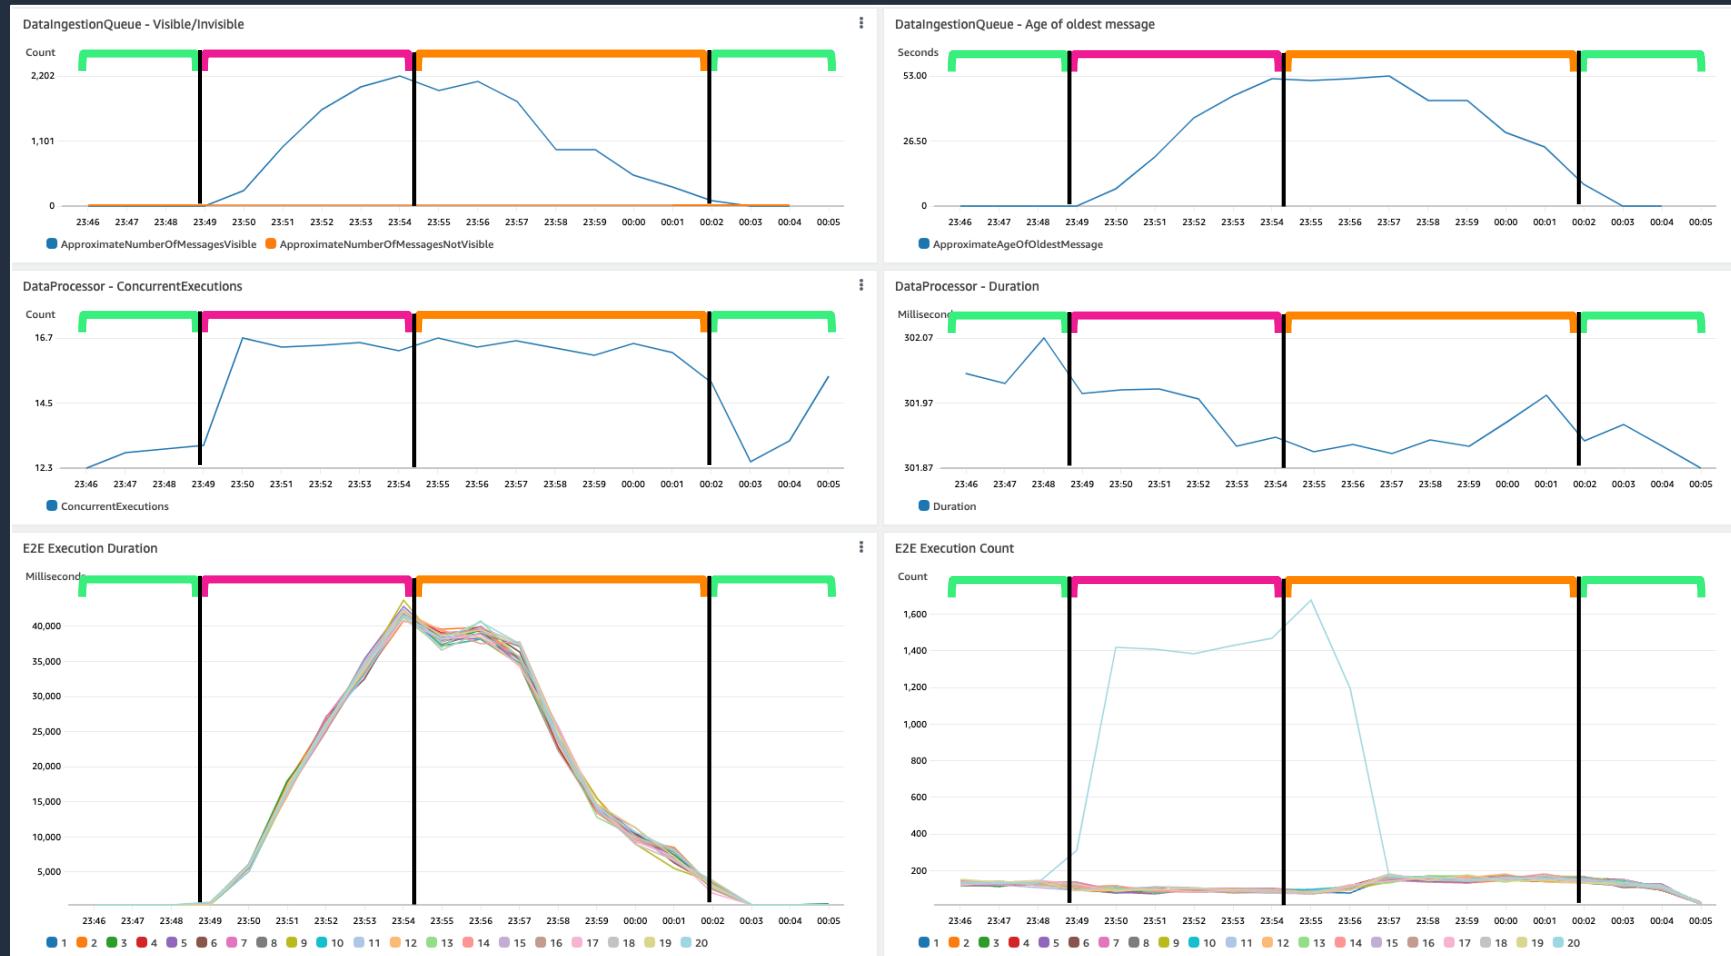
Step 2 – Noisy neighbor!

| Label | # Samples | Average | Median | 90% Line | 95% Line | 99% Line | Min | Maximum | Error % | Throughput | R |
|---------------------|-----------|---------|--------|----------|----------|----------|-----|---------|---------|------------|---|
| Request - Tenant 1 | 673 | 406 | 315 | 589 | 661 | 990 | 221 | 1282 | 0.00% | 2.0/sec | |
| Request - Tenant 10 | 653 | 402 | 311 | 580 | 665 | 864 | 225 | 1143 | 0.00% | 1.9/sec | |
| Request - Tenant 11 | 680 | 404 | 294 | 589 | 720 | 1084 | 228 | 1281 | 0.00% | 2.0/sec | |
| Request - Tenant 12 | 643 | 404 | 304 | 586 | 657 | 1032 | 232 | 1277 | 0.00% | 1.9/sec | |
| Request - Tenant 13 | 646 | 411 | 306 | 615 | 729 | 1043 | 230 | 1118 | 0.00% | 1.9/sec | |
| Request - Tenant 14 | 629 | 404 | 297 | 597 | 678 | 920 | 229 | 1213 | 0.00% | 1.9/sec | |
| Request - Tenant 15 | 664 | 393 | 291 | 578 | 654 | 865 | 222 | 1226 | 0.00% | 2.0/sec | |
| Request - Tenant 16 | 672 | 410 | 324 | 596 | 685 | 1011 | 227 | 1208 | 0.00% | 2.0/sec | |
| Request - Tenant 17 | 642 | 397 | 287 | 588 | 674 | 1051 | 228 | 1250 | 0.00% | 1.9/sec | |
| Request - Tenant 18 | 689 | 397 | 293 | 579 | 673 | 985 | 224 | 1141 | 0.00% | 2.0/sec | |
| Request - Tenant 19 | 655 | 405 | 293 | 590 | 711 | 1012 | 231 | 1150 | 0.00% | 1.9/sec | |
| Request - Tenant 20 | 9968 | 408 | 310 | 590 | 691 | 1060 | 220 | 1579 | 0.00% | 29.3/sec | |
| TOTAL | 22690 | 405 | 304 | 590 | 690 | 1039 | 220 | 2183 | 0.00% | 66.7/sec | |

Results

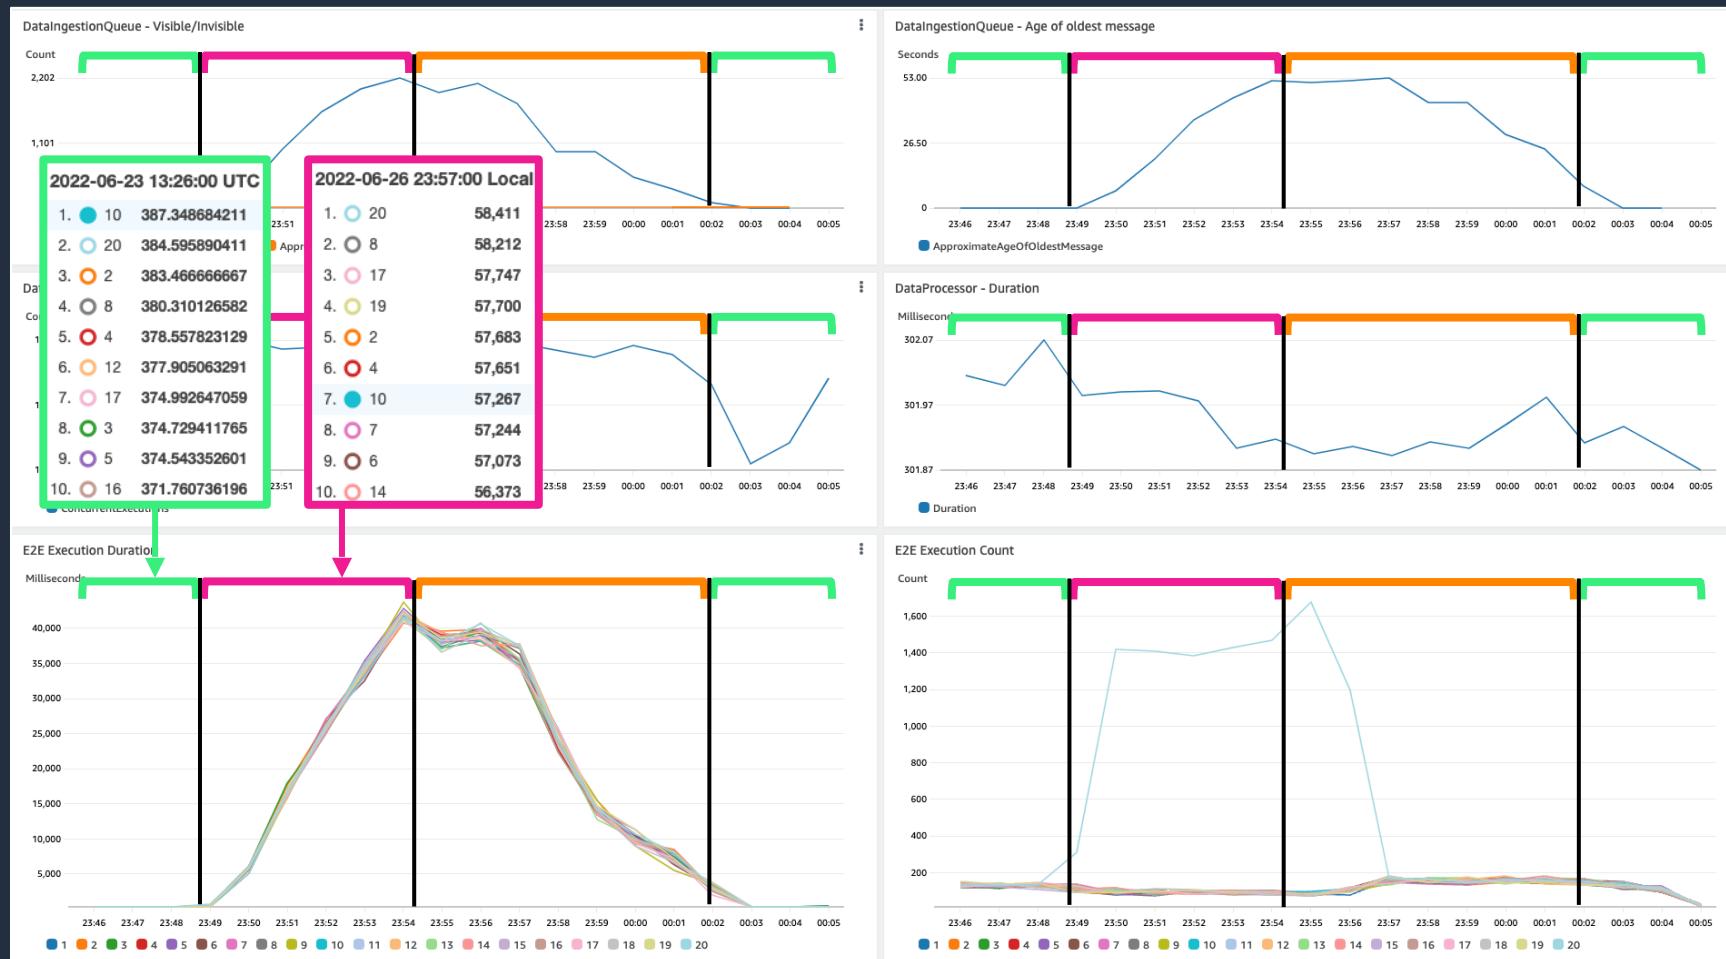


Results



Peace and quiet
Noisy neighbor
Cooldown

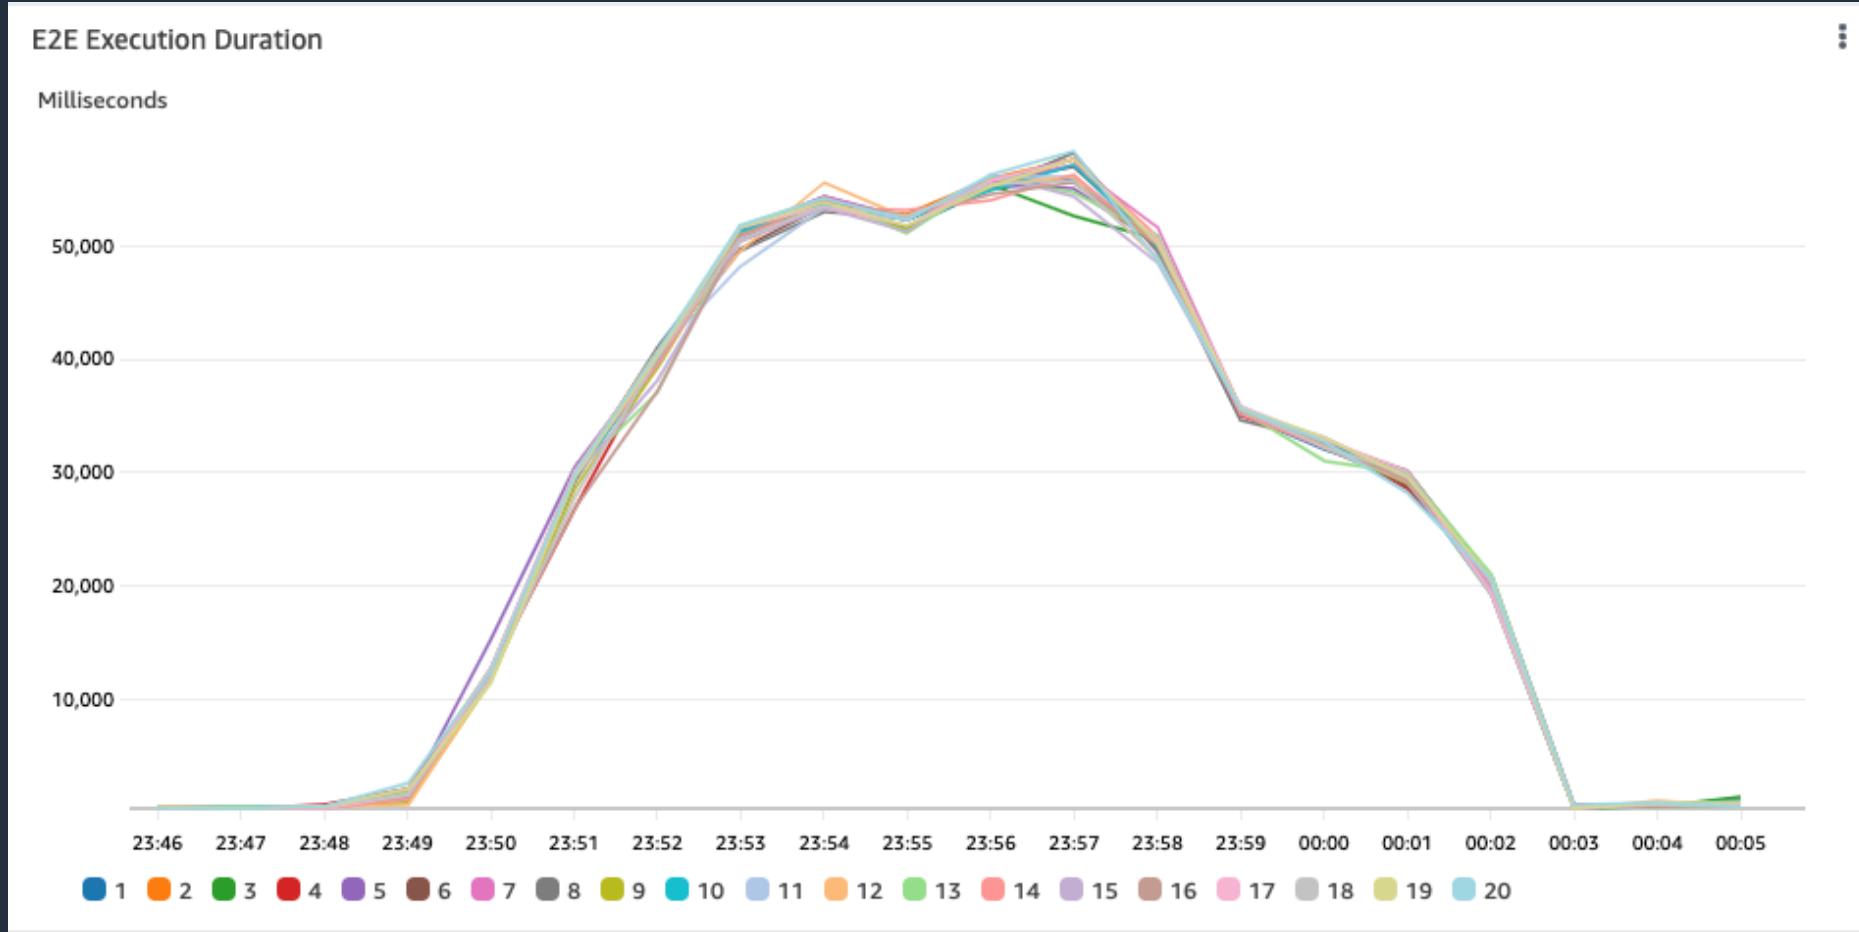
Results



Peace and quiet
Noisy neighbor
Cooldown

Results – P99

P99

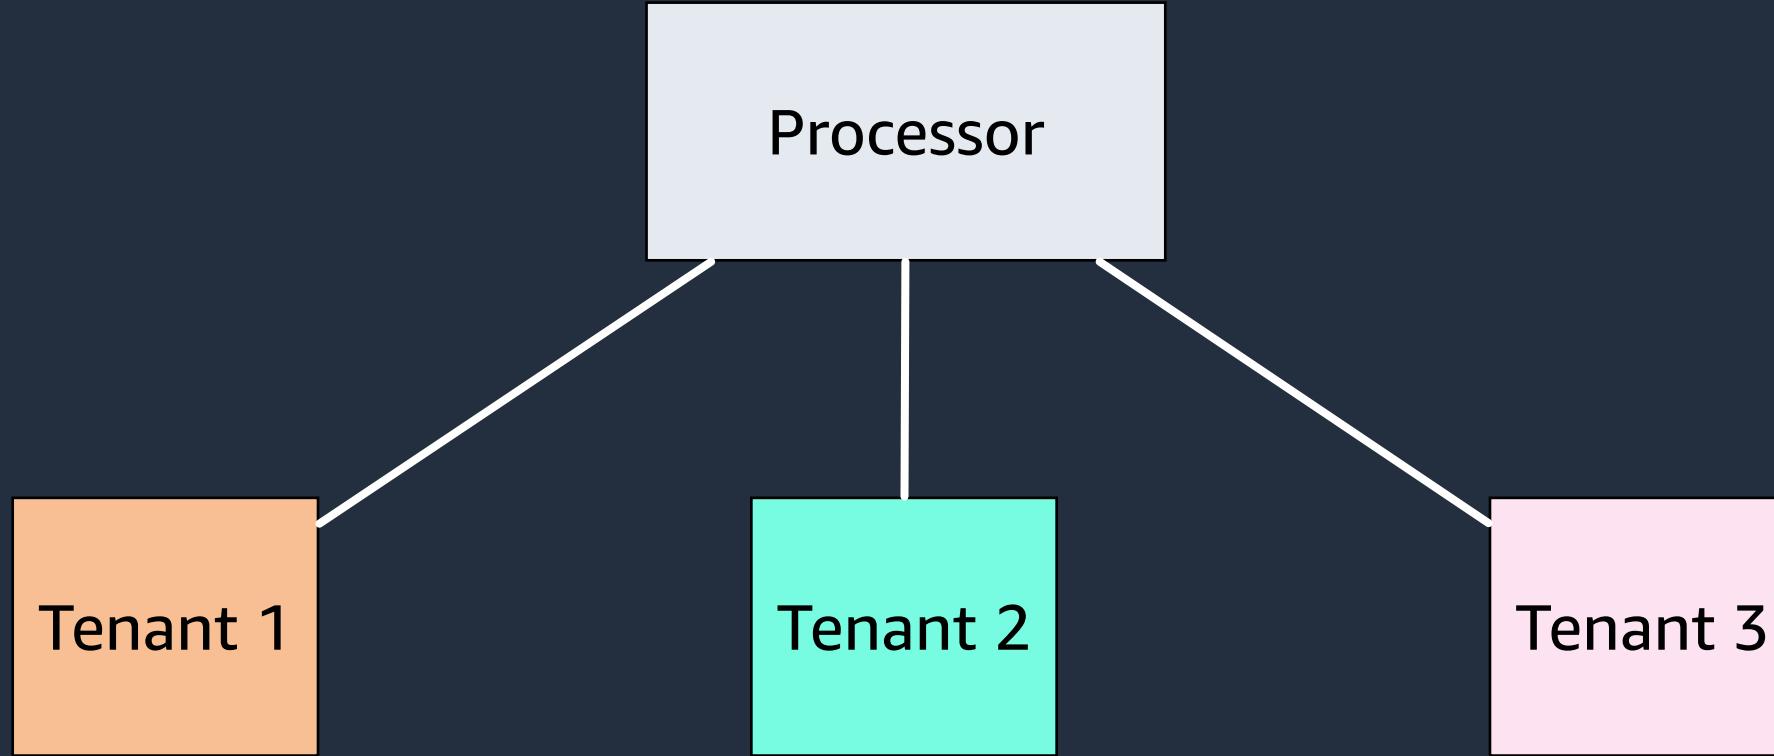


Conclusions

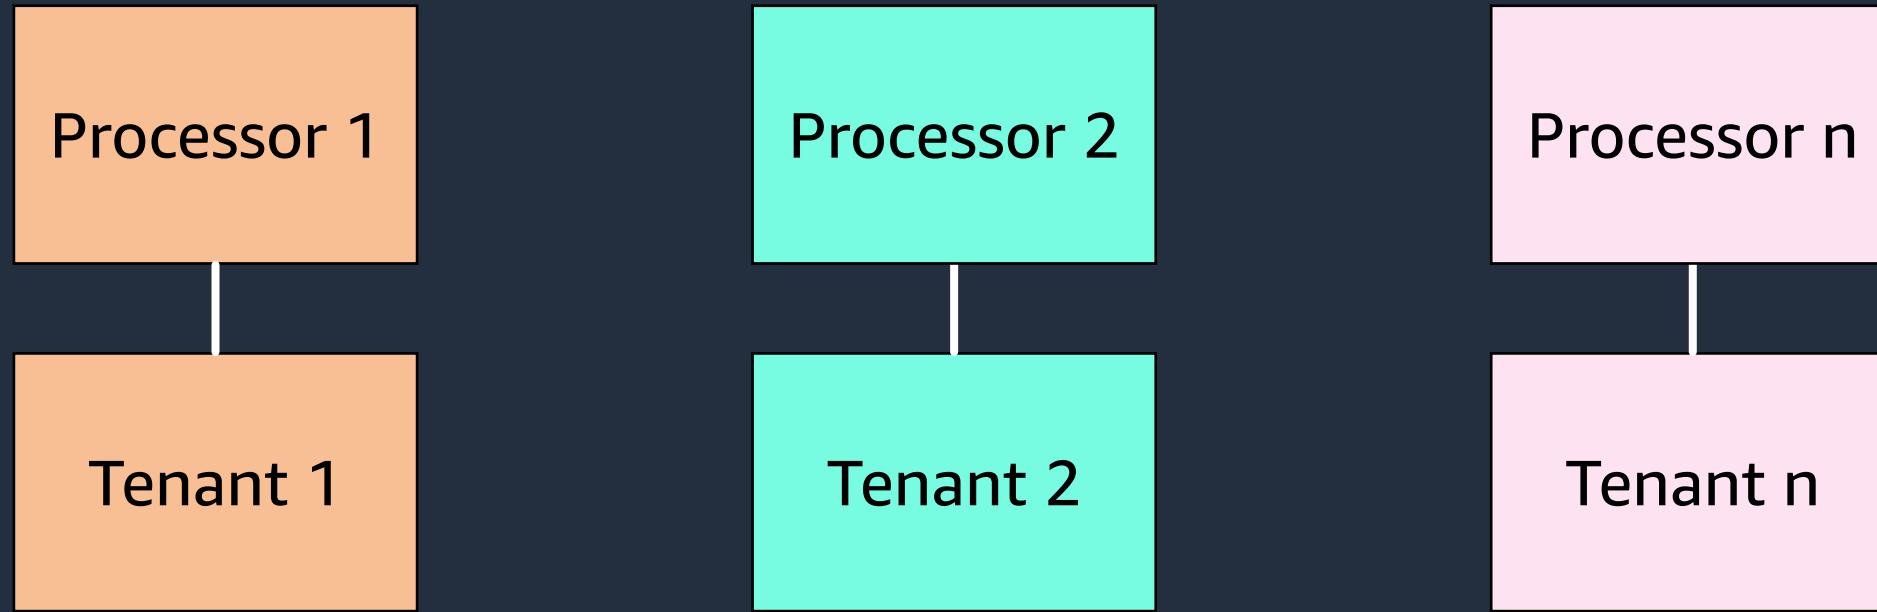
- A single noisy neighbor or a tenant sending “poisonous” requests can have adverse effect on every other tenant
- Building decoupled asynchronous systems with messaging services like SQS helps with vertical fault isolation
- Use API Gateway throttling mechanism where applicable
- Cooldown depends on the volume of spare capacity
- Consider fault isolation techniques when building multi-tenant SaaS systems

Brainstorming a solution

Single processor

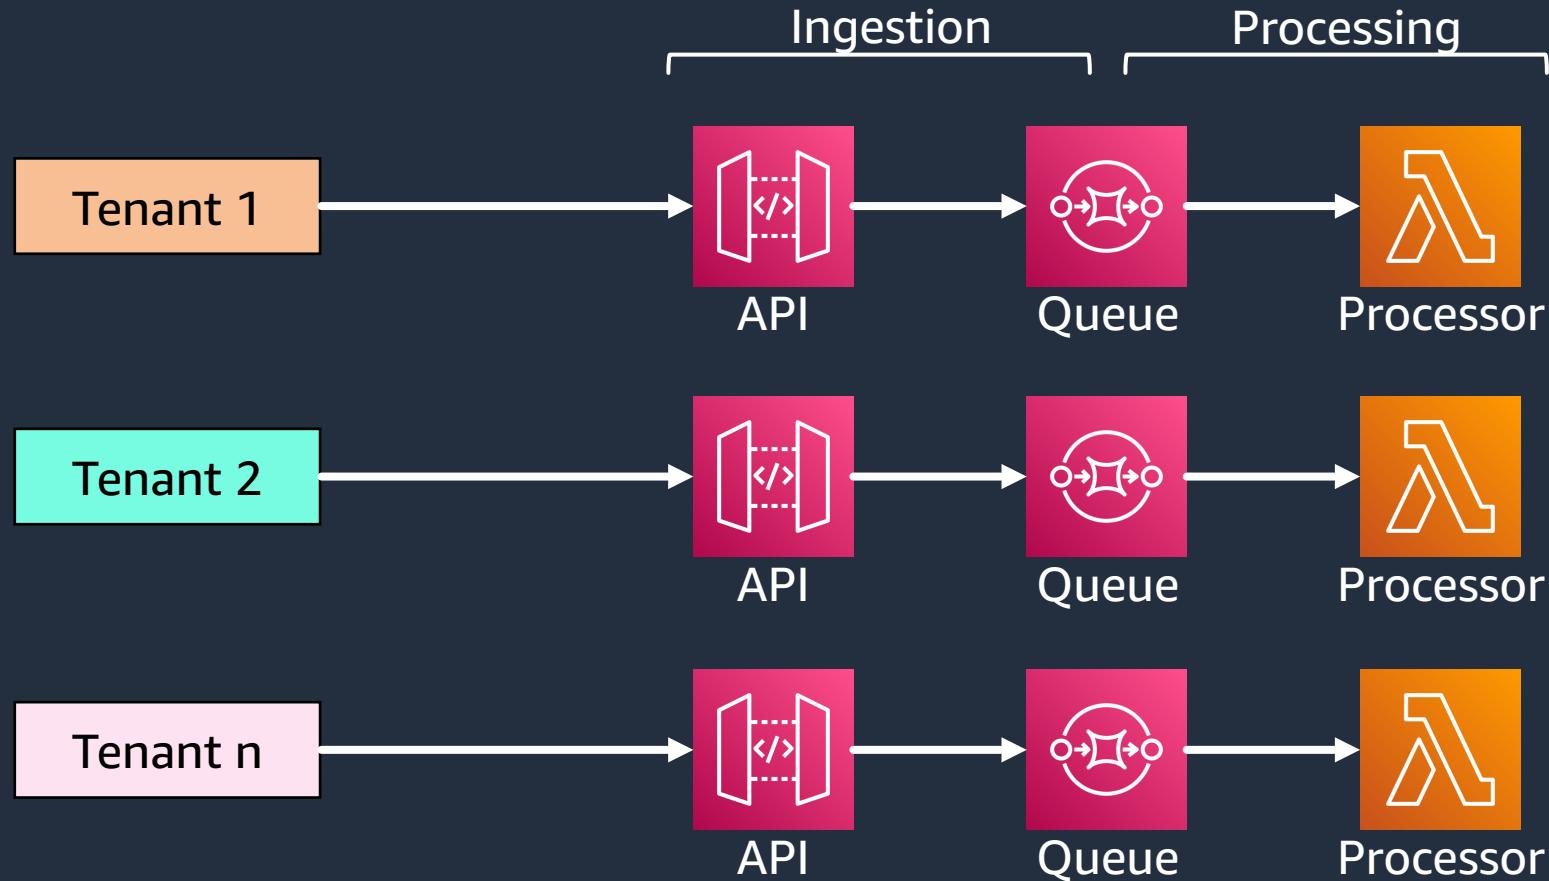


Processor per tenant

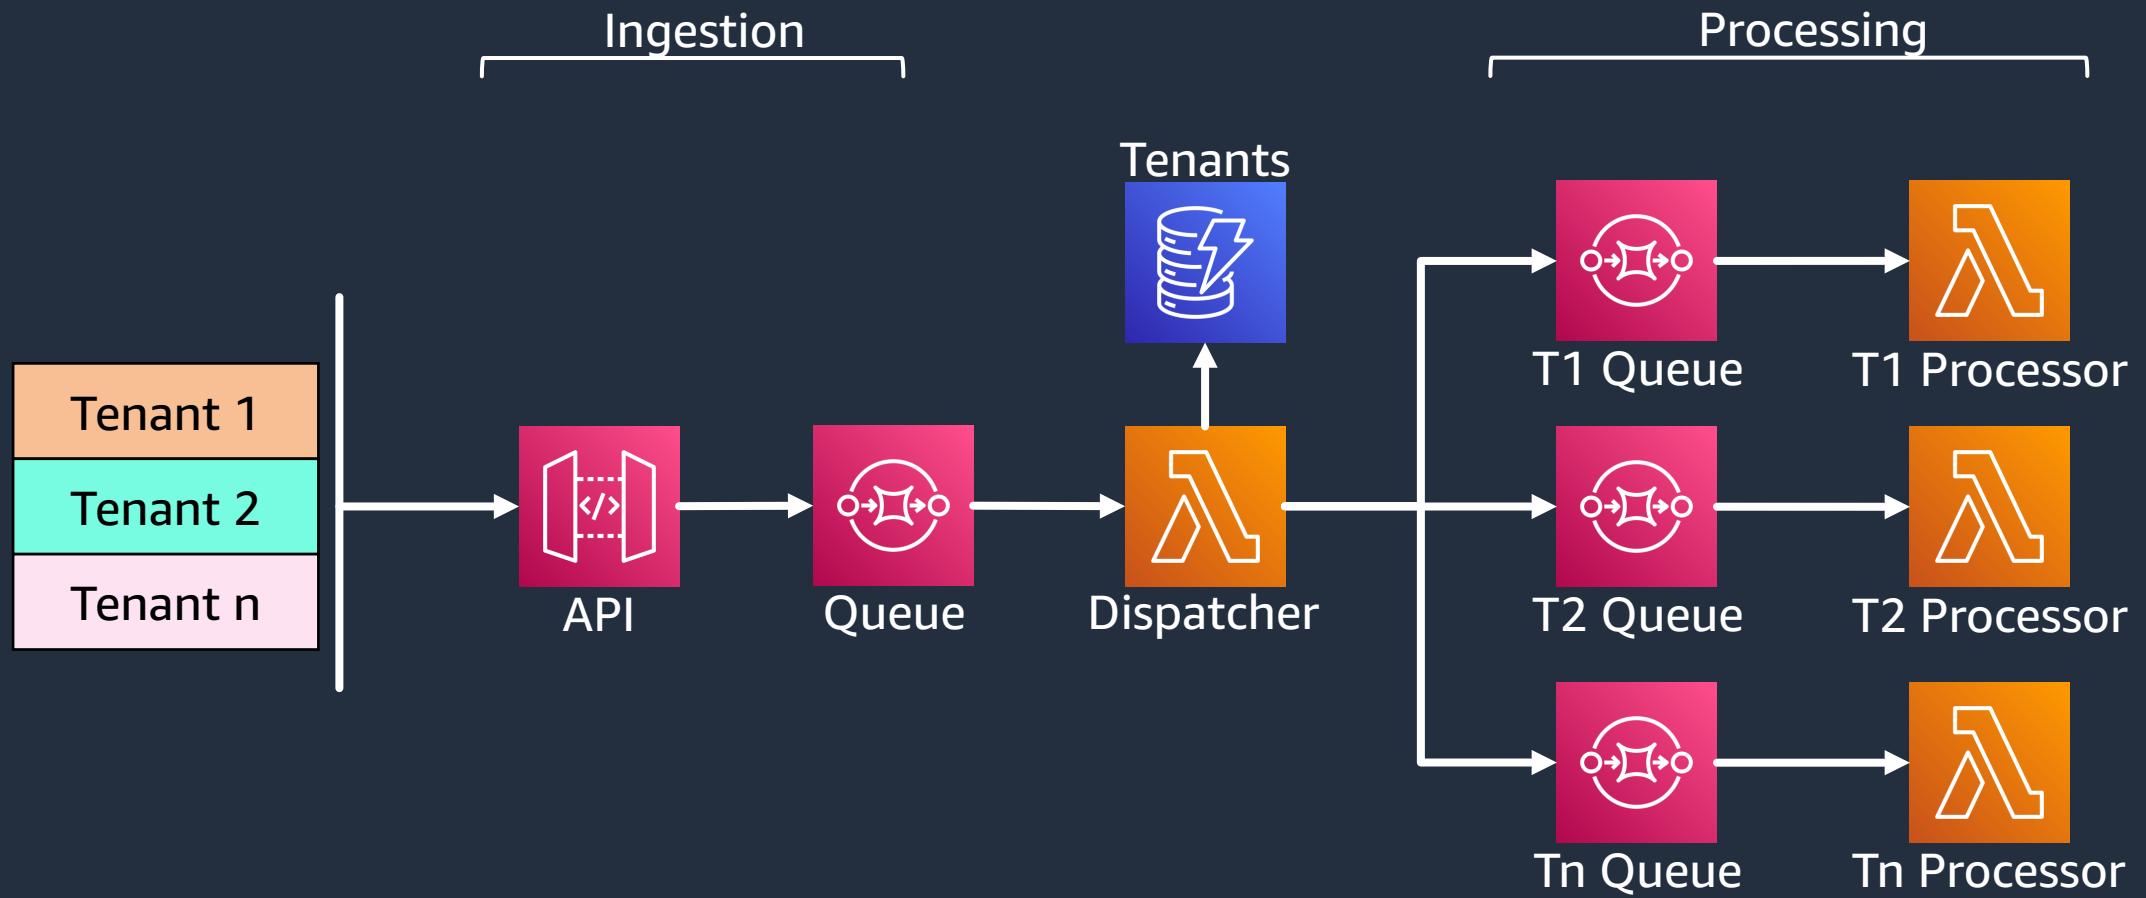


P==T

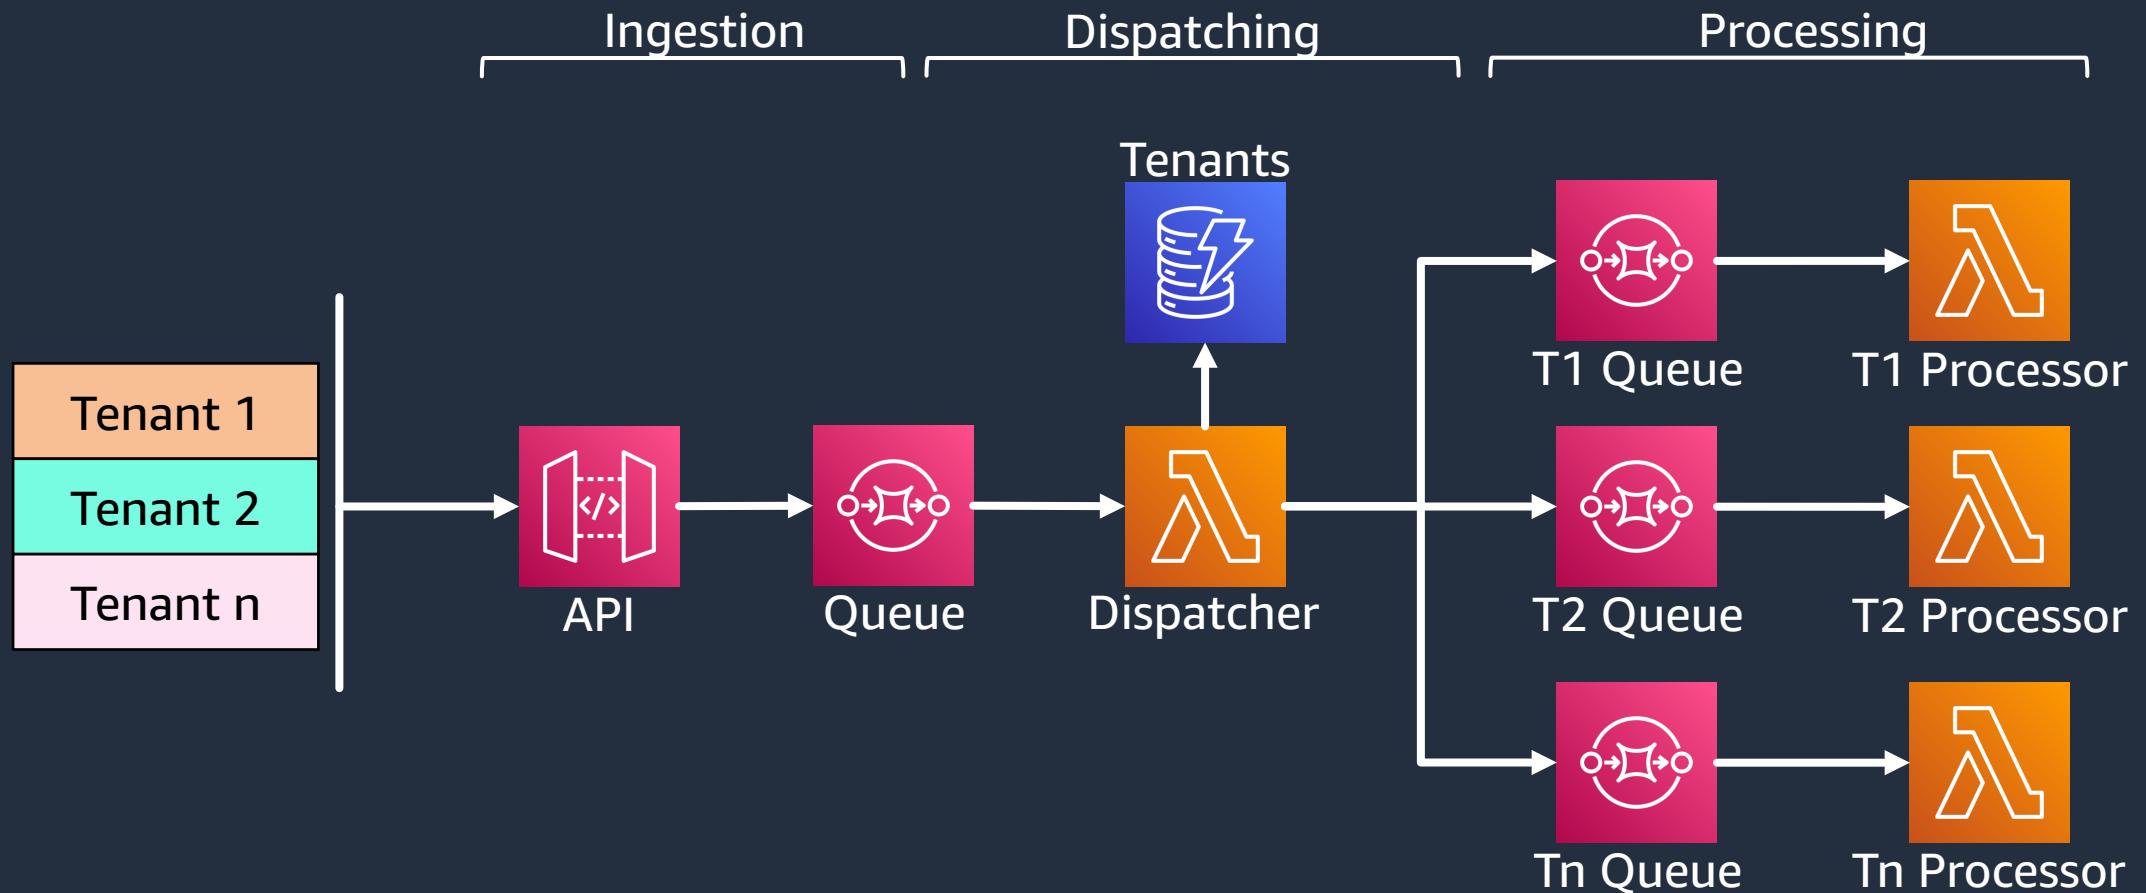
Processor per tenant – sharing nothing



Processor per tenant – sharing ingestion



Processor per tenant – sharing ingestion



Comparing the options

$P == 1$

One processor for all tenants

\$

Vertical scaling

Major noisy neighbor
problem

$P << T$

Number of processors is significantly
lower than number of tenants

\$\$

Sublinear scaling (good!)

But would it address the noisy
neighbor problem?

$P == T$

Number of processors equals
to number of tenants

\$\$\$\$

Linear scaling (bad!)

No noisy neighbor
problem!

Sharding 101

| ID | Name | Stock | Price |
|----|--------------|-------|--------|
| 1 | Refrigerator | 75 | \$1800 |
| 2 | Dish washer | 102 | \$700 |
| 3 | TV | 39 | \$1500 |
| 4 | Camera | 55 | \$400 |

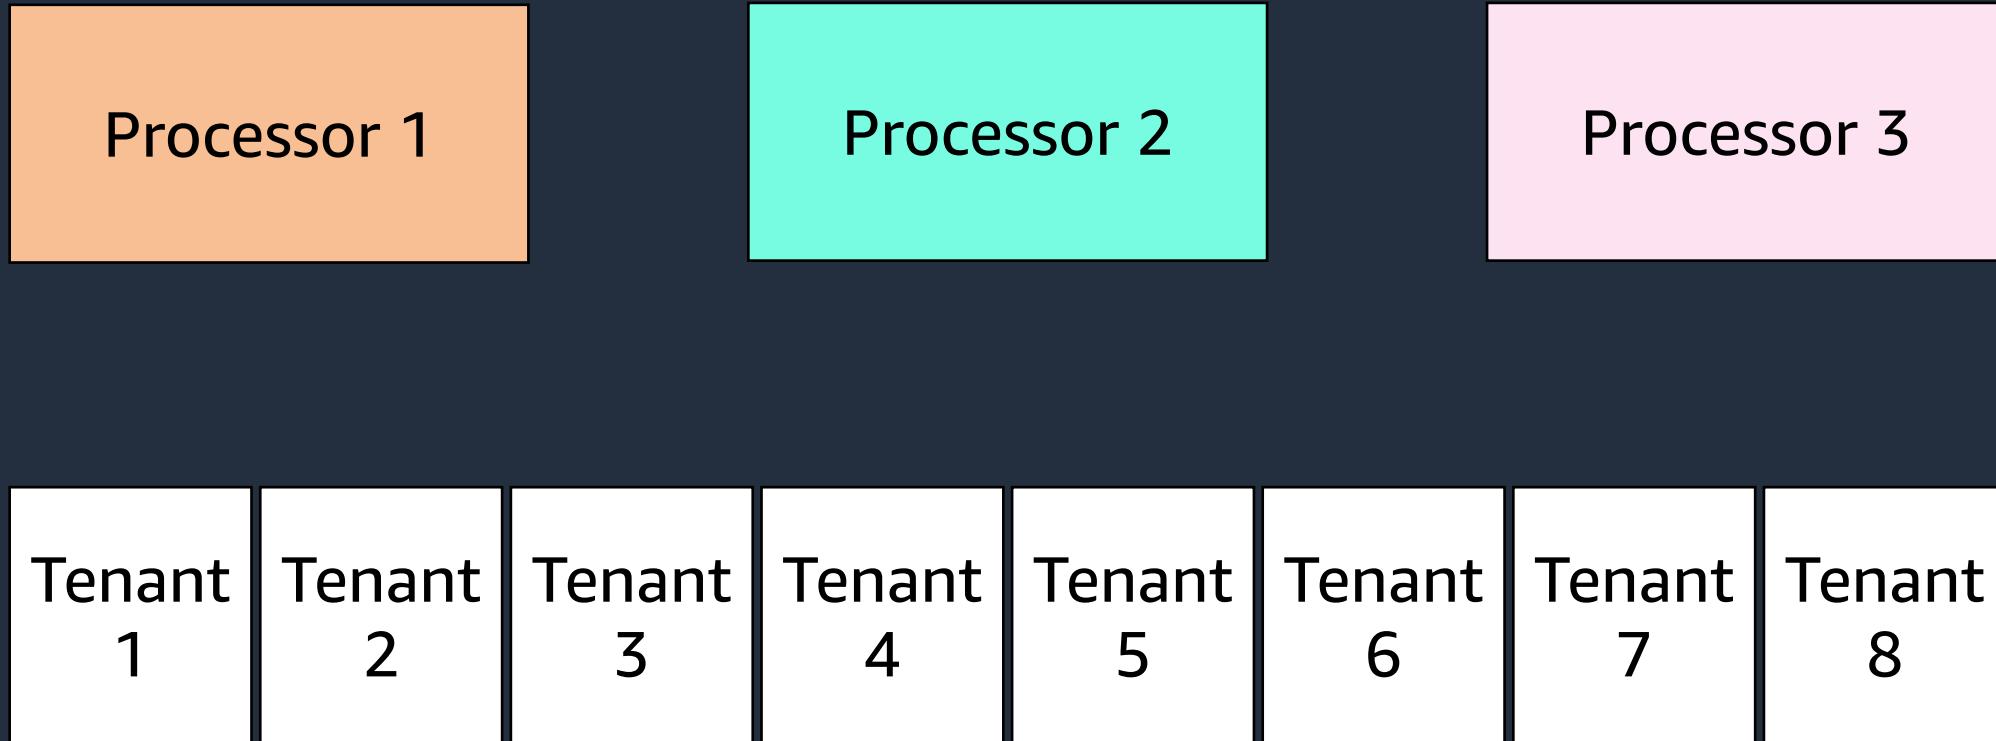
| Vertical Shard 1 | | |
|------------------|--------------|--------|
| ID | Name | Price |
| 1 | Refrigerator | \$1800 |
| 2 | Dish washer | \$700 |
| 3 | TV | \$1500 |
| 4 | Camera | \$400 |

| Vertical Shard 2 | |
|------------------|-------|
| ID | Stock |
| 1 | 75 |
| 2 | 102 |
| 3 | 39 |
| 4 | 55 |

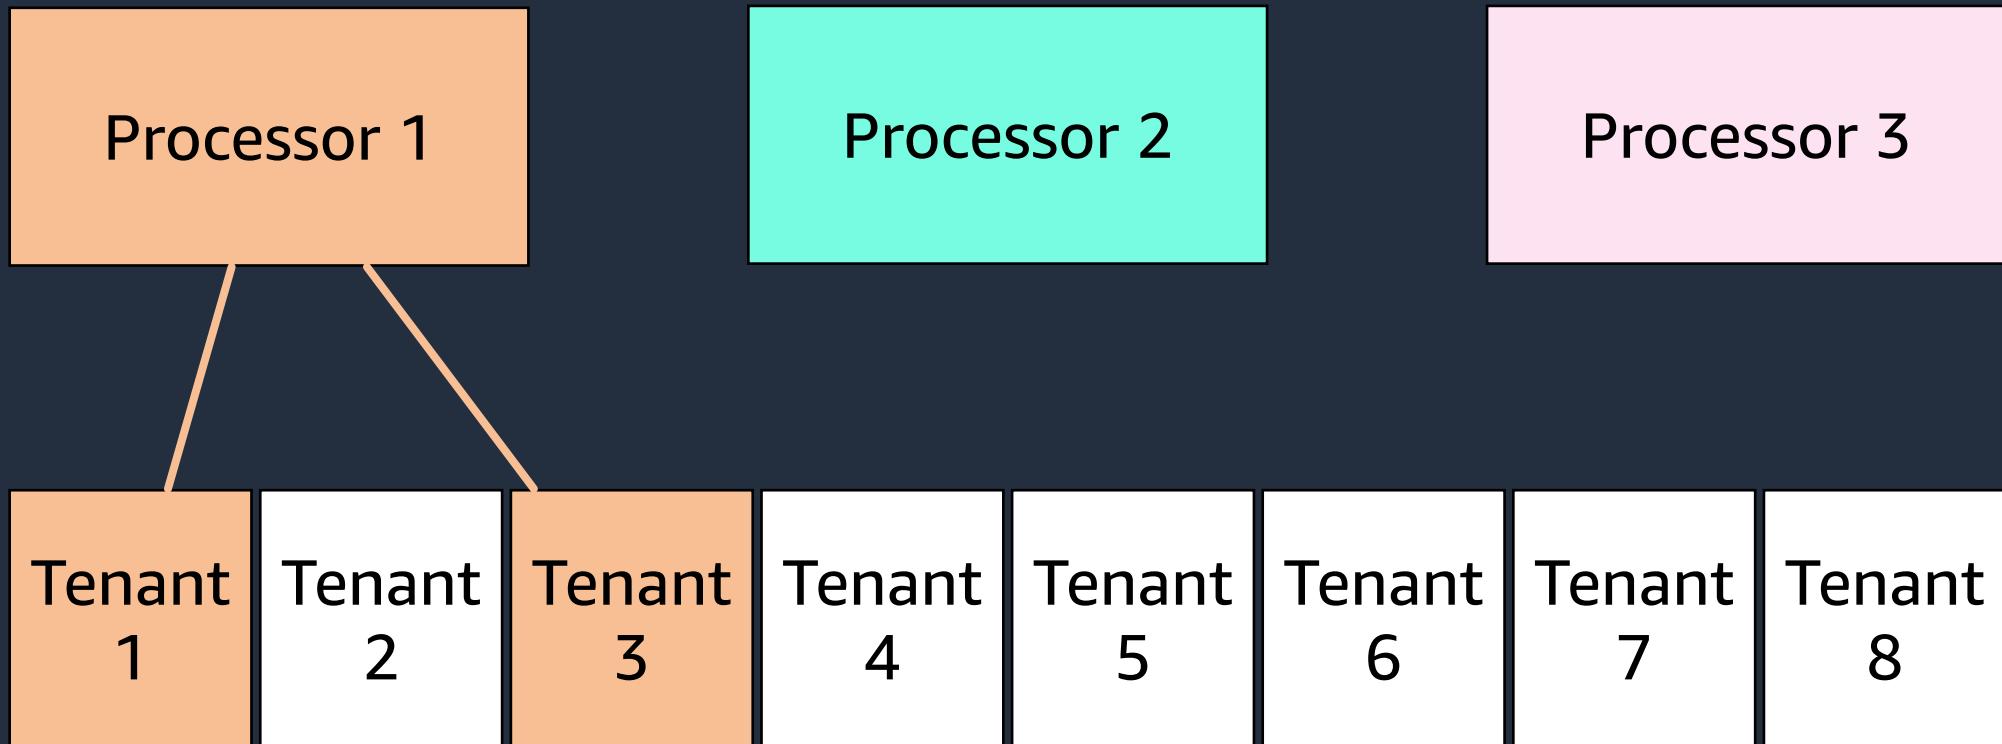
| Horizontal Shard 1 | | | |
|--------------------|--------------|-------|--------|
| ID | Name | Stock | Price |
| 1 | Refrigerator | 75 | \$1800 |
| 2 | Dish washer | 102 | \$700 |

| Horizontal Shard 2 | | | |
|--------------------|--------|-------|--------|
| ID | Name | Stock | Price |
| 3 | TV | 39 | \$1500 |
| 4 | Camera | 55 | \$400 |

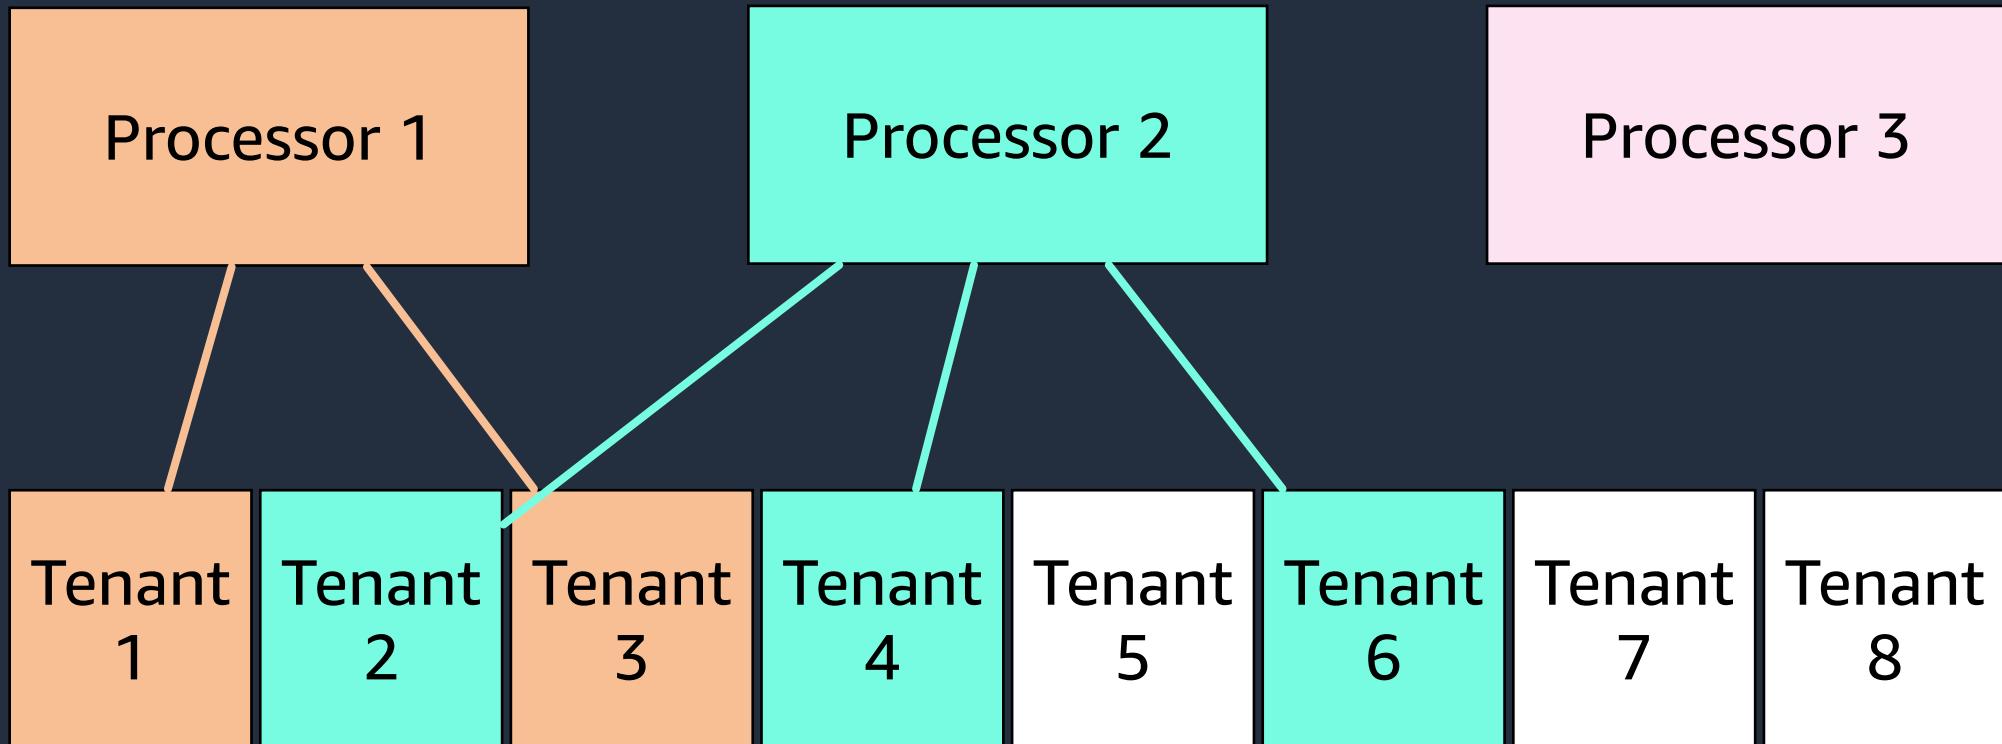
Simple sharding



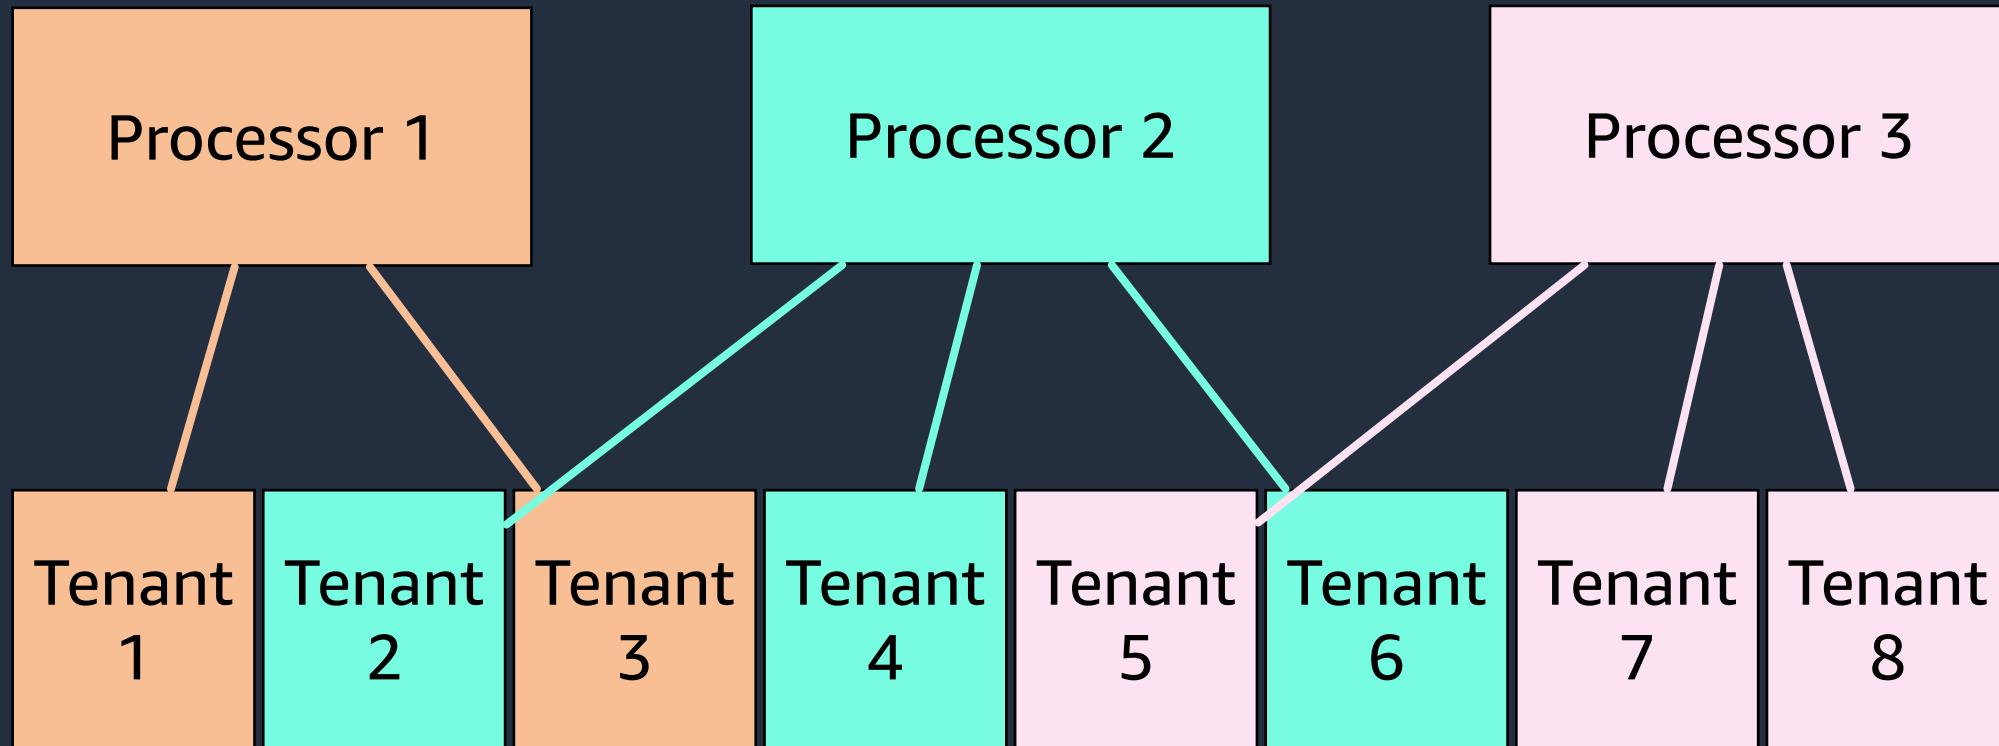
Simple sharding



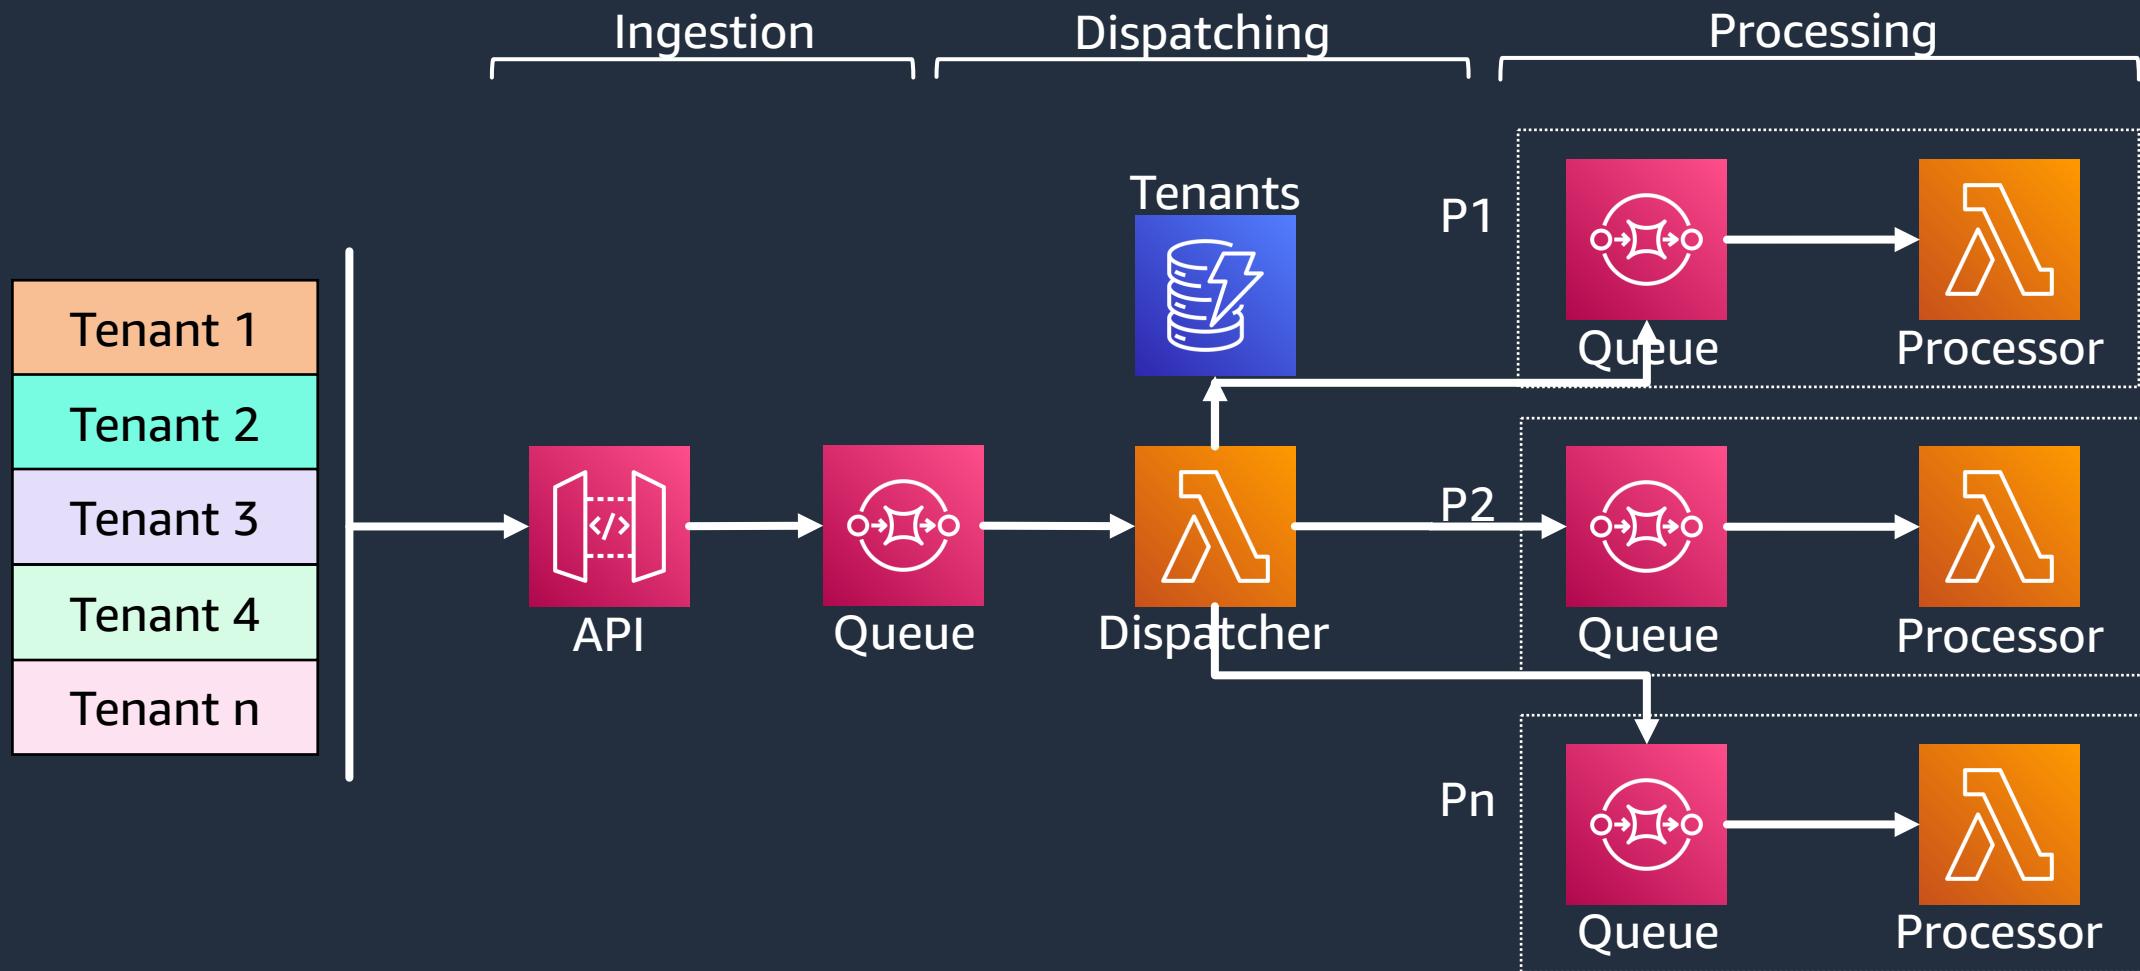
Simple sharding



Simple sharding



Simple sharding: $P \ll T$



Let's see this in action

Noisy neighbor!

| Label | # Samples | Average | Median | 90% Line | 95% Line | 99% Line | Min | Maximum | Error % | Throughput | R |
|---------------------|-----------|---------|--------|----------|----------|----------|-----|---------|---------|------------|---|
| Request - Tenant 1 | 673 | 406 | 315 | 589 | 661 | 990 | 221 | 1282 | 0.00% | 2.0/sec | |
| Request - Tenant 10 | 653 | 402 | 311 | 580 | 665 | 864 | 225 | 1143 | 0.00% | 1.9/sec | |
| Request - Tenant 11 | 680 | 404 | 294 | 589 | 720 | 1084 | 228 | 1281 | 0.00% | 2.0/sec | |
| Request - Tenant 12 | 643 | 404 | 304 | 586 | 657 | 1032 | 232 | 1277 | 0.00% | 1.9/sec | |
| Request - Tenant 13 | 646 | 411 | 306 | 615 | 729 | 1043 | 230 | 1118 | 0.00% | 1.9/sec | |
| Request - Tenant 14 | 629 | 404 | 297 | 597 | 678 | 920 | 229 | 1213 | 0.00% | 1.9/sec | |
| Request - Tenant 15 | 664 | 393 | 291 | 578 | 654 | 865 | 222 | 1226 | 0.00% | 2.0/sec | |
| Request - Tenant 16 | 672 | 410 | 324 | 596 | 685 | 1011 | 227 | 1208 | 0.00% | 2.0/sec | |
| Request - Tenant 17 | 642 | 397 | 287 | 588 | 674 | 1051 | 228 | 1250 | 0.00% | 1.9/sec | |
| Request - Tenant 18 | 689 | 397 | 293 | 579 | 673 | 985 | 224 | 1141 | 0.00% | 2.0/sec | |
| Request - Tenant 19 | 655 | 405 | 293 | 590 | 711 | 1012 | 231 | 1150 | 0.00% | 1.9/sec | |
| Request - Tenant 20 | 9968 | 408 | 310 | 590 | 691 | 1060 | 220 | 1579 | 0.00% | 29.3/sec | |
| TOTAL | 22690 | 405 | 304 | 590 | 690 | 1039 | 220 | 2183 | 0.00% | 66.7/sec | |

Before we begin

Affected tenants

Noisy neighbor

Tenants Table - DynamoDB

| | id | queueUrl |
|--|----|---|
| | 14 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Simple-P8-DataProcessorQueue433DB58A-IP8lbqqPb9Zc |
| | 5 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Simple-P7-DataProcessorQueueD4706613-E2yH2j28PFUy |
| | 7 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Simple-P7-DataProcessorQueueD4706613-E2yH2j28PFUy |
| | 20 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Simple-P7-DataProcessorQueueD4706613-E2yH2j28PFUy |
| | 8 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Simple-P6-DataProcessorQueueE3424521-9kjvvuXs0loU |

Before we begin

Affected tenants

Noisy neighbor

Tenants Table - DynamoDB

| | id | queueUrl |
|--|----|--|
| | 14 | https://sqs.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Simple-P8-DataProcessorQueue433DB58A-IP8lbqqPb9Zc |
| | 5 | https://sqs.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Simple-P7-DataProcessorQueueD4706613-E2yH2j28PFUy |
| | 7 | https://sqs.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Simple-P7-DataProcessorQueueD4706613-E2yH2j28PFUy |
| | 20 | https://sqs.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Simple-P7-DataProcessorQueueD4706613-E2yH2j28PFUy |
| | 8 | https://sqs.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Simple-P6-DataProcessorQueueE3424521-9kjvvuXs0loU |

Mental note

Noisy neighbor – 20
Affected tenants – 5,7
Affected processor – P7

Before we begin

Tenants Table - DynamoDB

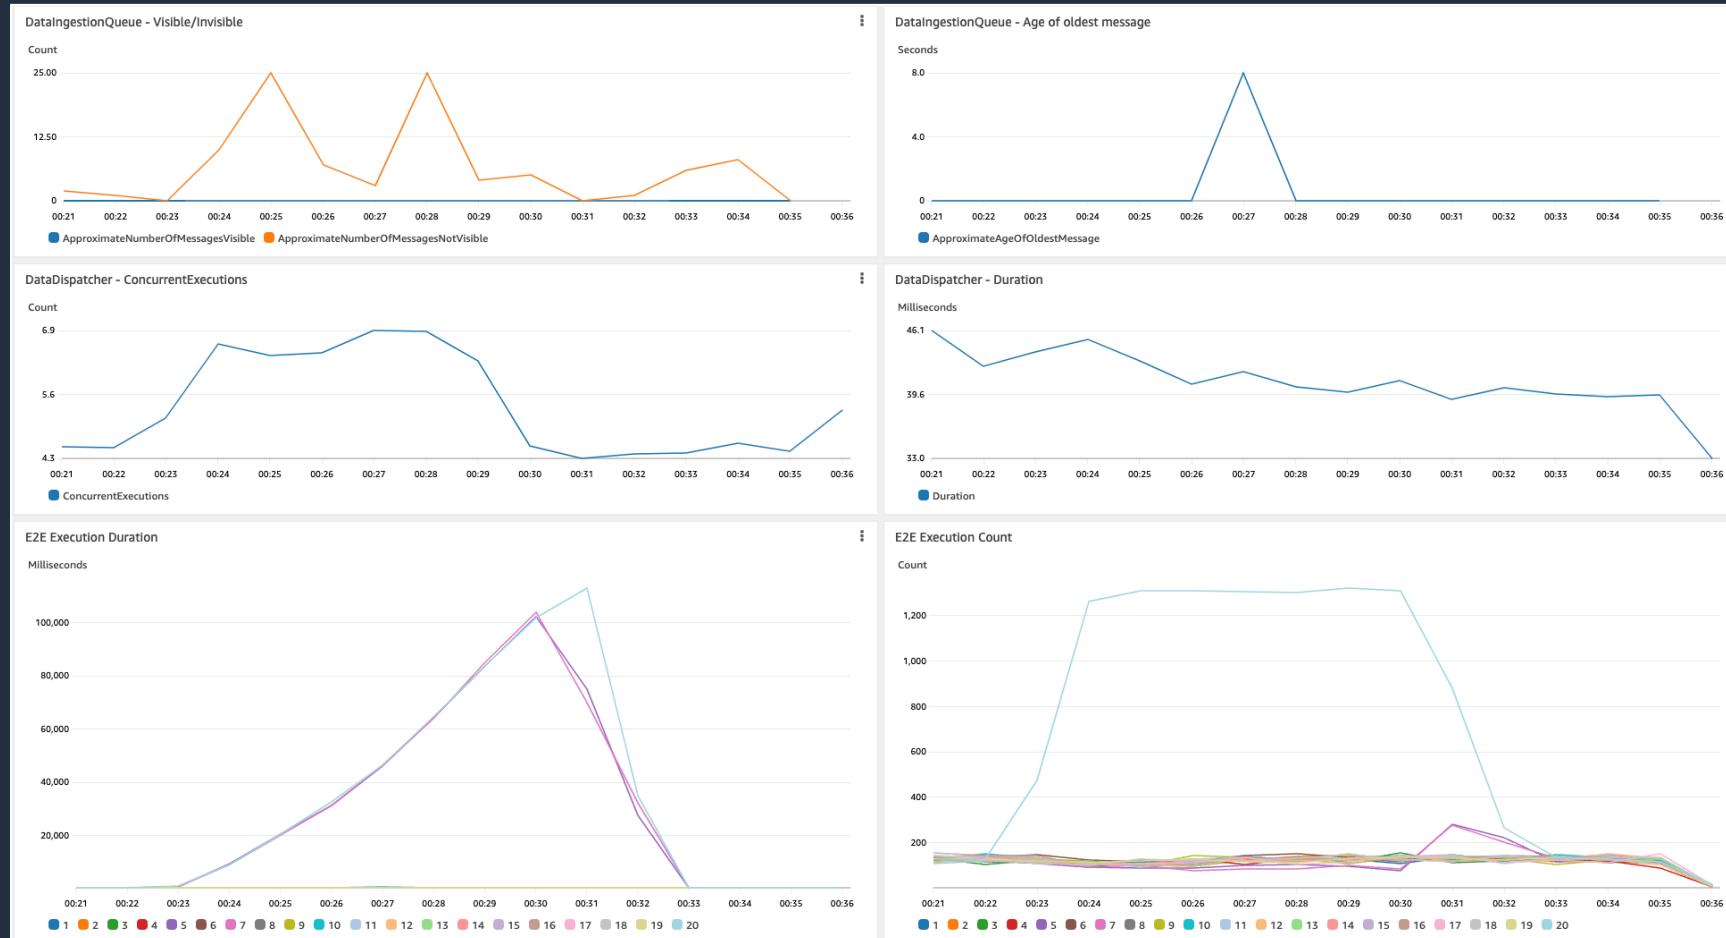
Affected tenants

Noisy neighbor

| | id | queueUrl |
|--|-----------|---|
| | 14 | https://sq...433DB58A-IP8lbqqPb9Zc |
| | 5 | https://sq...D4706613-E2yH2j28PFUy |
| | 7 | https://sq...D4706613-E2yH2j28PFUy |
| | 20 | https://sq...D4706613-E2yH2j28PFUy |
| | 8 | https://sq...E3424521-9kjvvuXs0loU |

Results

Mental note
 Noisy neighbor – 20
 Affected tenants – 5,7
 Affected processor – P7



Results

Mental note
 Noisy neighbor – 20
 Affected tenants – 5,7
 Affected processor – P7



Peace and quiet

Noisy neighbor

Coldown

Results

Mental note
 Noisy neighbor – 20
 Affected tenants – 5,7
 Affected processor – P7



Peace and quiet

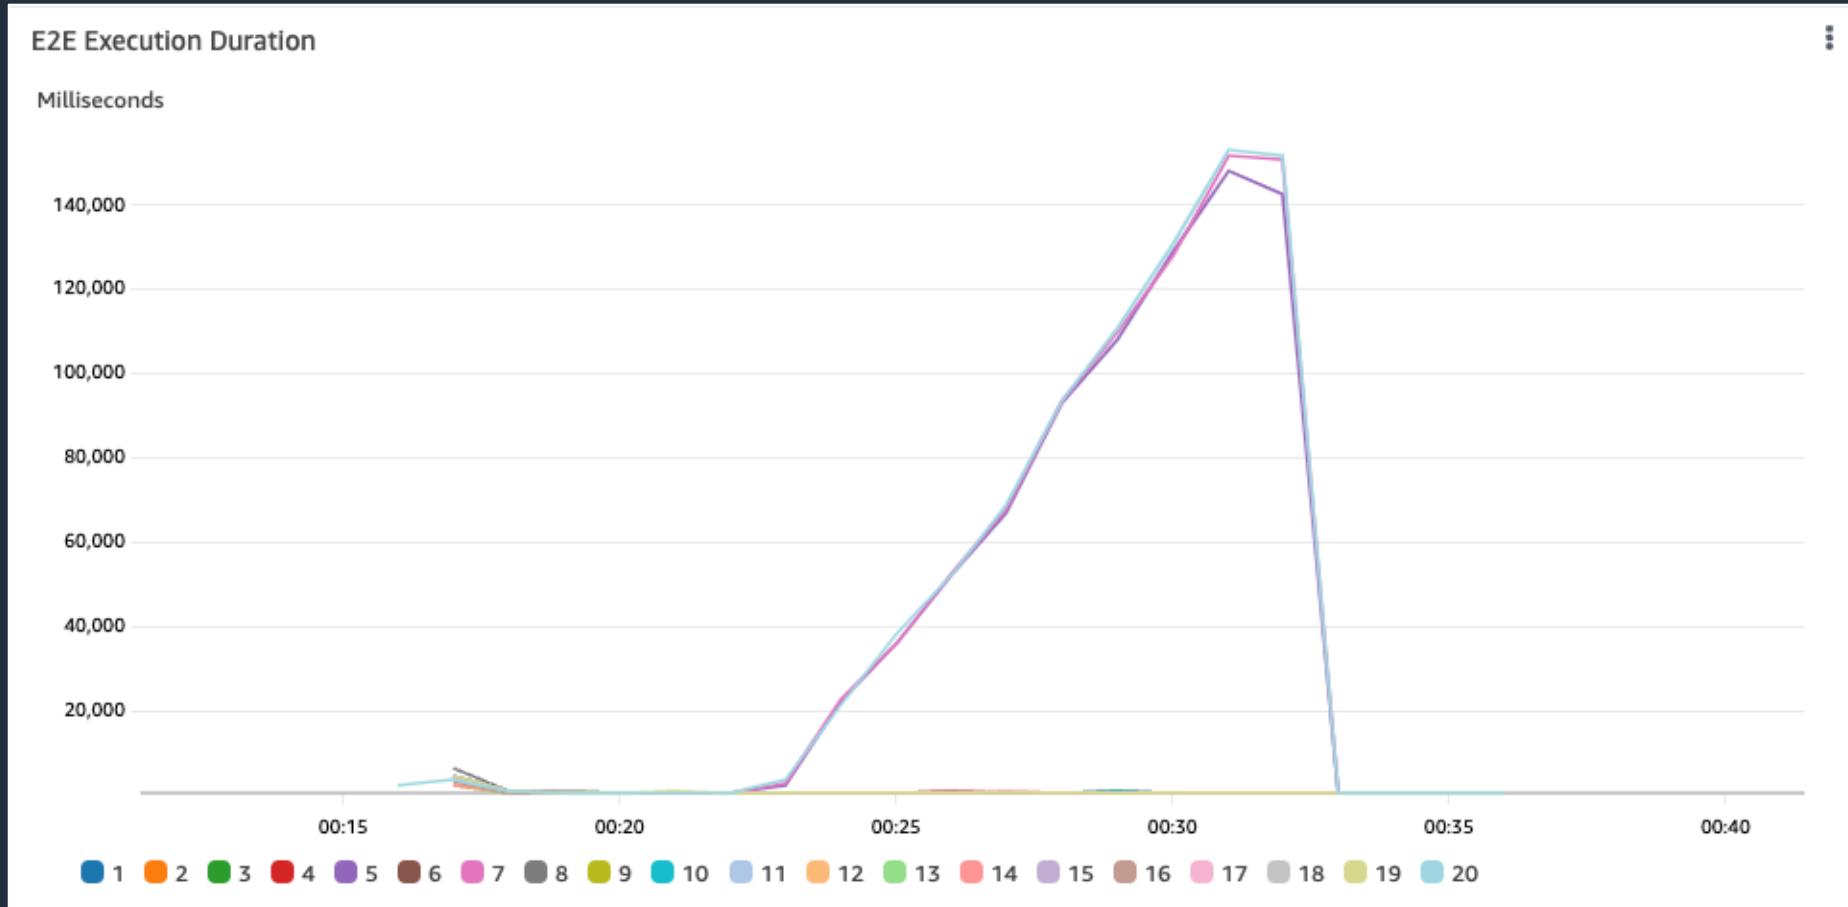
Noisy neighbor

Coldown

Results – P99

Mental note
Noisy neighbor – 20
Affected tenants – 5,7
Affected processor – P7

P99



Mental note
 Noisy neighbor – 20
 Affected tenants – 5,7
 Affected processor – P7

Results



Conclusions

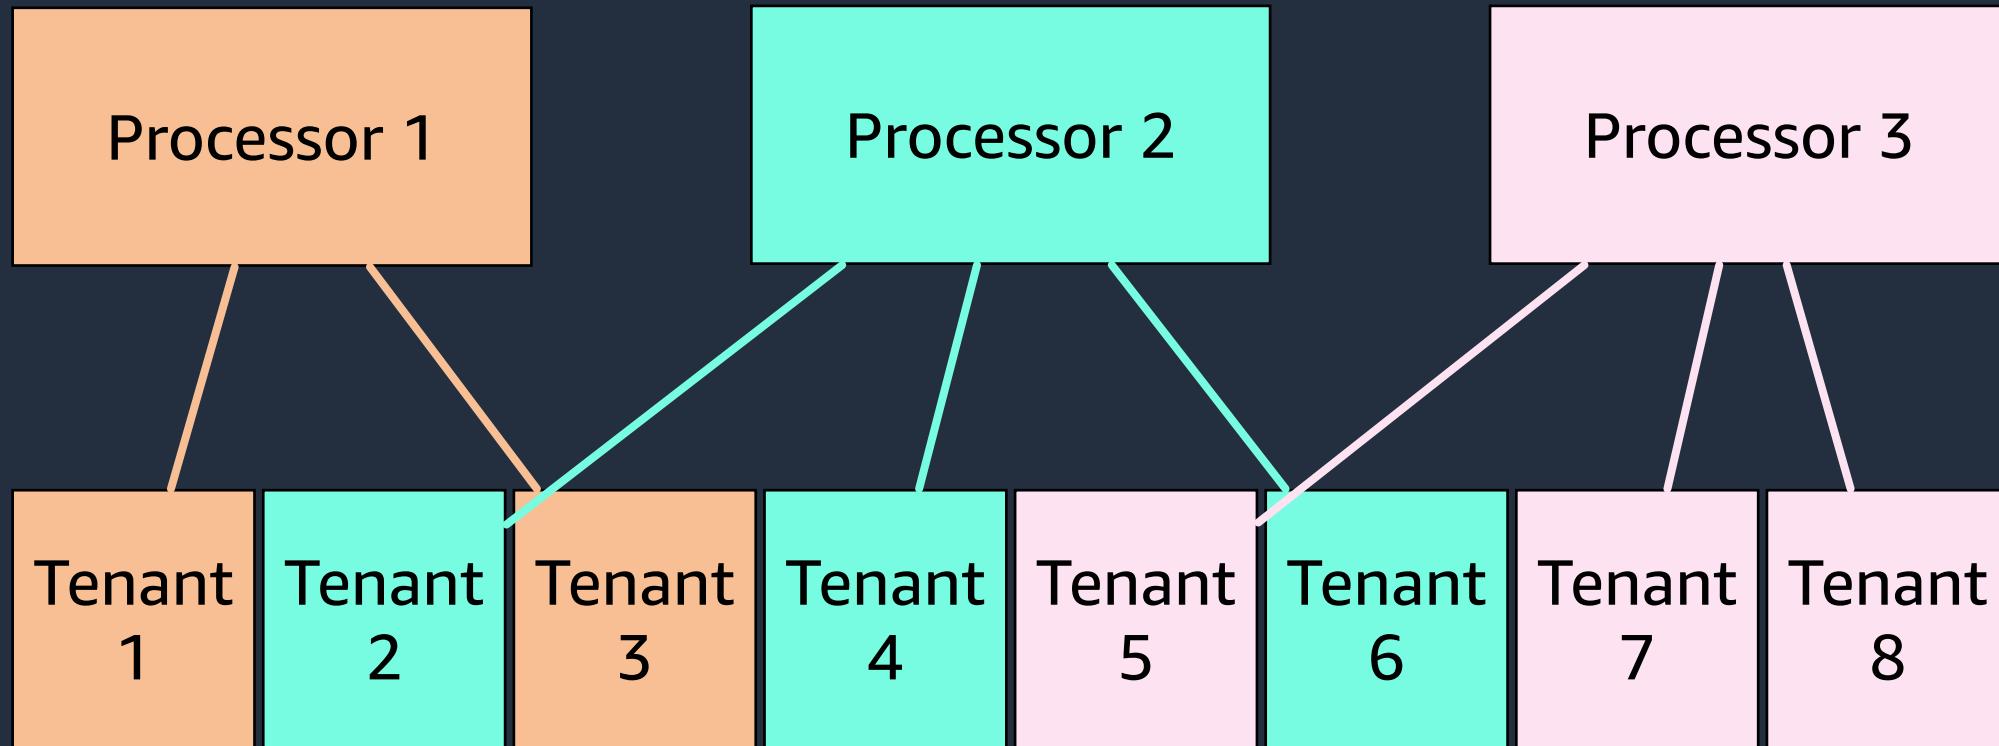
Number of affected tenants = roundup $\left(\frac{\text{number of tenants}}{\text{number of processors}}\right)$ = $\frac{20}{8}$
= **3**

Blast radius = $\frac{\text{number of affected tenants}}{\text{total number of tenants}}$ = $\frac{3}{20}$ = **15%**

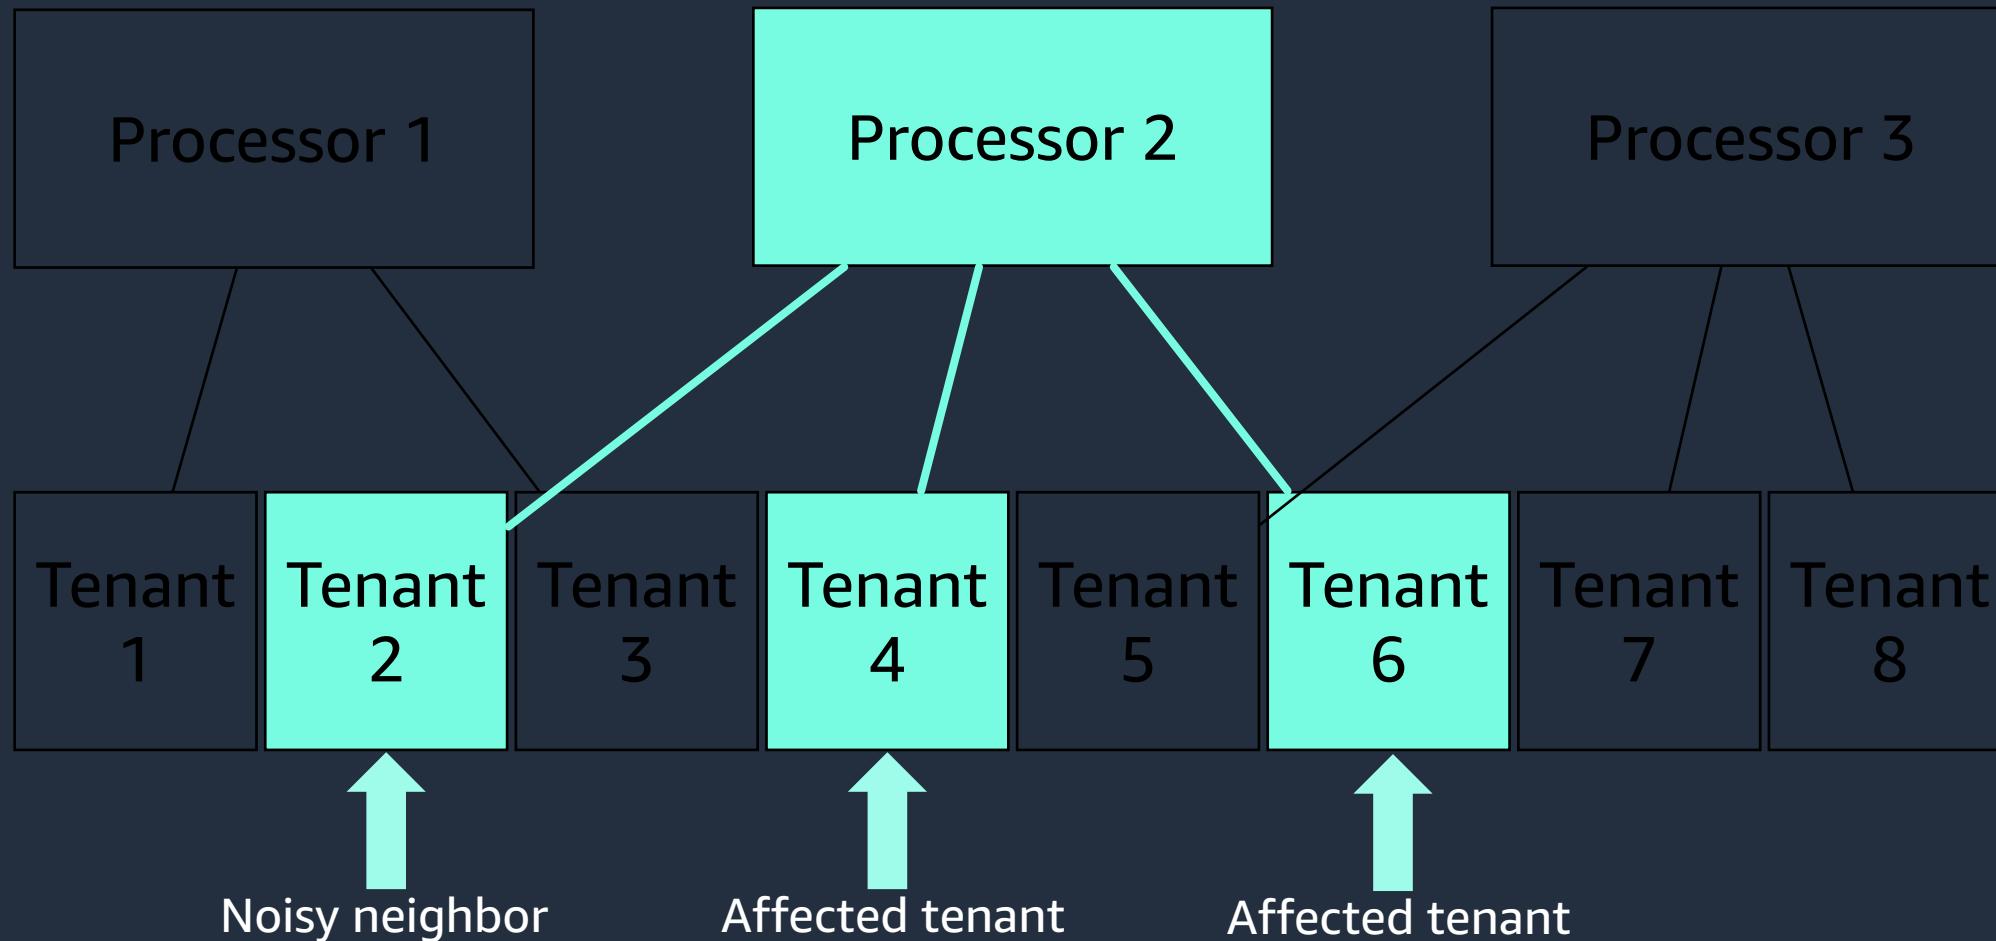
Not bad. But can we do even better?!

Let's brainstorm further

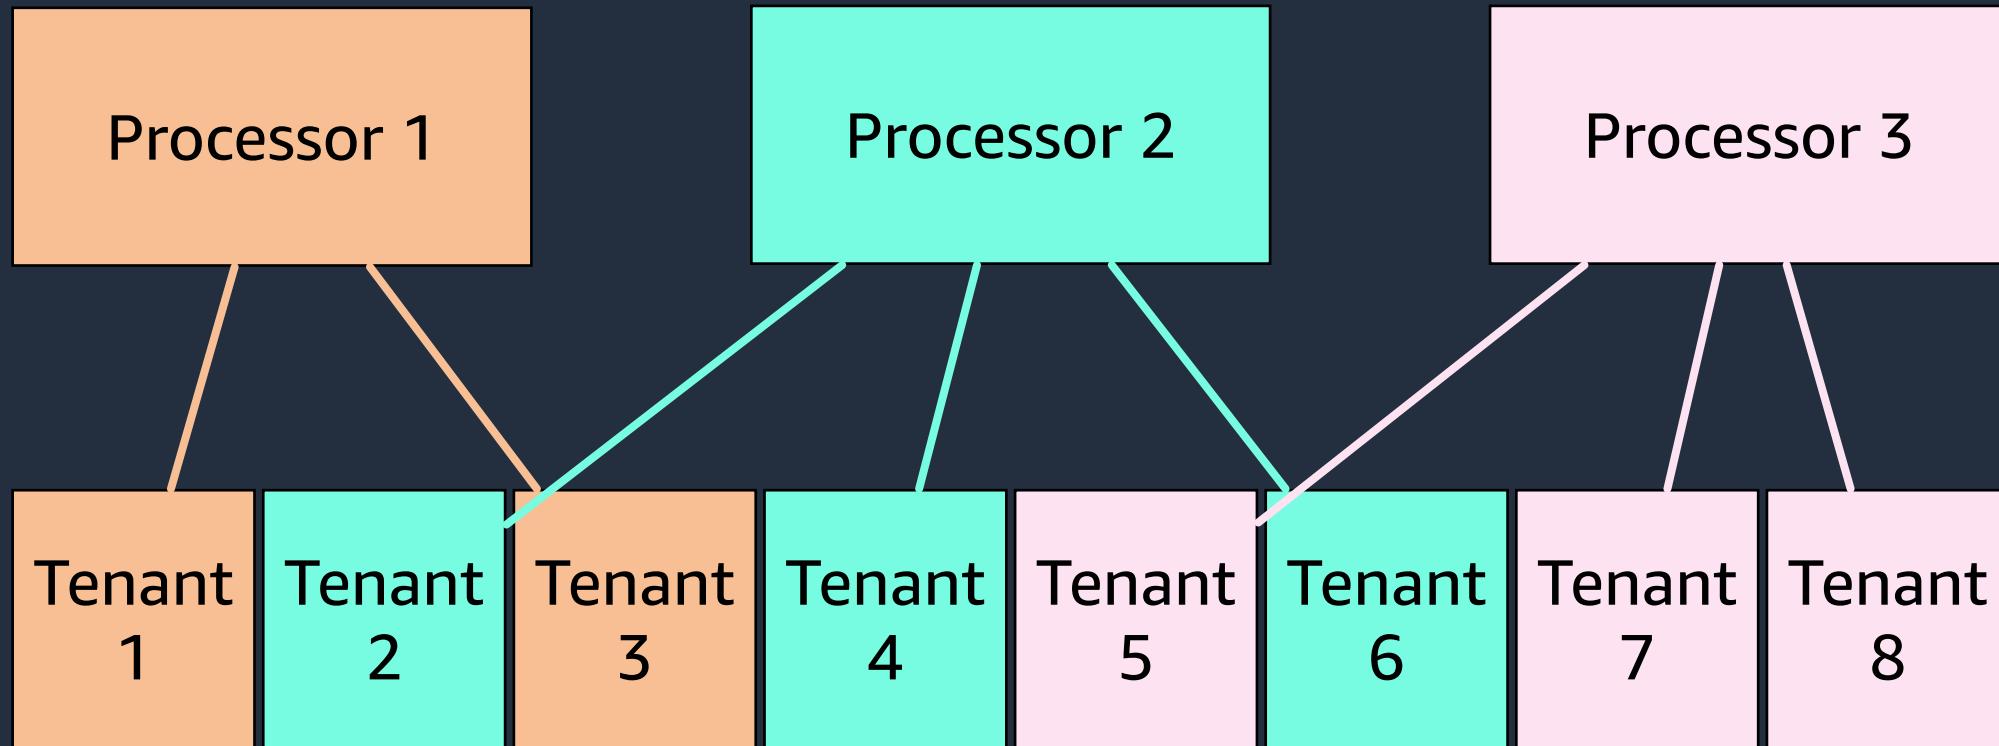
Let's analyze the problem



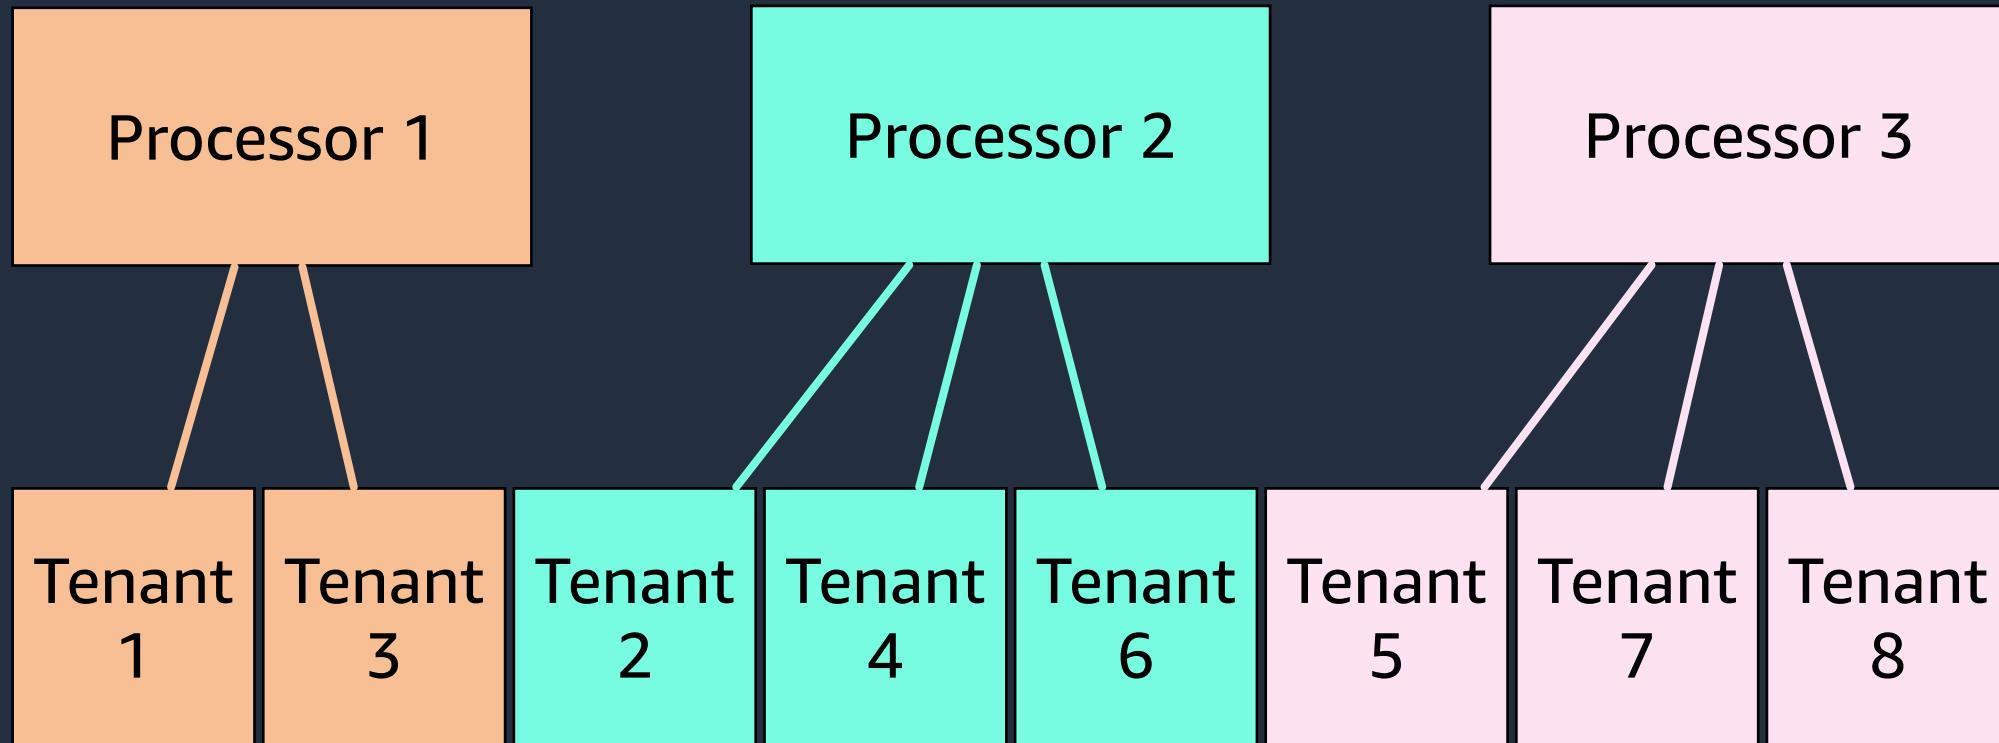
Let's analyze the problem



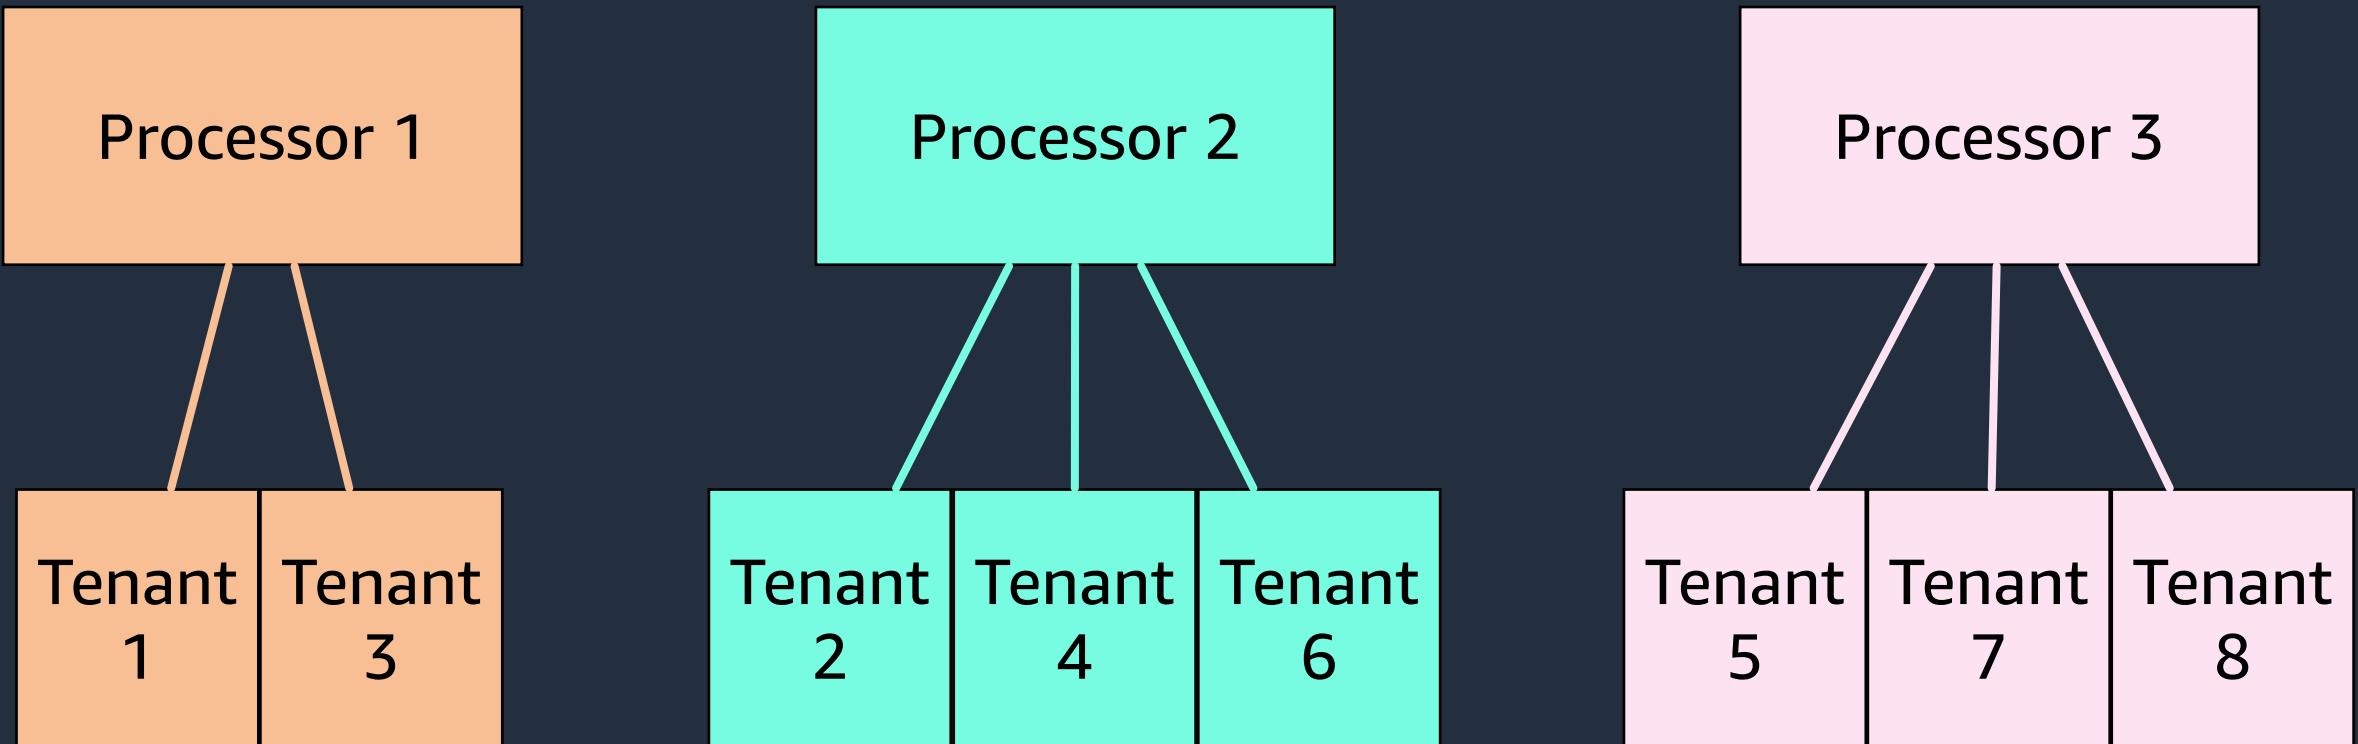
Let's analyze the problem



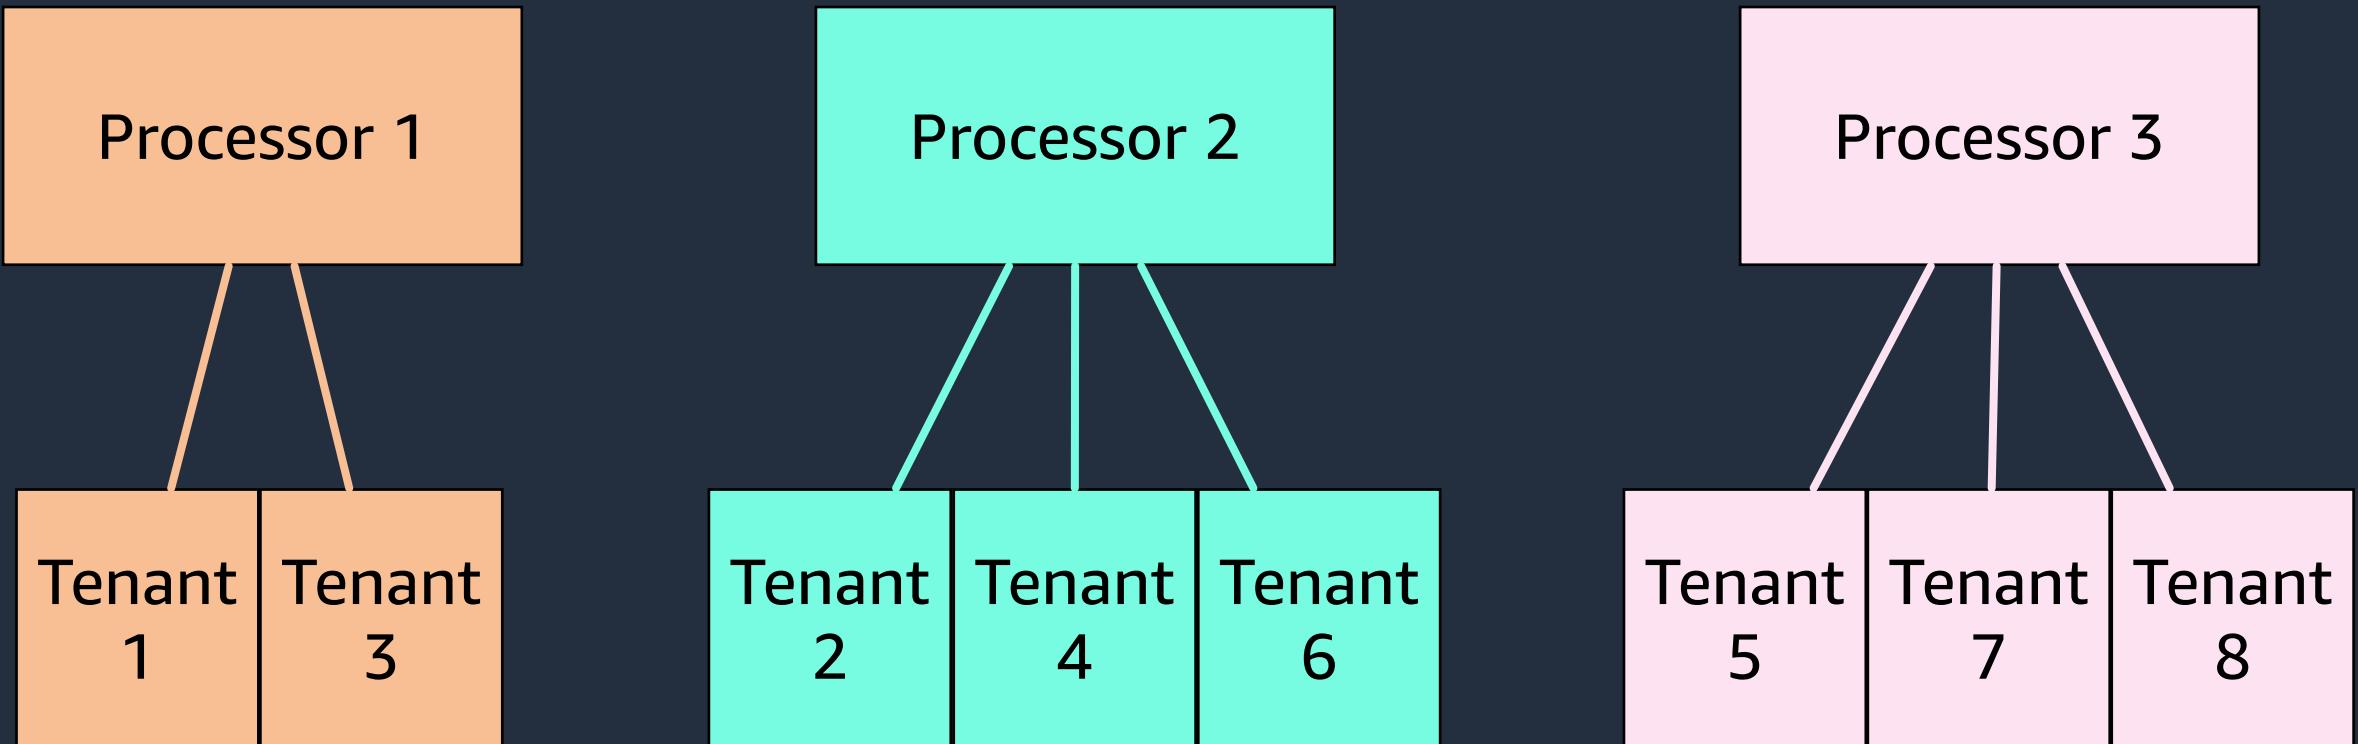
Let's analyze the problem



Let's analyze the problem

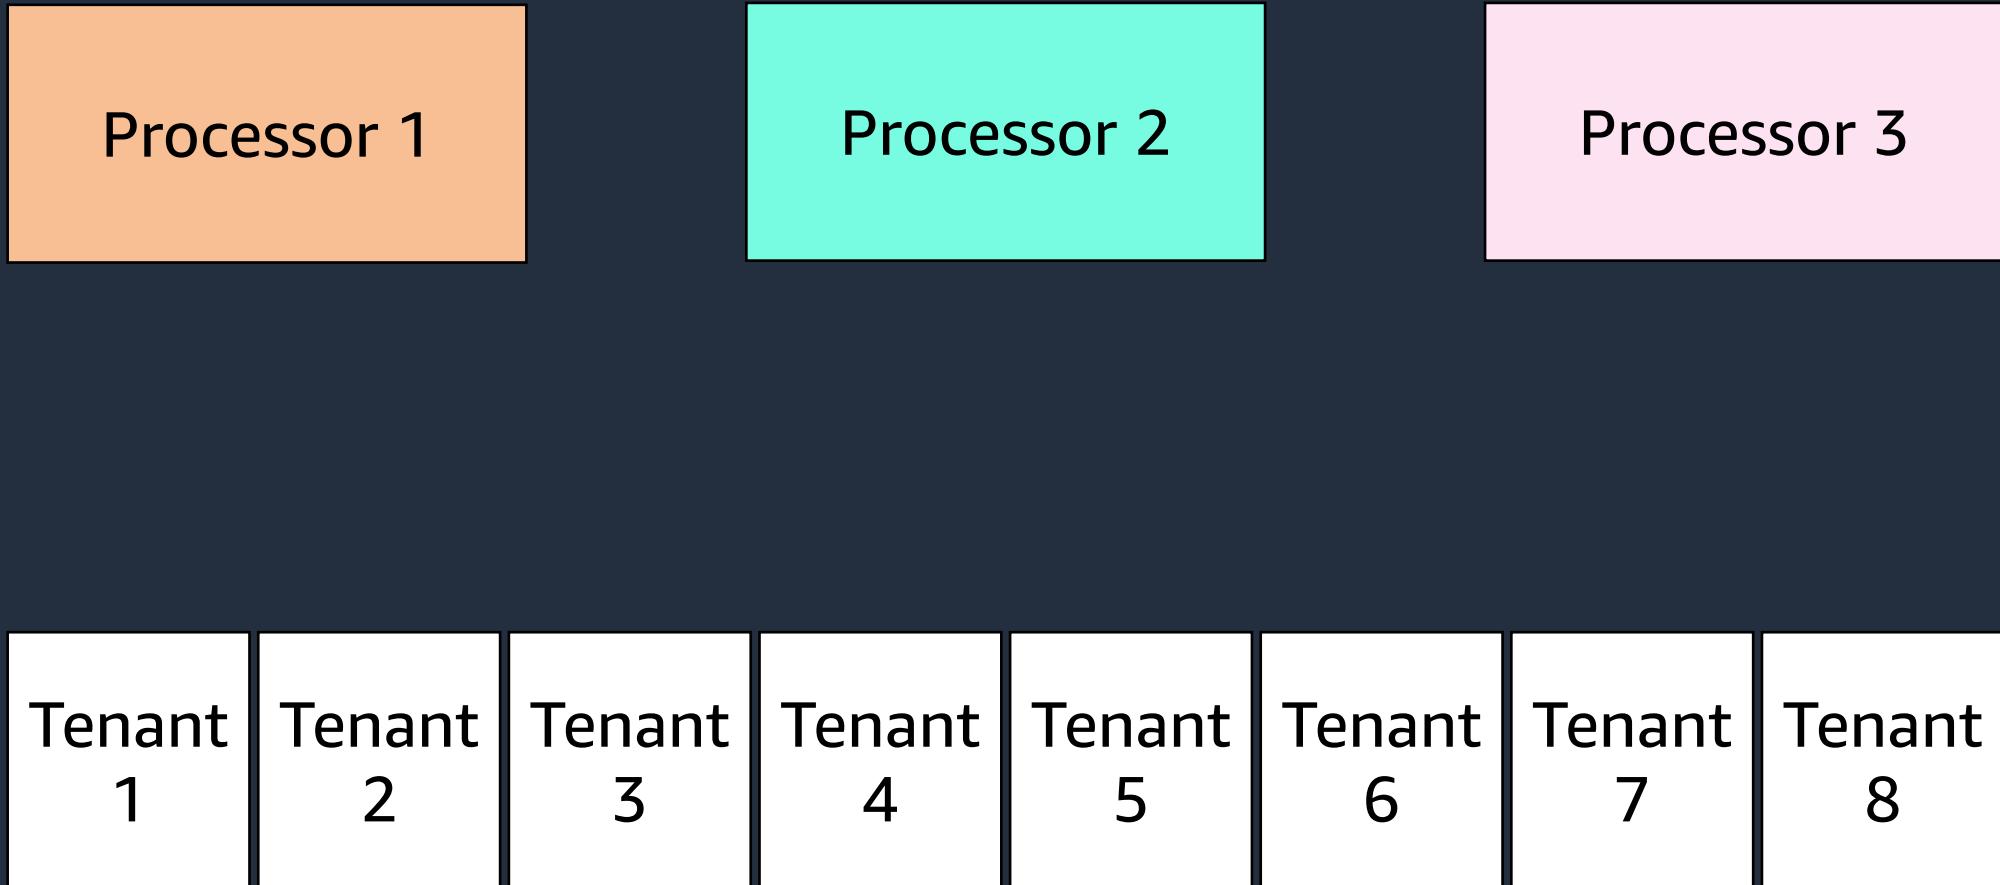


Let's analyze the problem

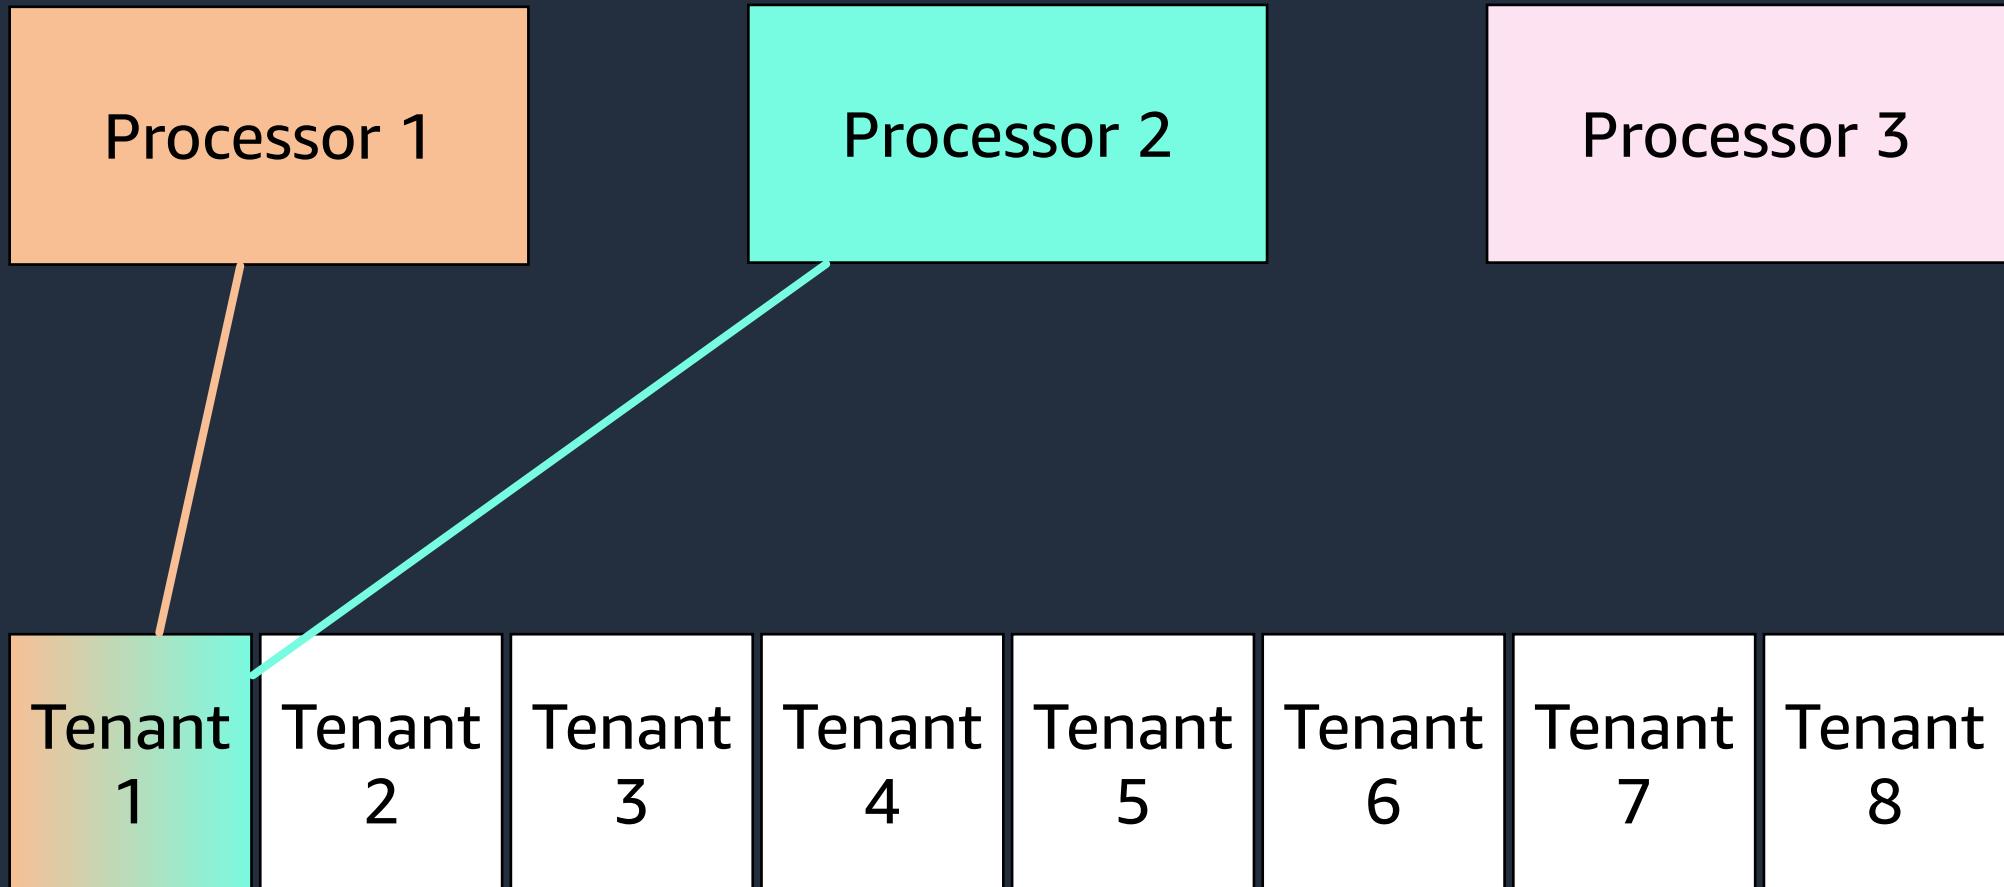


Same problem as before, but on a smaller scale

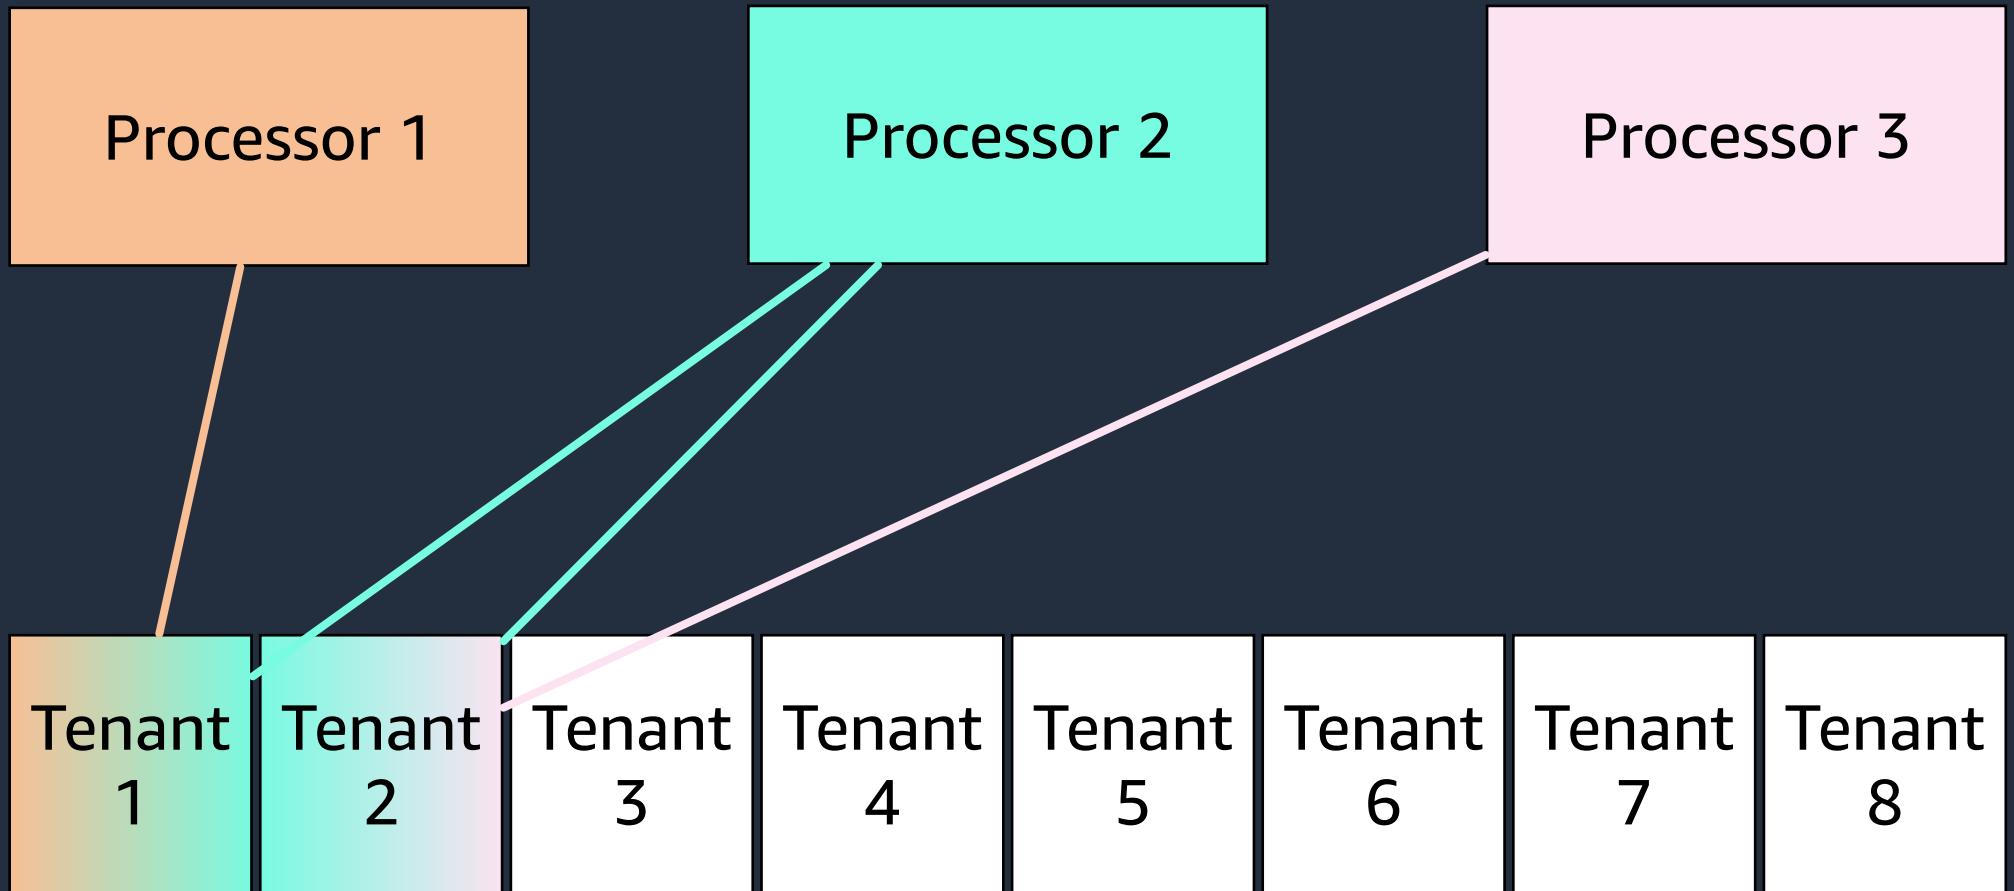
Shuffle sharding



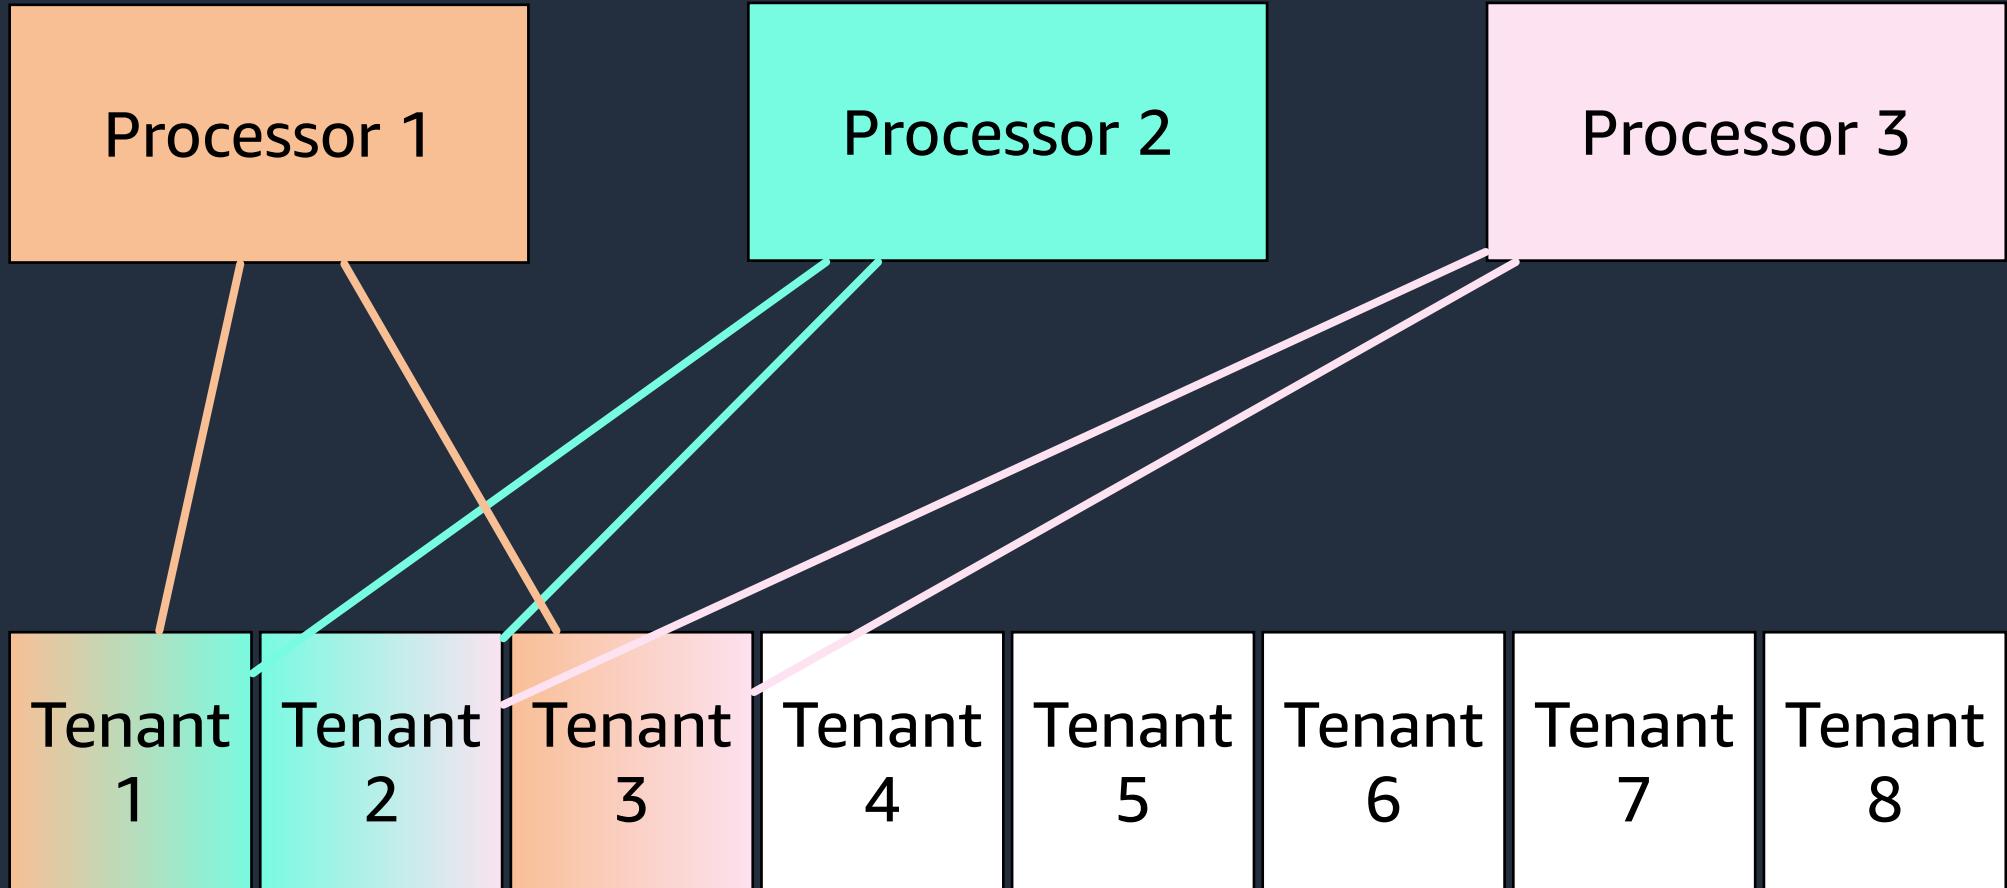
Shuffle sharding



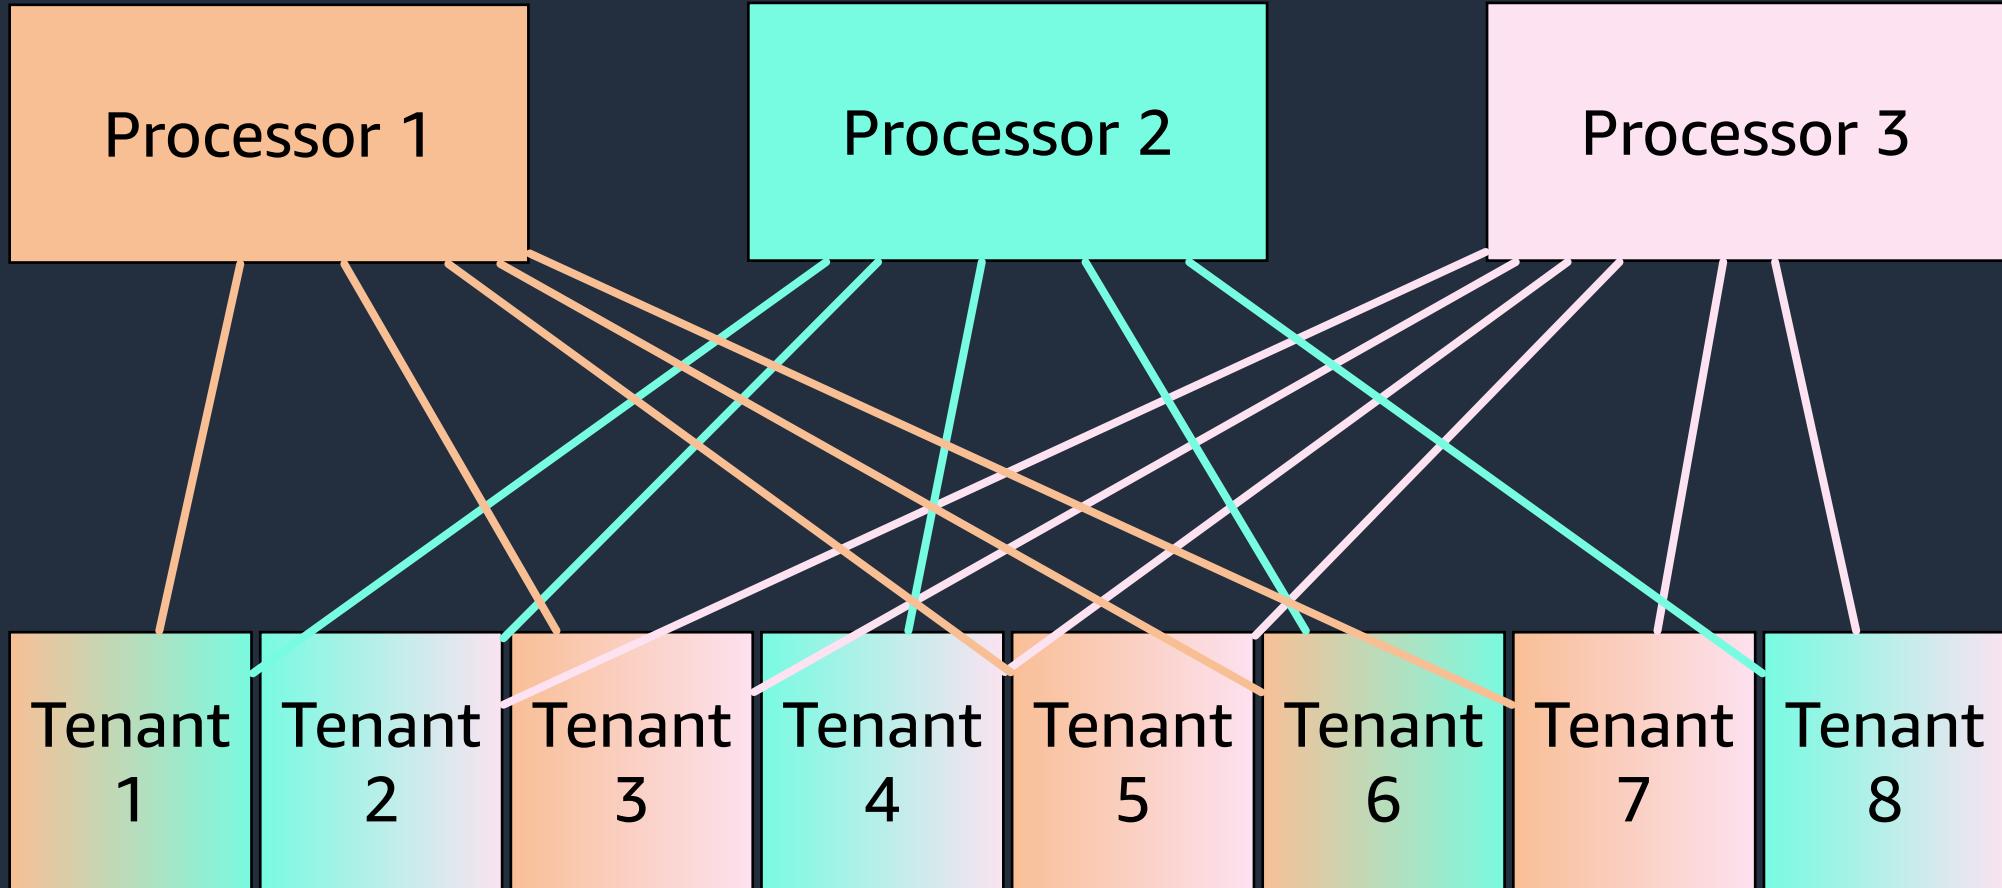
Shuffle sharding



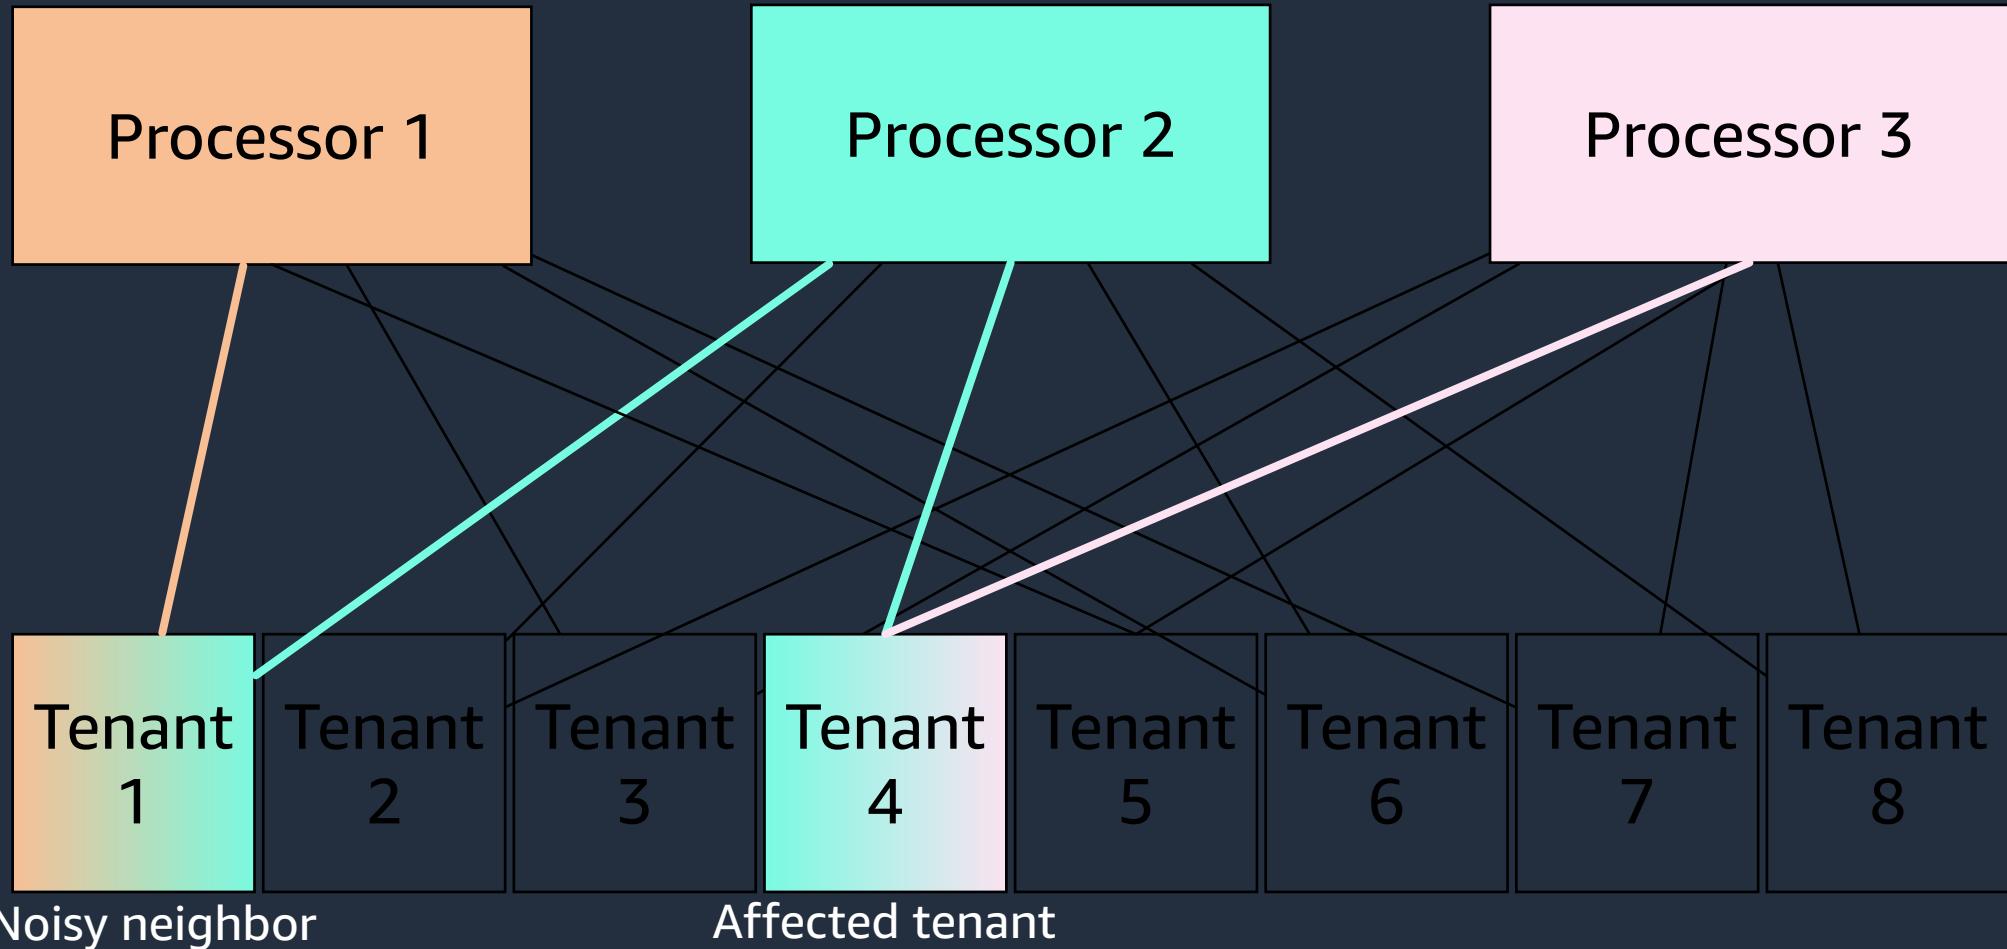
Shuffle sharding



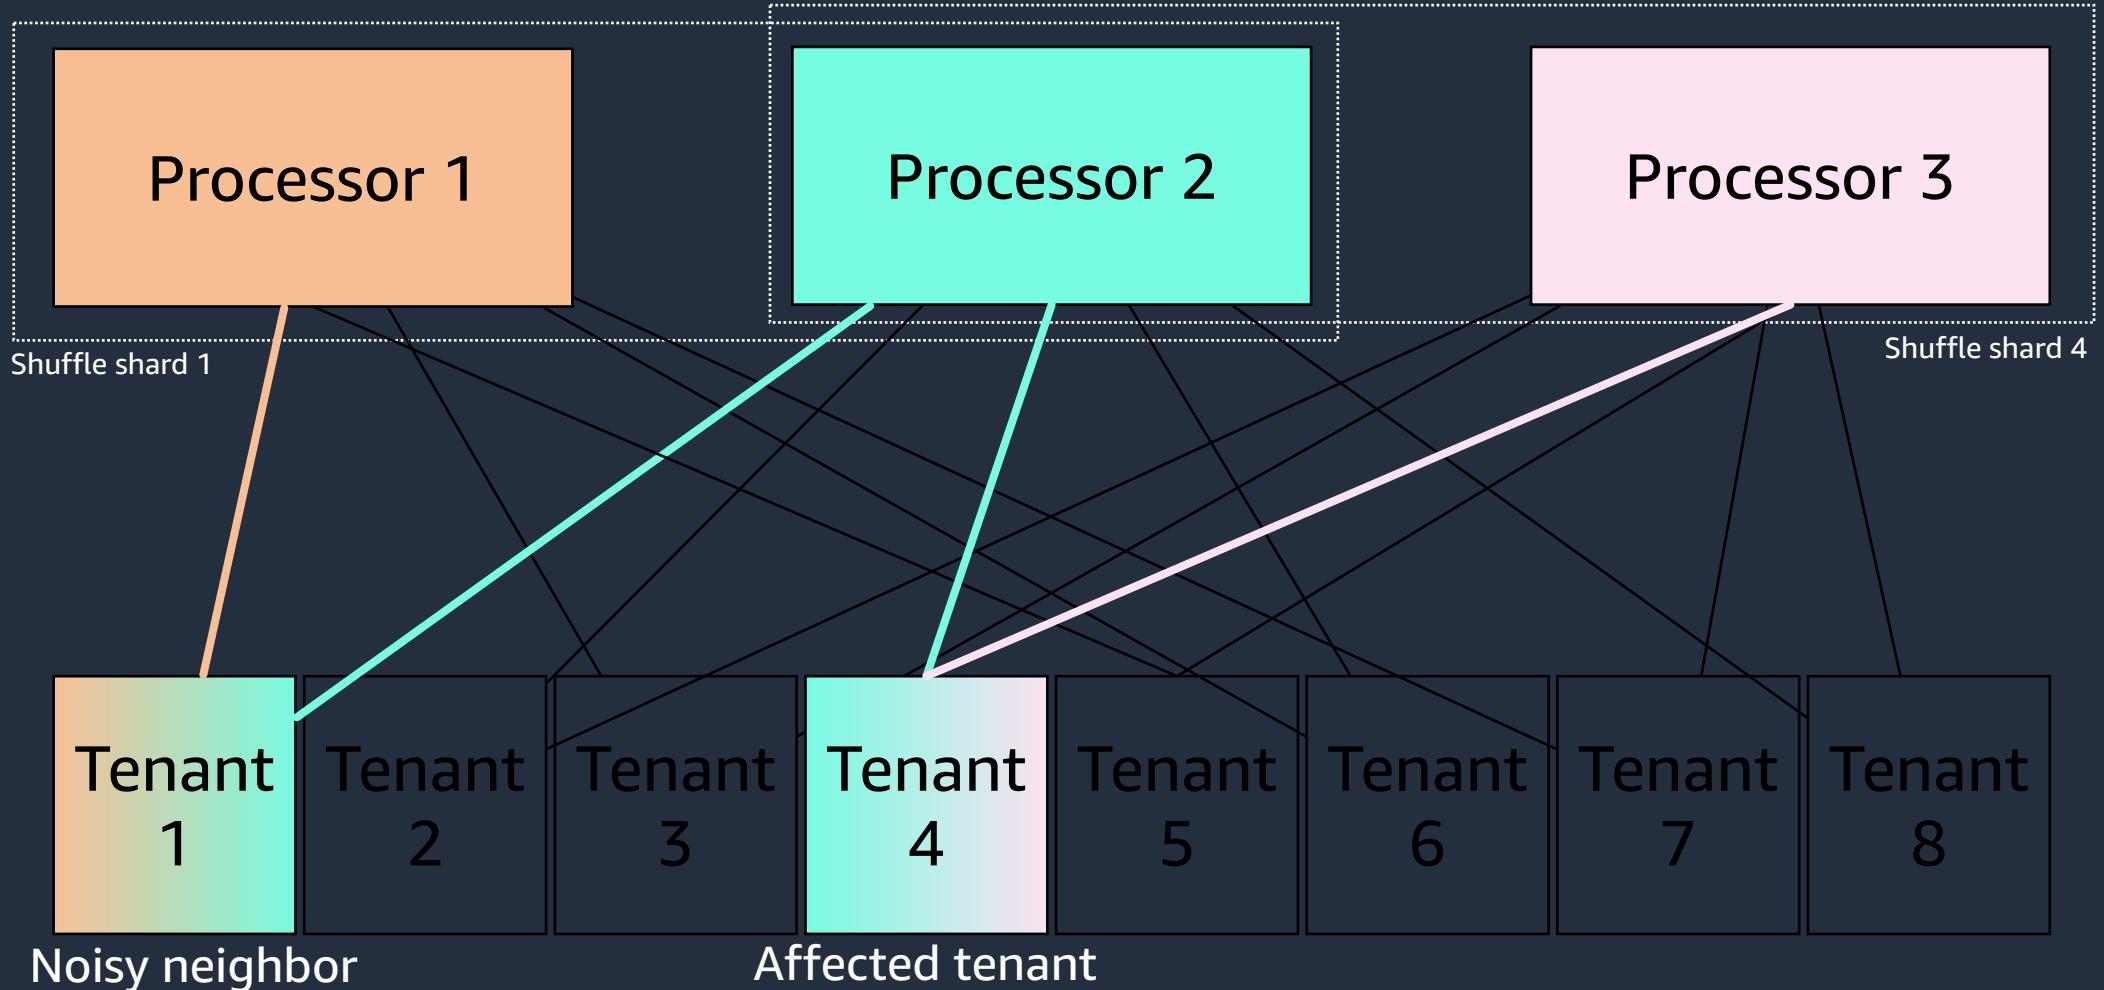
Shuffle sharding



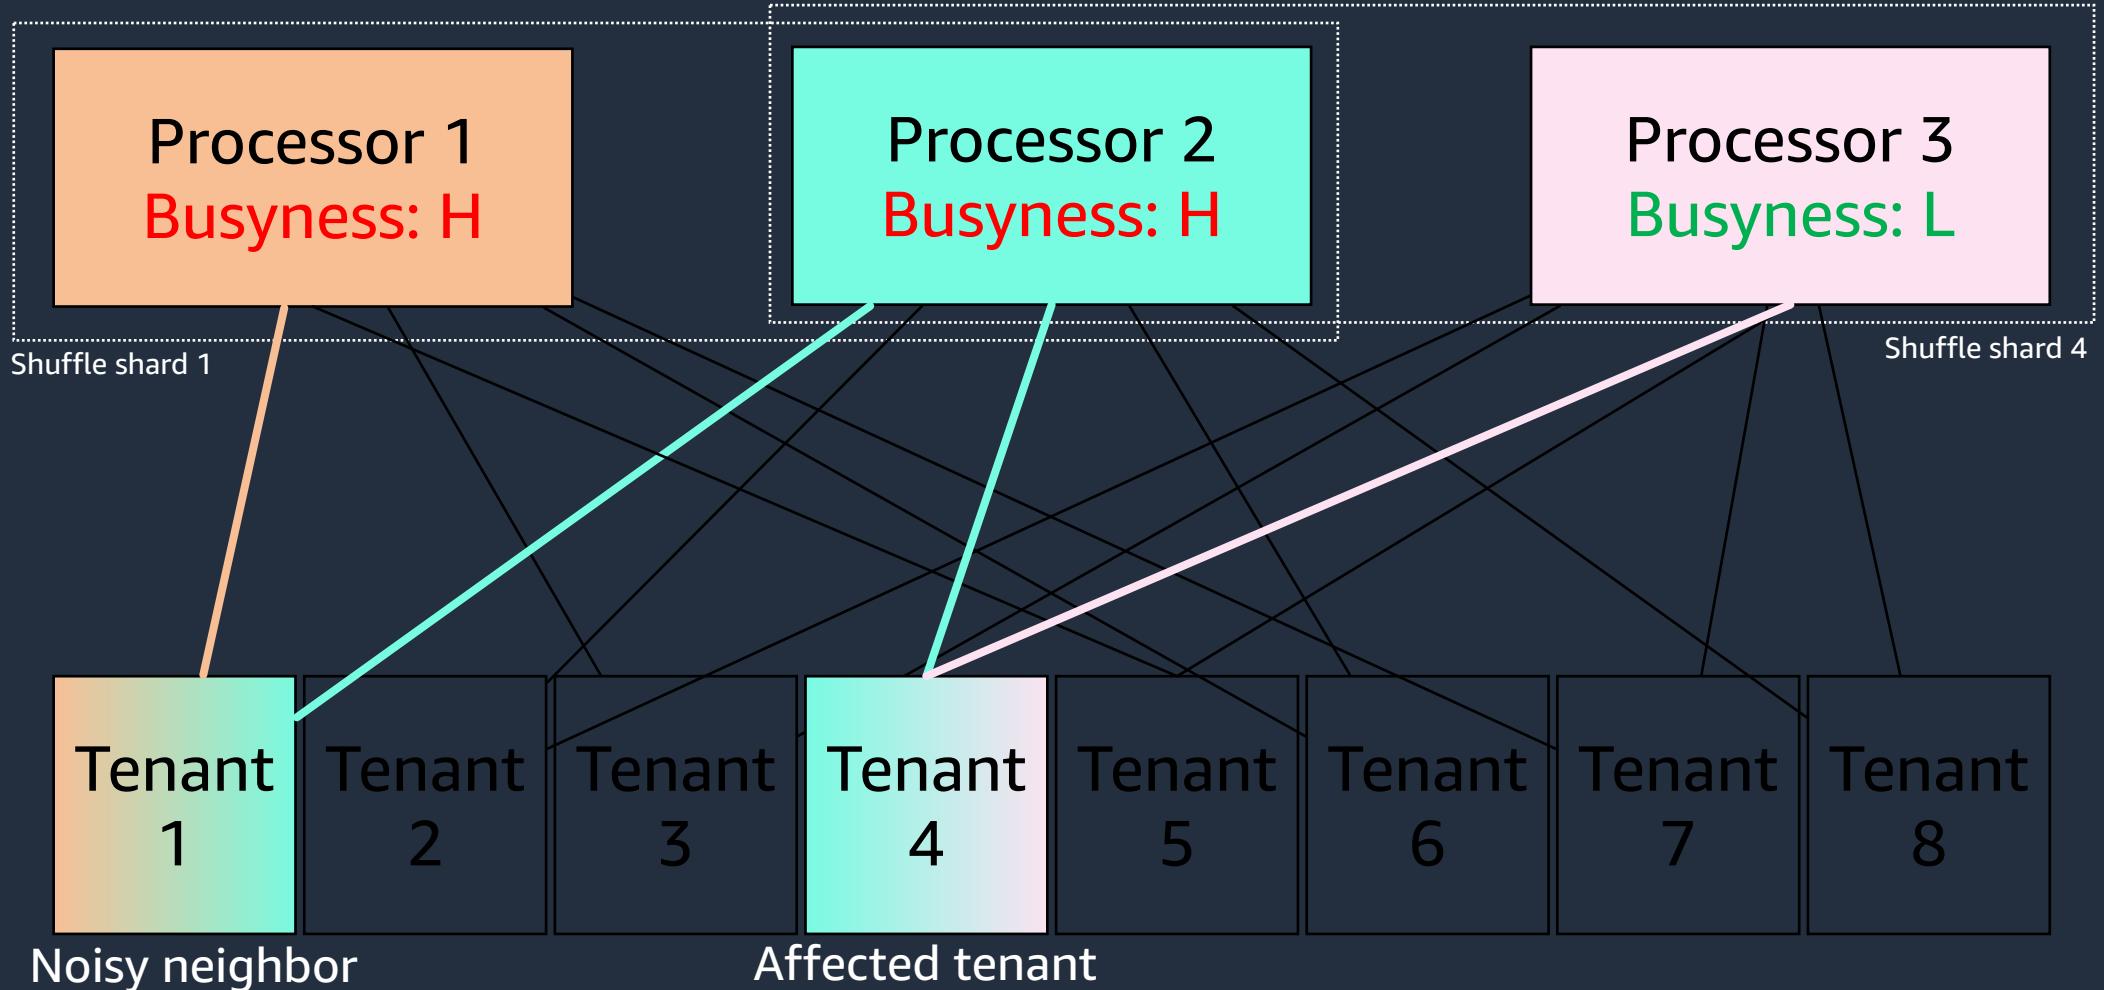
Shuffle sharding



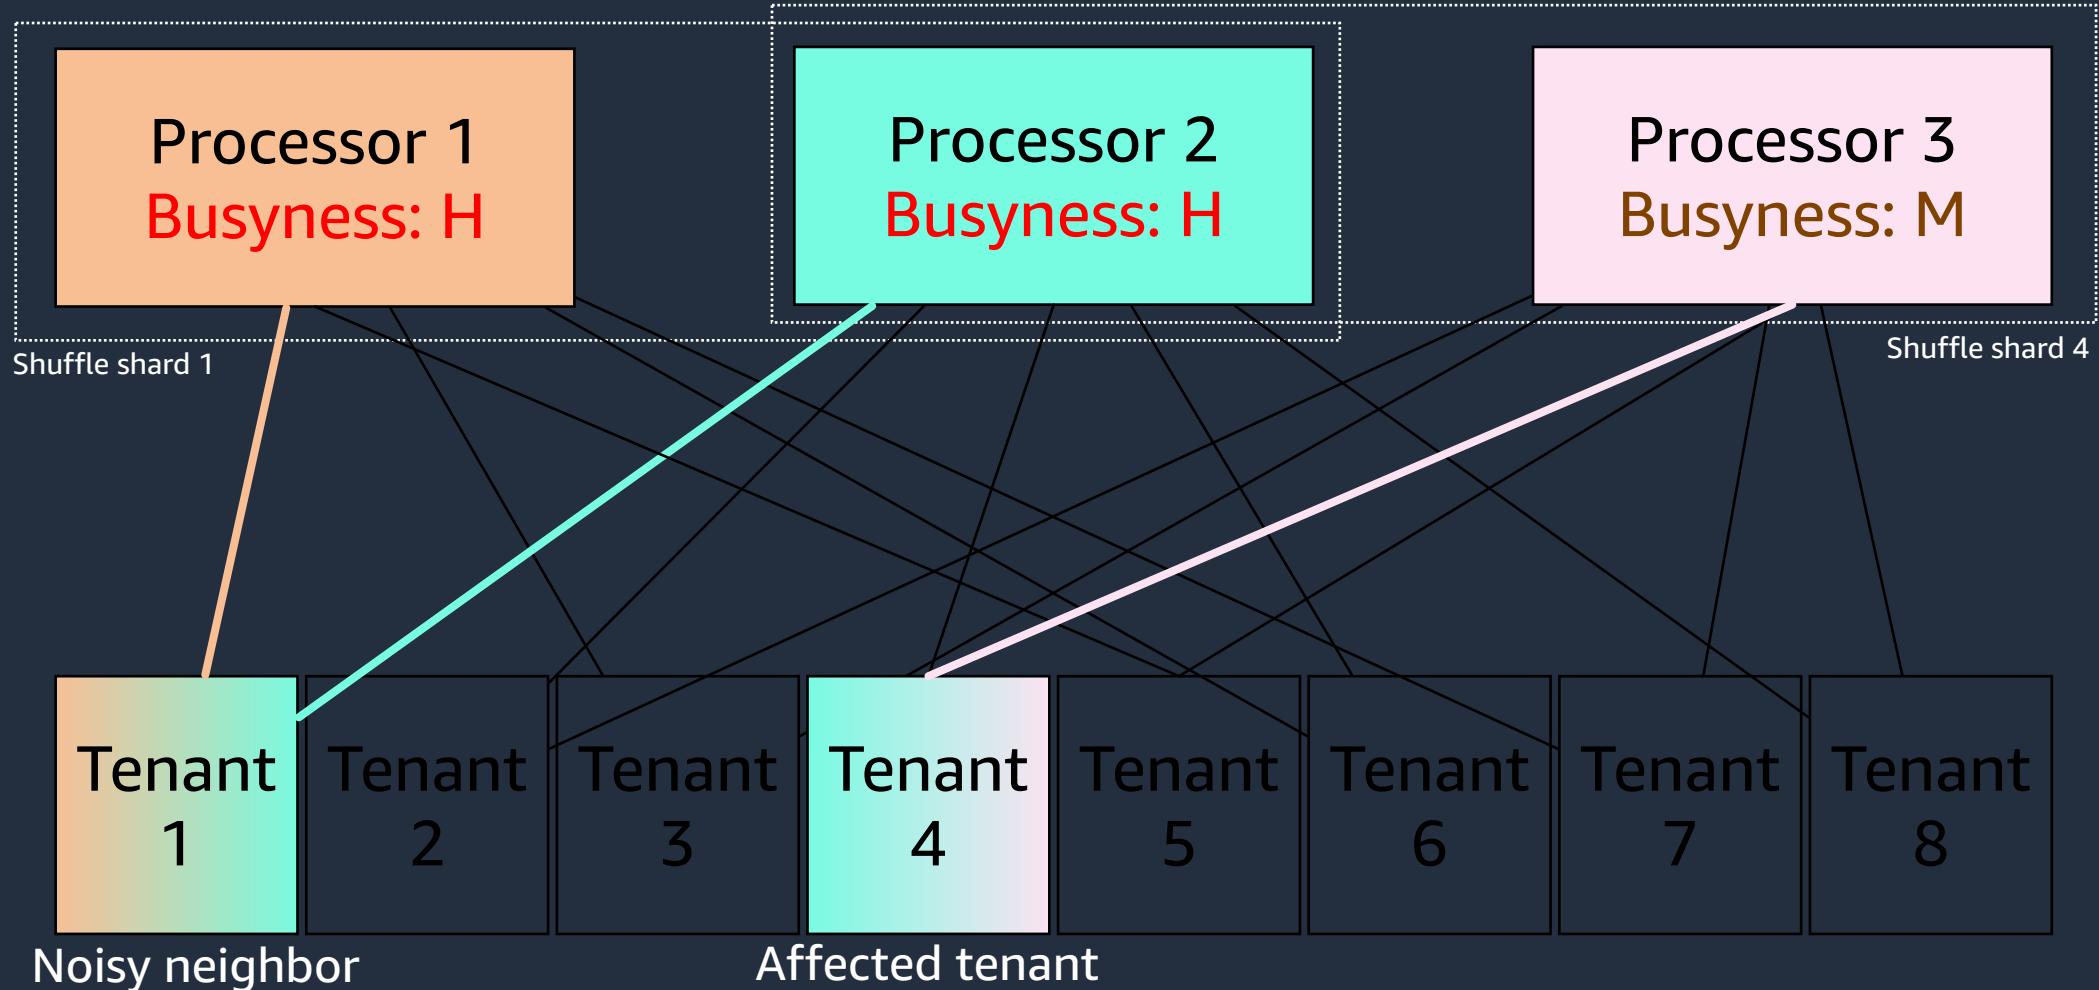
Shuffle sharding



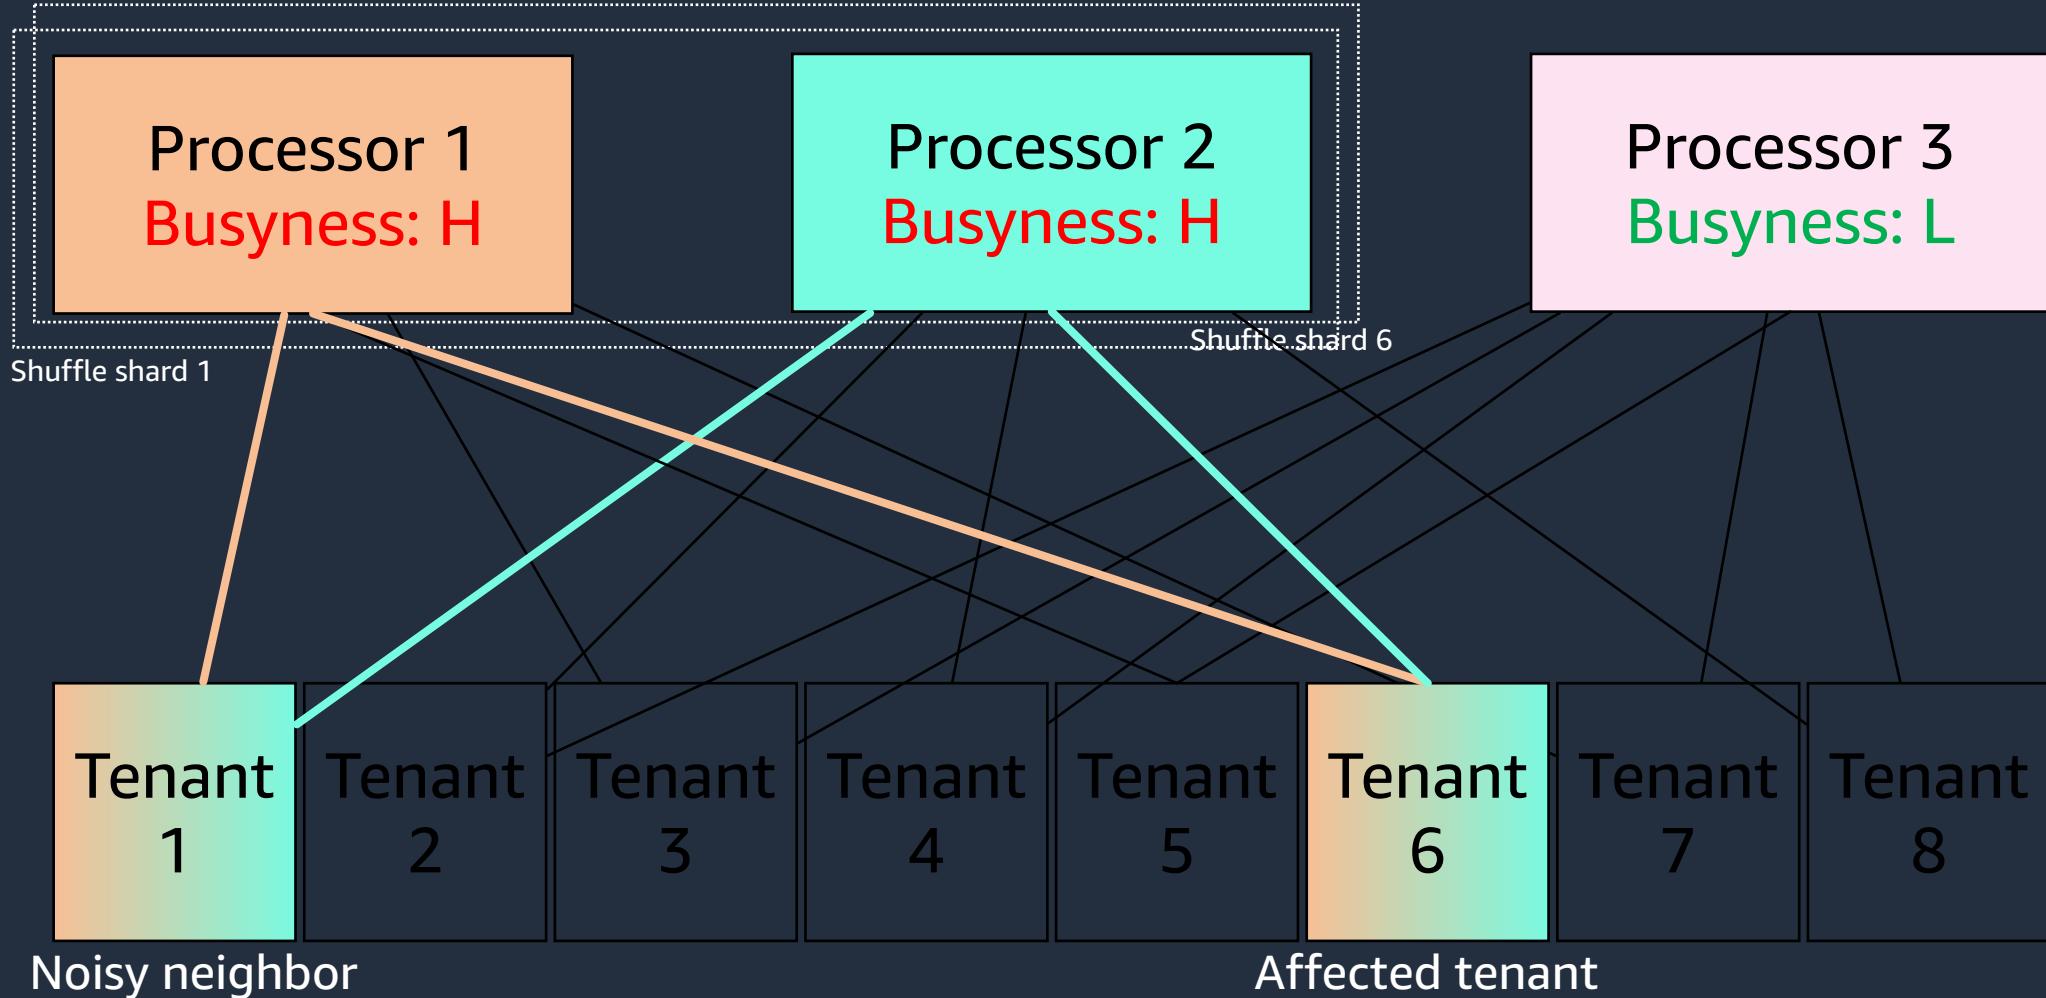
Shuffle sharding – adding a heuristic



Shuffle sharding – adding a heuristic



Shuffle sharding – tenants using the same processors



Probability of using the same processors set

p – number of processors

s – processors per tenant (shuffle shard size)

N = C(p, s) number of combinations

P₂ – probability of 2 tenants picking the same combination = (1/N)¹

P₃ – probability of 3 tenants picking the same combination = (1/N)²

P₄ – probability of 4 tenants picking the same combination = (1/N)³

(this is the most boring slide in the whole presentation)

Probability of using the same processors set

| Number of processors | 1 | 3 | 5 | 8 | 10 | 20 | 50 |
|------------------------------------|------|-----|------|---------|----------|-------------|-----------------|
| Processors per tenant (shard size) | 1 | 1 | 2 | 3 | 3 | 3 | 3 |
| Number of combinations | 1 | 3 | 10 | 56 | 120 | 1140 | 19600 |
| P (2 tenants using same shard) | 100% | 33% | 10% | 2% | 0.83% | 0.0877% | 0.00510% |
| P (3 tenants using same shard) | 100% | 11% | 1% | 0.0319% | 0.0069% | 0.0000769% | 0.000000260% |
| P (4 tenants using same shard) | 100% | 4% | 0.1% | 0.0006% | 0.00006% | 0.00000067% | 0.000000000133% |

Probability of using the same processors set

| Number of processors | 1 | 3 | 5 | 8 | 10 | 20 | 50 |
|------------------------------------|------|-----|------|---------|----------|-------------|-----------------|
| Processors per tenant (shard size) | 1 | 1 | 2 | 3 | 3 | 3 | 3 |
| Number of combinations | 1 | 3 | 10 | 56 | 120 | 1140 | 19600 |
| P (2 tenants using same shard) | 100% | 33% | 10% | 2% | 0.83% | 0.0877% | 0.00510% |
| P (3 tenants using same shard) | 100% | 11% | 1% | 0.0319% | 0.0069% | 0.0000769% | 0.000000260% |
| P (4 tenants using same shard) | 100% | 4% | 0.1% | 0.0006% | 0.00006% | 0.00000067% | 0.000000000133% |

Probability of using the same processors set

| Number of processors | 1 | 3 | 5 | 8 | 10 | 20 | 50 |
|------------------------------------|------|-----|------|---------|----------|-------------|-----------------|
| Processors per tenant (shard size) | 1 | 1 | 2 | 3 | 3 | 3 | 3 |
| Number of combinations | 1 | 3 | 10 | 56 | 120 | 1140 | 19600 |
| P (2 tenants using same shard) | 100% | 33% | 10% | 2% | 0.83% | 0.0877% | 0.00510% |
| P (3 tenants using same shard) | 100% | 11% | 1% | 0.0319% | 0.0069% | 0.0000769% | 0.000000260% |
| P (4 tenants using same shard) | 100% | 4% | 0.1% | 0.0006% | 0.00006% | 0.00000067% | 0.000000000133% |

Probability of using the same processors set

| Number of processors | 1 | 3 | 5 | 8 | 10 | 20 | 50 |
|------------------------------------|------|-----|------|---------|----------|-------------|-----------------|
| Processors per tenant (shard size) | 1 | 1 | 2 | 3 | 3 | 3 | 3 |
| Number of combinations | 1 | 3 | 10 | 56 | 120 | 1140 | 19600 |
| P (2 tenants using same shard) | 100% | 33% | 10% | 2% | 0.83% | 0.0877% | 0.00510% |
| P (3 tenants using same shard) | 100% | 11% | 1% | 0.0319% | 0.0069% | 0.0000769% | 0.000000260% |
| P (4 tenants using same shard) | 100% | 4% | 0.1% | 0.0006% | 0.00006% | 0.00000067% | 0.000000000133% |

Probability of using the same processors set

| Number of processors | 1 | 3 | 5 | 8 | 10 | 20 | 50 |
|------------------------------------|------|-----|------|---------|----------|-------------|-----------------|
| Processors per tenant (shard size) | 1 | 1 | 2 | 3 | 3 | 3 | 3 |
| Number of combinations | 1 | 3 | 10 | 56 | 120 | 1140 | 19600 |
| P (2 tenants using same shard) | 100% | 33% | 10% | 2% | 0.83% | 0.0877% | 0.00510% |
| P (3 tenants using same shard) | 100% | 11% | 1% | 0.0319% | 0.0069% | 0.0000769% | 0.000000260% |
| P (4 tenants using same shard) | 100% | 4% | 0.1% | 0.0006% | 0.00006% | 0.00000067% | 0.000000000133% |

Probability of using the same processors set

| | | | | | | | |
|------------------------------------|------|-----|--|---------|----------|-------------|-----------------|
| Number of processors | 1 | 3 | <u>Statistically</u> | | | | |
| Processors per tenant (shard size) | 1 | 1 | For ~20,000 tenants, a noisy neighbor should not have major impact on any other tenants | | | | |
| Number of combinations | 1 | 3 | For ~100,000 tenants, a noisy neighbor, should only have major impact on 4 other tenants | | | | |
| P (2 tenants using same shard) | 100% | 33% | 10% | 2% | 0.83% | 0.0877% | 0.00510% |
| P (3 tenants using same shard) | 100% | 11% | 1% | 0.0319% | 0.0069% | 0.0000769% | 0.000000260% |
| P (4 tenants using same shard) | 100% | 4% | 0.1% | 0.0006% | 0.00006% | 0.00000067% | 0.000000000133% |

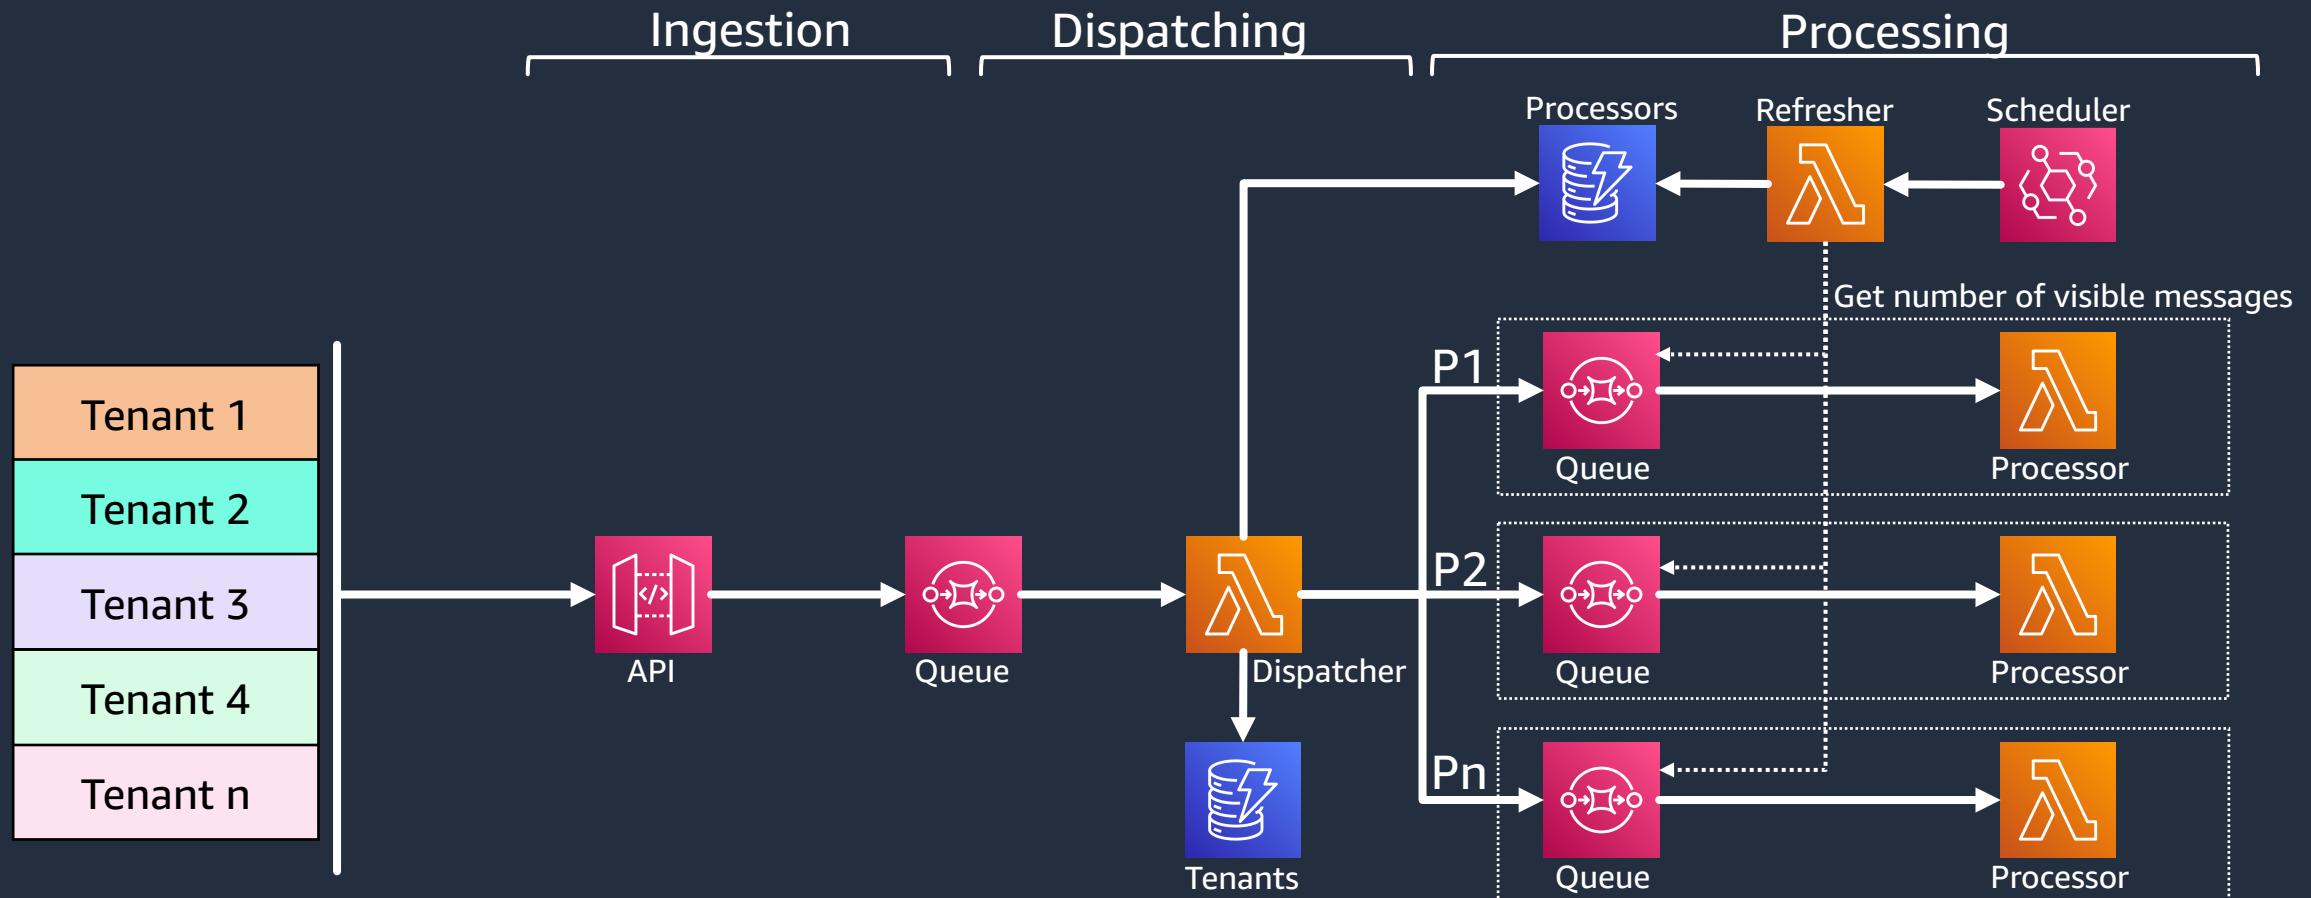
Shuffle sharding and Route 53

Amazon Route 53 and shuffle sharding

How does all of this help Amazon Route 53? With Route 53, we decided to arrange our capacity into a total of 2048 virtual name servers. These servers are virtual because they don't correspond to the physical servers hosting Route 53. We can move them around to help manage capacity. We then assign every customer domain to a shuffle shard of four virtual name servers. With those numbers, there are a staggering 730 billion possible shuffle shards. We have so many possible shuffle shards that we can assign a unique shuffle shard to every domain. Actually, we can go further, and ensure that no customer domain will ever share more than two virtual name servers with any other customer domain.

<https://aws.amazon.com/builders-library/workload-isolation-using-shuffle-sharding/>

Shuffle sharding



Let's see this in action

Noisy neighbor!

| Label | # Samples | Average | Median | 90% Line | 95% Line | 99% Line | Min | Maximum | Error % | Throughput | R |
|---------------------|-----------|---------|--------|----------|----------|----------|-----|---------|---------|------------|---|
| Request - Tenant 1 | 673 | 406 | 315 | 589 | 661 | 990 | 221 | 1282 | 0.00% | 2.0/sec | |
| Request - Tenant 10 | 653 | 402 | 311 | 580 | 665 | 864 | 225 | 1143 | 0.00% | 1.9/sec | |
| Request - Tenant 11 | 680 | 404 | 294 | 589 | 720 | 1084 | 228 | 1281 | 0.00% | 2.0/sec | |
| Request - Tenant 12 | 643 | 404 | 304 | 586 | 657 | 1032 | 232 | 1277 | 0.00% | 1.9/sec | |
| Request - Tenant 13 | 646 | 411 | 306 | 615 | 729 | 1043 | 230 | 1118 | 0.00% | 1.9/sec | |
| Request - Tenant 14 | 629 | 404 | 297 | 597 | 678 | 920 | 229 | 1213 | 0.00% | 1.9/sec | |
| Request - Tenant 15 | 664 | 393 | 291 | 578 | 654 | 865 | 222 | 1226 | 0.00% | 2.0/sec | |
| Request - Tenant 16 | 672 | 410 | 324 | 596 | 685 | 1011 | 227 | 1208 | 0.00% | 2.0/sec | |
| Request - Tenant 17 | 642 | 397 | 287 | 588 | 674 | 1051 | 228 | 1250 | 0.00% | 1.9/sec | |
| Request - Tenant 18 | 689 | 397 | 293 | 579 | 673 | 985 | 224 | 1141 | 0.00% | 2.0/sec | |
| Request - Tenant 19 | 655 | 405 | 293 | 590 | 711 | 1012 | 231 | 1150 | 0.00% | 1.9/sec | |
| Request - Tenant 20 | 9968 | 408 | 310 | 590 | 691 | 1060 | 220 | 1579 | 0.00% | 29.3/sec | |
| TOTAL | 22690 | 405 | 304 | 590 | 690 | 1039 | 220 | 2183 | 0.00% | 66.7/sec | |

Before we begin

Processor Queues - SQS

| | processorId | approximateNumberOfMessagesVisible | queueUrl |
|--|-------------|------------------------------------|---|
| | 1 | 0 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Shuffle-P1DataProcessorQueue1D1C95DF-LKCok95fqW5V |
| | 2 | 0 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Shuffle-P2DataProcessorQueue880A1E3E-Vial6DMprNSp |
| | 3 | 0 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Shuffle-P3DataProcessorQueue2DACC5EB-6WE276nFSVqd |
| | 4 | 0 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Shuffle-P4DataProcessorQueue8B458582-roCPPSu23Z2Q |
| | 5 | 0 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Shuffle-P5DataProcessorQueue799FB189-vHvHj55E5zHO |
| | 6 | 0 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Shuffle-P6DataProcessorQueueC6D8E8E5-AjcPjl4ZHtQk |
| | 7 | 0 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Shuffle-P7DataProcessorQueueF298AEC3-e9tJJ8zo93sA |
| | 8 | 0 | https://SQS.us-east-1.amazonaws.com/281024298475/SaaS-Sharding-Shuffle-P8DataProcessorQueue1F68FCB8-Xd7m6712kW1d |

Before we begin

Refresher trigger - EventBridge

The image displays two side-by-side screenshots of the AWS EventBridge Rule configuration interface.

Screenshot 1: Rule details for SaasSharding-Shuffle-RefresherTriggerRule

- Rule details:** Shows the rule name "SaasSharding-Shuffle-RefresherTriggerRuleA577C3B5-TEPZRRGV7XML", status "Enabled", and a description field.
- Event schedule:** Set to "Fixed rate of 1 minute".
- Targets:** Not visible in this screenshot.

Screenshot 2: Rule details for SaasSharding-Shuffle-RefresherTriggerRuleA577C3B5-TEPZRRGV7XML

- Rule details:** Shows the rule name "SaasSharding-Shuffle-RefresherTriggerRuleA577C3B5-TEPZRRGV7XML", status "Enabled", and a description field.
- Event schedule:** Not visible in this screenshot.
- Targets:** Visible and highlighted with an orange border. It shows one target named "SaasSharding-Shuffle-Refresher59C38116-kCVNSYjc9V9g" of type "Lambda function".
 - Input to target:** "Matched event"
 - Additional parameters:** "-"
 - Dead-letter queue (DLQ):** "-"

Before we begin

Refresher function - Lambda

```
function findBestProcessor(tenantId){  
    console.log('> findBestProcessor for tenant', tenantId);  
    const tenantProcessorIds = tenants[tenantId];  
  
    // Get processors for this tenantId  
    let tenantProcessors = tenantProcessorIds.map(id=>processors[id]);  
  
    // Sort processors by number of messages visible  
    tenantProcessors.sort((a,b)=>a.approximateNumberOfMessagesVisible - b.approximateNumberOfMessagesVisible);  
  
    // Get the smallest number of messages visible number  
    const min = tenantProcessors[0].approximateNumberOfMessagesVisible;  
  
    // Filter out all processors with number of messages !== min  
    const filteredProcessors = tenantProcessors.filter(i=>i.approximateNumberOfMessagesVisible==min);  
  
    // Pick random processor from the remaining list  
    const randomIdx = Math.floor(Math.random() * filteredProcessors.length);  
    const bestProcessor = filteredProcessors[randomIdx];  
  
    return bestProcessor;  
}
```

Before we begin

| tenantid | processordids |
|----------|---------------|
| 1 | [2,1,6] |
| 10 | [2,4,7] |
| 11 | [3,2,1] |
| 12 | [1,5,8] |
| 13 | [1,4,5] |
| 14 | [7,8,2] |
| 15 | [1,8,6] |
| 16 | [2,6,7] |
| 17 | [4,5,2] |
| 18 | [1,2,5] |
| 19 | [4,5,3] |
| 2 | [4,8,3] |
| 20 | [7,8,6] |
| 3 | [6,8,5] |
| 4 | [1,3,6] |
| 5 | [3,2,8] |
| 6 | [2,7,4] |
| 7 | [5,1,4] |
| 8 | [3,5,6] |
| 9 | [4,3,7] |

| | T1 | T2 | T3 | T4 | T5 | T6 | T7 | T8 | T9 | T10 | T11 | T12 | T13 | T14 | T15 | T16 | T17 | T18 | T19 | T20 |
|----|----|----|----|----|----|----|----|----|----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| P1 | X | | | X | | | X | | | | X | X | X | | X | | | X | | |
| P2 | X | | | | X | X | | | | X | X | | | X | | X | X | X | | |
| P3 | | X | | X | X | | | X | X | | X | | | | | | | X | | |
| P4 | | X | | | | X | X | | X | X | | | | X | | | X | X | | |
| P5 | | | X | | | | X | X | | | | X | X | | | | X | X | | |
| P6 | X | | | X | X | | | X | | | | | | | X | X | | | X | |
| P7 | | | | | | X | | | X | X | | | | X | | X | X | | X | |
| P8 | | | X | X | | X | | | | | | X | | X | X | | | | X | |

Before we begin

| tenantid | processordids |
|----------|---------------|
| 1 | [2,1,6] |
| 10 | [2,4,7] |
| 11 | [3,2,1] |
| 12 | [1,5,8] |
| 13 | [1,4,5] |
| 14 | [7,8,2] |
| 15 | [1,8,6] |
| 16 | [2,6,7] |
| 17 | [4,5,2] |
| 18 | [1,2,5] |
| 19 | [4,5,3] |
| 2 | [4,8,3] |
| 20 | [7,8,6] |
| 3 | [6,8,5] |
| 4 | [1,3,6] |
| 5 | [3,2,8] |
| 6 | [2,7,4] |
| 7 | [5,1,4] |
| 8 | [3,5,6] |
| 9 | [4,3,7] |

| | T1 | T2 | T3 | T4 | T5 | T6 | T7 | T8 | T9 | T10 | T11 | T12 | T13 | T14 | T15 | T16 | T17 | T18 | T19 | T20 |
|----|----|----|----|----|----|----|----|----|----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| P1 | X | | | X | | | X | | | | X | X | X | | X | | | X | | |
| P2 | X | | | | X | X | | | | X | X | | | X | | X | X | X | | |
| P3 | | X | | X | X | | | X | X | | X | | | | | | | | X | |
| P4 | | X | | | | X | X | | X | X | | | | X | | | X | | X | |
| P5 | | | X | | | | X | X | | | | | X | X | | | X | X | X | |
| P6 | X | | | X | X | | | X | | | | | | | | X | X | | | X |
| P7 | | | | | | X | | | X | X | | | | | X | | X | | | X |
| P8 | | | X | X | | X | | | | | | | X | | X | X | | | | X |

Before we begin

| tenantid | processordids |
|----------|---------------|
| 1 | [2,1,6] |
| 10 | [2,4,7] |
| 11 | [3,2,1] |
| 12 | [1,5,8] |
| 13 | [1,4,5] |
| 14 | [7,8,2] |
| 15 | [1,8,6] |
| 16 | [2,6,7] |
| 17 | [4,5,2] |
| 18 | [1,2,5] |
| 19 | [4,5,3] |
| 2 | [4,8,3] |
| 20 | [7,8,6] |
| 3 | [6,8,5] |
| 4 | [1,3,6] |
| 5 | [3,2,8] |
| 6 | [2,7,4] |
| 7 | [5,1,4] |
| 8 | [3,5,6] |
| 9 | [4,3,7] |

| | T1 | T2 | T3 | T4 | T5 | T6 | T7 | T8 | T9 | T10 | T11 | T12 | T13 | T14 | T15 | T16 | T17 | T18 | T19 | T20 |
|----|----|----|----|----|----|----|----|----|----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| P1 | X | | | X | | | X | | | | X | X | X | | X | | | X | | |
| P2 | X | | | | X | X | | | | X | X | | | X | | X | X | X | | |
| P3 | | X | | X | X | | | X | X | | X | | | | | | | | X | |
| P4 | | X | | | | X | X | | X | X | | | | X | | | X | X | | |
| P5 | | | X | | | | X | X | | | | X | X | | | | X | X | X | |
| P6 | X | | X | X | | | | X | | | | | | | X | X | | | X | |
| P7 | | | | | | X | | | X | X | | | | X | | X | X | | X | |
| P8 | | | X | X | | X | | | | | | X | | X | X | | | | X | |

- 6 tenants not affected
- 13 tenants affected

Before we begin

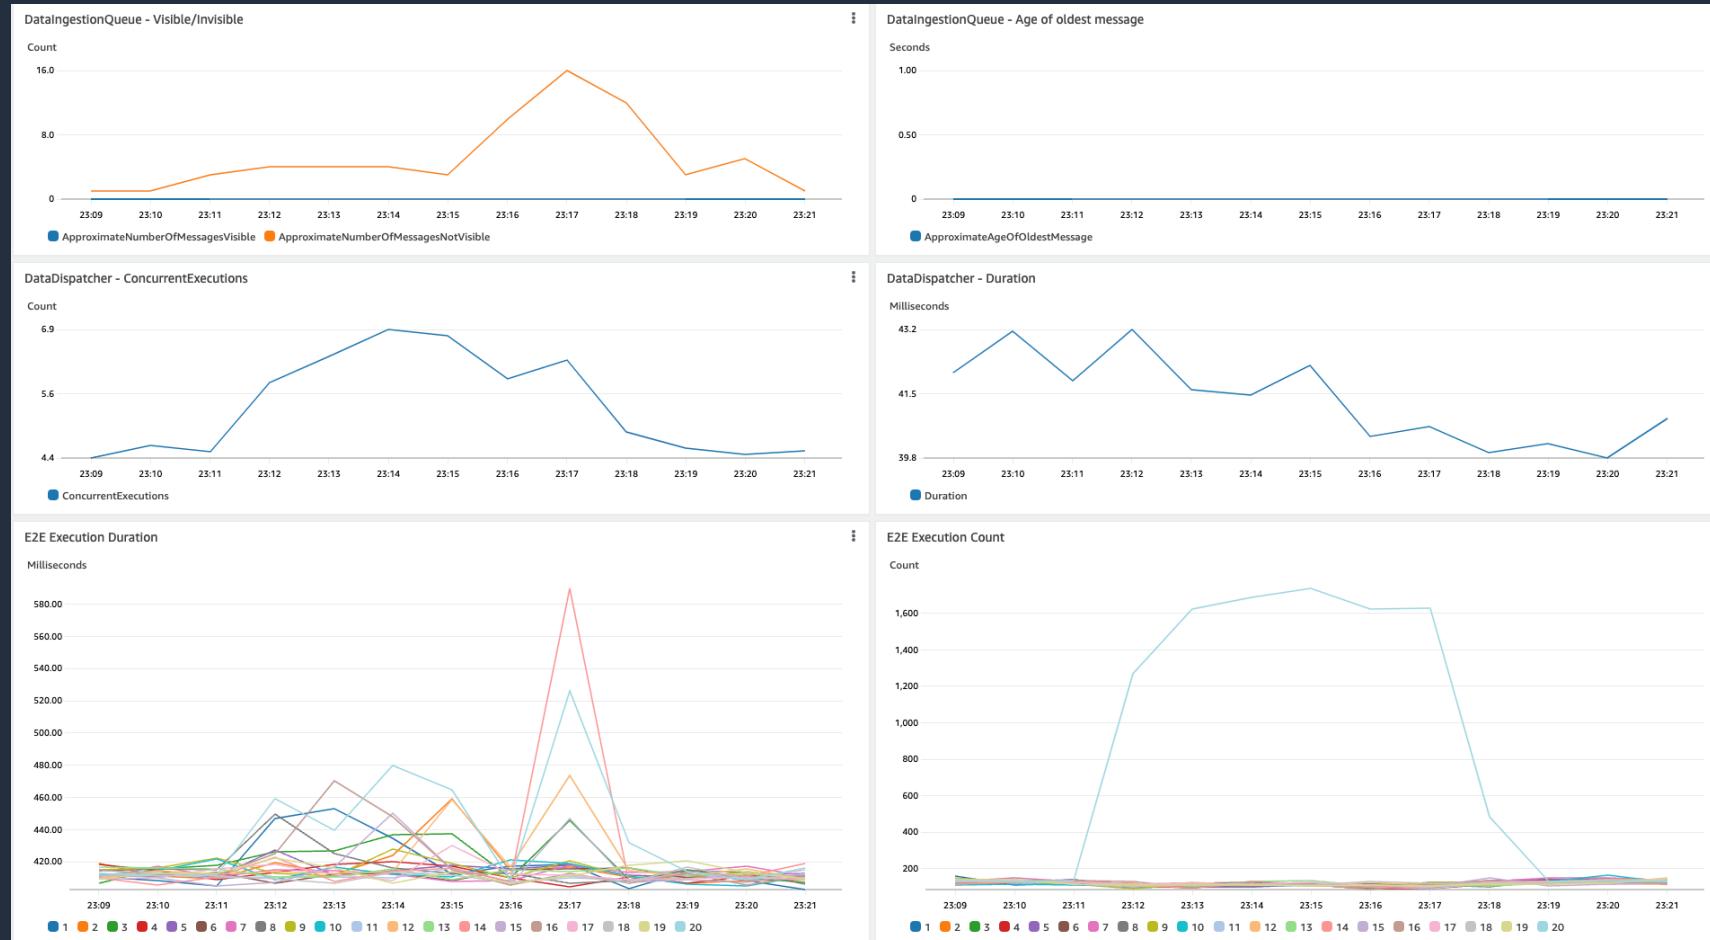
| tenantid | processordids |
|----------|---------------|
| 1 | [2,1,6] |
| 10 | [2,4,7] |
| 11 | [3,2,1] |
| 12 | [1,5,8] |
| 13 | [1,4,5] |
| 14 | [7,8,2] |
| 15 | [1,8,6] |
| 16 | [2,6,7] |
| 17 | [4,5,2] |
| 18 | [1,2,5] |
| 19 | [4,5,3] |
| 2 | [4,8,3] |
| 20 | [7,8,6] |
| 3 | [6,8,5] |
| 4 | [1,3,6] |
| 5 | [3,2,8] |
| 6 | [2,7,4] |
| 7 | [5,1,4] |
| 8 | [3,5,6] |
| 9 | [4,3,7] |

| | T1 | T2 | T3 | T4 | T5 | T6 | T7 | T8 | T9 | T10 | T11 | T12 | T13 | T14 | T15 | T16 | T17 | T18 | T19 | T20 |
|----|----|----|----|----|----|----|----|----|----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| P1 | X | | | X | | | X | | | X | X | X | X | X | | | X | | | |
| P2 | X | | | | X | X | | | | X | X | | | X | X | X | X | X | | |
| P3 | | X | | X | X | | | X | X | | X | | | | | | | X | | |
| P4 | | X | | | | X | X | | X | X | | | | | | | X | X | | |
| P5 | | | X | | | | X | X | | | | X | X | | | | X | X | X | |
| P6 | X | | X | X | | | | X | | | | | | | X | X | | | X | |
| P7 | | | | | | X | | | X | X | | | | X | X | X | | | X | |
| P8 | | X | X | | X | | | | | | X | | | X | X | | | | X | |

- 6 tenants not affected
- 13 tenants affected

Mental note
 Noisy neighbor – 20
 Affected processors – P6, P7, P8
 Affected tenants – T1, T2, T3, T4, T5, T6, T8, T9, T10, T11, T14, T15, T16

Results



Mental note
 Noisy neighbor – 20
 Affected processors – P6, P7, P8
 Affected tenants – T1, T2, T3, T4, T5, T6, T8, T9, T10, T11, T14, T15, T16

Results



Mental note
 Noisy neighbor – 20
 Affected processors – P6, P7, P8
 Affected tenants – T1, T2, T3, T4, T5, T6, T8, T9, T10, T11, T14, T15, T16

Peace and quiet

Noisy neighbor

Cooldown

Results



Mental note
 Noisy neighbor – 20
 Affected processors – P6, P7, P8
 Affected tenants – T1, T2, T3, T4, T5, T6, T8, T9, T10, T11, **T14, T15, T16**

Peace and quiet

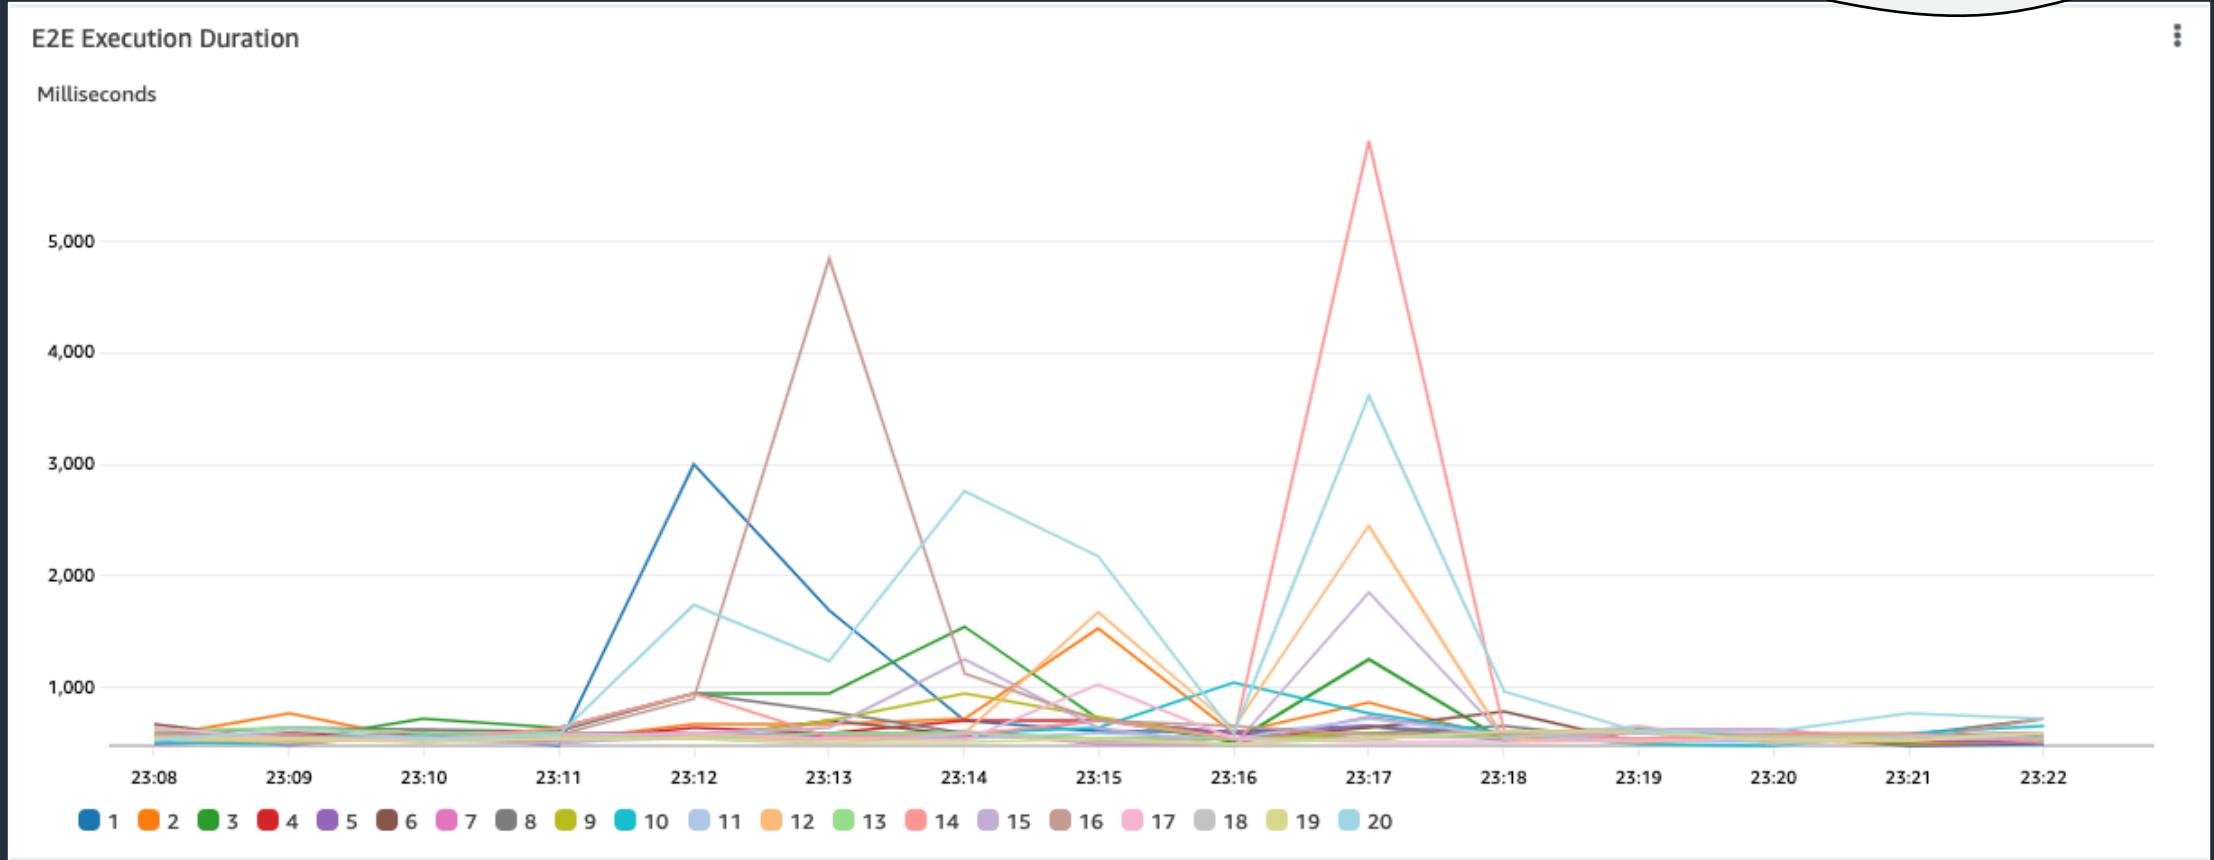
Noisy neighbor

Cooldown

Results

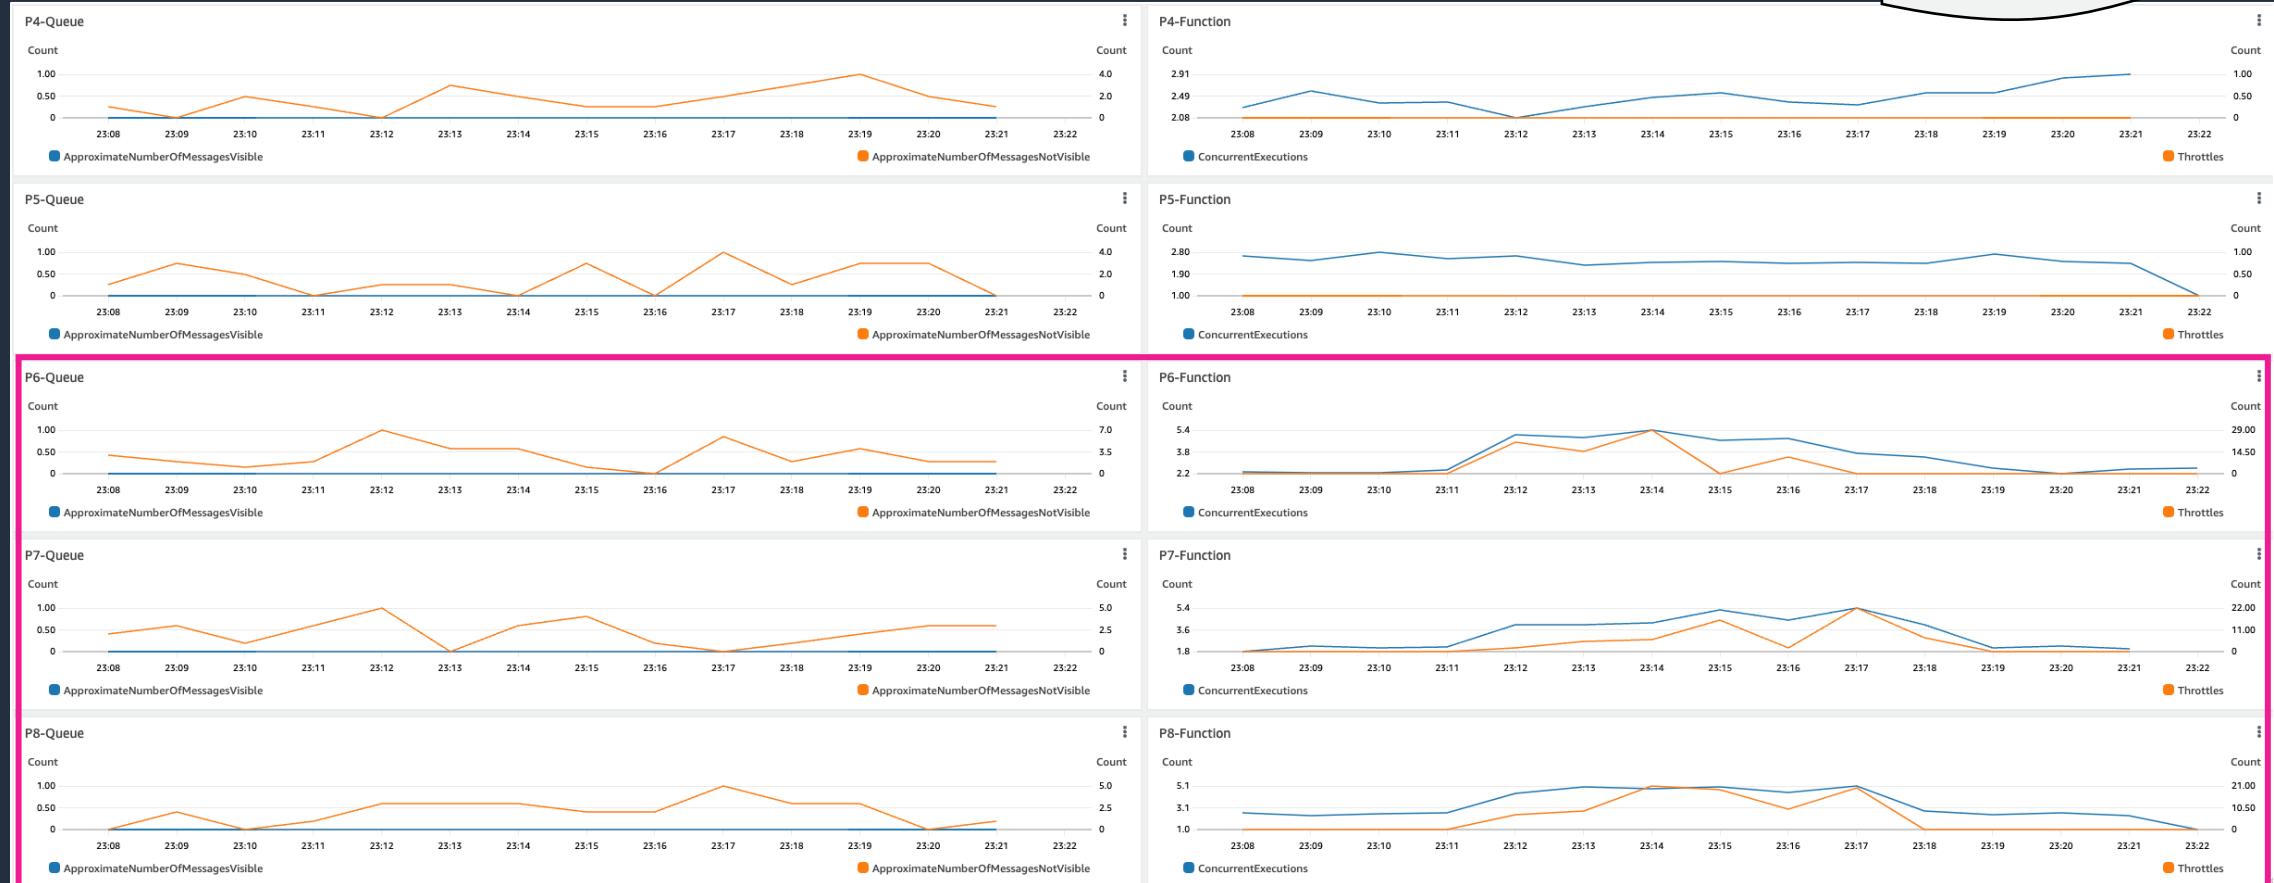
P99

Mental note
 Noisy neighbor – 20
 Affected processors – P6, P7, P8
 Affected tenants – T1, T2, T3, T4, T5, T6, T8, T9, T10, T11, T14, T15, T16



Results

Mental note
 Noisy neighbor – 20
 Affected processors – P6, P7, P8
 Affected tenants – T1, T2, T3, T4, T5, T6, T8, T9, T10, T11, T14, T15, T16



Conclusions

- Shuffle sharding helps by adding more resiliency to the system
- Built for containing blast radius and isolating failures
- Random tenant-to-shard assignment limits customer impact
- Using heuristics for rebalancing can help to build a self-healing system
- Do more with less

Summary

Action items

- Sharding is a great technique for fault isolation and achieving higher workflow resiliency
- Shuffle sharding helps to limit the blast radius exponentially
- Hope for success, build for failure

Additional Shuffle Sharding content

- Workload isolation using shuffle-sharding
<https://tinyurl.com/shufflesharding1>
- Shuffle Sharding: Massive and Magical Fault Isolation
<https://tinyurl.com/shufflesharding2>
- Implement shuffle sharding lab <https://tinyurl.com/shufflesharding3>
- Improving Workload Resiliency Using Shuffle-Sharding
<https://youtu.be/TVdxPF7KL1c>
- How to scale beyond limits with cell-based architectures
<https://youtu.be/HUwz8uko7HY>



Thank you!



Anton Aleksandrov
Specialist Solution Architect, Serverless
AWS